

# Advanced Cluster Programming with MPI

Bill Gropp

*Mathematics and Computer Science Division*

*Argonne National Laboratory*

## Thanks to

---

- Rusty Lusk  
Rob Ross  
Rajeev Thakur

## Tutorial Outline

---

- Background
- Tuning for Performance on Clusters
- Parallel I/O in MPI-2
- Introduction to dynamic process management in MPI-2
- MPI and threads

3

## About This Tutorial

---

- Some of this material is taken from *Using MPI-2*, by Gropp, Lusk, and Thakur, MIT Press, 1999.
- Assumes familiarity with MPI but no deep knowledge.
- Most of the time is spent on those features most widely implemented at this time.
- Language-independent, but most examples are in C.
- Complete MPI-2 specification is on the web <http://www.mpi-forum.org> and in *MPI: The Complete Reference*, 2 volumes from MIT Press.

4

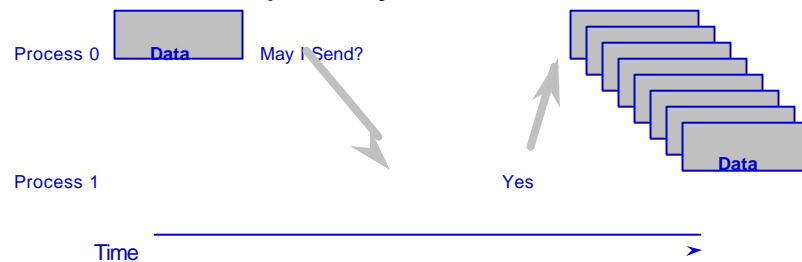
## Contents of MPI-2

- Extensions to the message-passing model
  - **Parallel I/O**
  - One-sided operations
  - **Dynamic process management**
- Making MPI more robust and convenient
  - C++ and Fortran 90 bindings
  - External interfaces, handlers
  - Extended collective operations
  - Language interoperability
  - **MPI interaction with threads**

5

## What is message passing?

- Data transfer plus synchronization



- Requires cooperation of sender and receiver
- Cooperation not always apparent in code

MPI Protocols

## Quick review of MPI Message passing

- Basic terms
  - nonblocking - Operation does not wait for completion
  - synchronous - Completion of send *requires* initiation (but not completion) of receive
  - ready - *Correct* send requires a matching receive
  - asynchronous - communication and computation take place simultaneously, **not** an MPI concept (implementations *may* use asynchronous methods)


## Message protocols

- Message consists of “envelope” and data
  - Envelope contains tag, communicator, length, source information, plus impl. private data
- Short
  - Message data (message for short) sent with envelope
- Eager
  - Message sent assuming destination can store
- Rendezvous
  - Message not sent until destination oks

## Message Protocol Details

- Eager is not Rsend, rendezvous is not Ssend resp., but related
- User versus system buffer space
- Packetization
- Collective operations
- Datatypes, particularly non-contiguous
  - Handling of important special cases
    - Constant stride
    - Contiguous structures

## Eager Protocol




- Data delivered to process 1
  - No matching receive may exist; process 1 must then buffer and copy.

## Eager Features

- Reduces synchronization delays
- Simplifies programming (just MPI\_Send)
- Requires significant buffering
- May require active involvement of CPU to drain network at receiver's end
- May introduce additional copy (buffer to final destination)

## Rendezvous Protocol



- Envelope delivered first
- Data delivered when user-buffer available
  - Only buffering of envelopes required

## Rendezvous Features

---

- Robust and safe
  - (except for limit on the number of envelopes...)
- May remove copy (user to user direct)
- More complex programming (waits/tests)
- May introduce synchronization delays (waiting for receiver to ok send)

## Short Protocol

---

- Data is part of the envelope
- Otherwise like eager protocol
- May be performance optimization in interconnection system for short messages, particularly for networks that send fixed-length packets (or cache lines)

## User and System Buffering

- Where is data stored (or staged) while being sent?
  - User's memory
    - Allocated on the fly
    - Preallocated
  - System memory
    - May be limited
    - Special memory may be faster

## Implementing MPI\_Isend

- Simplest implementation is to always use rendezvous protocol:
  - MPI\_Isend delivers a request-to-send control message to receiver
  - Receiving process responds with an ok-to-send
    - May or may not have matching MPI receive; only needs buffer space to store incoming message
  - Sending process transfers data
- Wait for MPI\_Isend request
  - wait for ok-to-send message from receiver
  - wait for data transfer to be complete on sending side



## Alternatives for MPI\_Isend

- Use a short protocol for small messages
  - No need to exchange control messages
  - Need guaranteed (but small) buffer space on destination for short message envelope
  - Wait becomes a no-op
- Use eager protocol for modest sized messages
  - Still need guaranteed buffer space for both message envelope and eager data on destination
  - Avoids exchange of control messages

17

## Implementing MPI\_Send

- Can't use eager always because this could overwhelm the receiving process

```
if (rank != 0) MPI_Send( 100 MB of data )
else receive 100 MB from each process
```
- Would like to exploit the blocking nature (can wait for receive)
- Would like to be fast
- Select protocol based on message size (and perhaps available buffer space at destination)
  - Short and/or eager for small messages
  - Rendezvous for longer messages

18

## Implementing MPI\_Rsend

---

- Just use MPI\_Send; no advantage for users
- Use eager always (or short if small)
  - even for long messages

19

## Choosing MPI Alternatives

---

- MPI offers many ways to accomplish the same task
- Which is best?
  - Just like everything else, it depends on the vendor, system architecture
  - Like C and Fortran, MPI provides the programmer with the tools to achieve high performance without sacrificing portability
- The best choice depends on the use:
  - Consider choices based on system and MPI implementation
  - Example: Experiments with a Jacobi relaxation example

## Tuning for MPI's Send/Receive Protocols

- Aggressive Eager
  - Performance problem: extra copies
  - Possible deadlock for inadequate eager buffering
  - Ensure that receives are posted before sends
  - MPI\_Issend can be used to express “wait until receive is posted”
- Rendezvous with sender push
  - Extra latency
  - Possible delays while waiting for sender to begin
- Rendezvous with receiver pull
  - Possible delays while waiting for receiver to begin

21

## Rendezvous Blocking

- What happens once sender and receiver rendezvous?
  - Sender (push) or receiver (pull) may complete operation
  - May block other operations while completing
- Performance tradeoff
  - If operation does *not* block (by checking for other requests), it adds latency or reduces bandwidth.
- Can reduce performance if a receiver, having acknowledged a send, must wait for the sender to complete a separate operation that it has started.

22

## Tuning for Rendezvous with Sender Push

- Ensure receives posted before sends
  - better, ensure receives match sends before computation starts; may be better to do sends before receives
- Ensure that sends have time to start transfers
- Can use short control messages
- Beware of the cost of extra messages
  - Intel i860 encouraged use of control messages with ready send (force type)

23

## Tuning for Rendezvous with Receiver Pull

- Place MPI\_Isends before receives
- Use short control messages to ensure matches
- Beware of the cost of extra messages

24


## Experiments with MPI Implementations

- Multiparty data exchange
- Jacobi iteration in 2 dimensions
  - Model for PDEs, Sparse matrix-vector products, and algorithms with surface/volume behavior
  - Issues are similar to unstructured grid problems (but harder to illustrate)
- Others at <http://www.mcs.anl.gov/mpi/tutorials/perf>

25

## Jacobi Iteration (Fortran Ordering)

- Simple parallel data structure



- Processes exchange columns with neighbors
- Local part declared as `xlocal(m,0:n+1)`

27

## Background to Tests

- Goals
  - Identify better performing idioms for the same communication operation
  - Understand these by understanding the underlying MPI process
  - Provide a starting point for evaluating additional options (there are many ways to write even simple codes)

28

## Some Send/Receive Approaches

- Based on operation hypothesis. Most of these are for polling mode. Each of the following is a *hypothesis* that the experiments test
  - Better to start receives first
  - Ensure recvs posted before sends
  - Ordered (no overlap)
- More details at <http://www.mcs.anl.gov/mpi/tutorial/perf/mpiexmpl/src3/runs.html>

29

## Send and Recv

- Simplest use of send and recv

integer status(MPI\_STATUS\_SIZE)

```
call MPI_Send( xlocal(1,1), m, MPI_DOUBLE_PRECISION, &  
              left_nbr, 0, ring_comm, ierr )
```

```
call MPI_Recv( xlocal(1,0), m, MPI_DOUBLE_PRECISION, &  
              right_nbr, 0, ring_comm, status, ierr )
```

```
call MPI_Send( xlocal(1,n), m, MPI_DOUBLE_PRECISION, &  
              right_nbr, 0, ring_comm, ierr )
```

```
call MPI_Recv( xlocal(1,n+1), m, MPI_DOUBLE_PRECISION, &  
              left_nbr, 0, ring_comm, status, ierr )
```

31

## Performance of Simplest Code

- Very poor performance for large m
  - Rendezvous sequentializes sends/receives
- Good to OK performance for modest m
  - eager operations

32

## Better to start sends first

- Irecv, Isend, Waitall - ok performance

integer statuses(MPI\_STATUS\_SIZE,4), requests(4)

```
call MPI_Isend( xlocal(1,n), m, MPI_DOUBLE_PRECISION, &
               right_nbr, ring_comm, requests(1), ierr )
call MPI_Isend( xlocal(1,1), m, MPI_DOUBLE_PRECISION, &
               left_nbr, ring_comm, requests(3), ierr )
call MPI_Irecv( xlocal(1,0), m, MPI_DOUBLE_PRECISION, &
               left_nbr, ring_comm, requests(2), ierr )
call MPI_Irecv( xlocal(1,n+1), m, MPI_DOUBLE_PRECISION, &
               right_nbr, ring_comm, requests(4), ierr )
call MPI_Waitall( 4, requests, statuses, ierr )
```

33

## Summary of Results

- Better to start sends before receives
  - Most implementations use rendezvous protocols for long messages (Cray, IBM, SGI, MPICH, LAM)
  - Sending the envelope immediately is the best way to improve performance
- Short messages even better
  - Eager messages have lowest latency

34



## MPI I/O

- Goals of this part
  - introduce the important features of MPI I/O in the form of example programs, following the outline of the Parallel I/O chapter in *Using MPI-2*
  - focus on how to achieve high performance
- What can you expect from this session?
  - learn how to use MPI I/O and, hopefully, like it
  - be able to go back home and immediately use MPI I/O in your applications
  - get much higher I/O performance than what you have been getting so far using other techniques

35

## MPI I/O

- Why do I/O in MPI?
  - Why not just Posix?
    - Performance
    - Single file output (instead of one file / process)
- Example of MPI I/O: Non-parallel I/O from an MPI program
  - Emphasizes similarities with conventional I/O
  - High performance parallel I/O later

36


## Why MPI is a Good Setting for Parallel I/O

- Writing is like sending and reading is like receiving.
- Any parallel I/O system will need:
  - collective operations
  - user-defined datatypes to describe both memory and file layout
  - communicators to separate application-level message passing from I/O-related message passing
  - non-blocking operations
- I.e., lots of MPI-like machinery

37

## Introduction to Parallel I/O

- First example: non-parallel I/O from an MPI program



38

## Next few examples have:

```
#include "mpi.h"
#include <stdio.h>
#define BUFSIZE 1000

int main(int argc, char *argv[])
{
    int i, myrank, numprocs, buf[BUFSIZE];

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    ....
    ....
    MPI_Finalize();
    return 0;
}
```

39

## Non-Parallel I/O from MPI Program

```
MPI_Status status;
FILE *myfile;

for (i=0; i<BUFSIZE; i++)
    buf[i] = myrank * BUFSIZE + i;
if (myrank != 0)
    MPI_Send(buf, BUFSIZE, MPI_INT, 0, 99,
             MPI_COMM_WORLD);
else {
    myfile = fopen("testfile", "w");
    fwrite(buf, sizeof(int), BUFSIZE, myfile);
    for (i=1; i<numprocs; i++) {
        MPI_Recv(buf, BUFSIZE, MPI_INT, i, 99,
                 MPI_COMM_WORLD, &status);
        fwrite(buf, sizeof(int), BUFSIZE, myfile);
    }
    fclose(myfile);
}
```

40


## Pros and Cons of Sequential I/O

- Pros:
  - parallel machine may support I/O from only one process
  - I/O libraries (e.g. HDF-4, SILO, etc.) not parallel
  - resulting single file is handy for `ftp`, `mv`
  - big blocks improve performance
  - short distance from original, serial code
- Cons:
  - lack of parallelism limits scalability, performance

41

## Non-MPI Parallel I/O

- Each process writes to a separate file



- Pro: parallelism
- Con: lots of small files to manage

42

## Non-MPI Parallel I/O

---

```
char filename[128];
FILE *myfile;

for (i=0; i<BUFSIZE; i++)
    buf[i] = myrank * BUFSIZE + i;
sprintf(filename, "testfile.%d", myrank);
myfile = fopen(filename, "w");
fwrite(buf, sizeof(int), BUFSIZE, myfile);
fclose(myfile);
```

43

## MPI I/O to Separate Files

---

- Same pattern as previous example
- MPI I/O replaces Unix I/O in a straightforward way
- Easy way to start with MPI I/O
- Does not exploit advantages of MPI I/O
  - in producing single file
  - in allowing collective operations
- Note files cannot be read conveniently by a different number of processes

44

## MPI I/O to Separate Files

```
char filename[128];
MPI_File myfile;

for (i=0; i<BUFSIZE; i++)
    buf[i] = myrank * BUFSIZE + i;
sprintf(filename, "testfile.%d", myrank);
MPI_File_open(MPI_COMM_SELF, filename,
               MPI_MODE_WRONLY | MPI_MODE_CREATE,
               MPI_INFO_NULL, &myfile);
MPI_File_write(myfile, buf, BUFSIZE, MPI_INT,
               MPI_STATUS_IGNORE);
MPI_File_close(&myfile);
```

45

## MPI Versions of Unix I/O

- Unix

```
FILE myfile;
myfile =
    fopen(...)
```

```
fread(...)
fwrite(...)
```

```
fclose
```

- MPI

```
MPI_File myfile;
MPI_File_open(...)
    takes info, comm
    args
```

```
MPI_File_read/write(.
..) take (addr,
count, datatype)
```

```
MPI_File_close
```

46

## Parallel I/O in MPI

- MPI provides effective ways to describe and perform high-performance parallel I/O
  - Requires specifying *all* data to move
  - Natural once you get used to it

47

## Beyond POSIX I/O with MPI

- Why do I/O in MPI?
- non-parallel I/O from an MPI program
- non-MPI parallel I/O to separate files
- parallel I/O to shared file with MPI I/O
- Fortran-90 version
- Reading a file with a different number of processes
- C++ version
- Survey of advanced features in MPI I/O


48

## Parallel I/O with MPI

- Nonparallel I/O shown earlier
  - Simple but
    - Poor performance (single process writes to one file) or
    - Awkward and not interoperable with other tools (each process writes a separate file)
- Parallel I/O
  - Provides high performance
  - Can provide a single file that can be used with other tools (such as visualization programs)

49

## Using MPI for Simple I/O



Each process needs to read a chunk of data from a common file

50



## Using Individual File Pointers

```
MPI_File fh;
MPI_Status status;

MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &nprocs);

bufsize = FILESIZE/nprocs;
nints = bufsize/sizeof(int);

MPI_File_open(MPI_COMM_WORLD, "/pfs/datafile",
              MPI_MODE_RDONLY, MPI_INFO_NULL, &fh);
MPI_File_seek(fh, rank * bufsize, MPI_SEEK_SET);
MPI_File_read(fh, buf, nints, MPI_INT, &status);
MPI_File_close(&fh);
```

51

## Using Explicit Offsets

```
include 'mpif.h'

integer status(MPI_STATUS_SIZE)
integer (kind=MPI_OFFSET_KIND) offset
C in F77, see implementation notes (might be integer*8)

call MPI_FILE_OPEN(MPI_COMM_WORLD, '/pfs/datafile', &
                  MPI_MODE_RDONLY, MPI_INFO_NULL, fh, ierr)
nints = FILESIZE / (nprocs*INTSIZE)
offset = rank * nints * INTSIZE
call MPI_FILE_READ_AT(fh, offset, buf, nints,
                    MPI_INTEGER, status, ierr)
call MPI_GET_COUNT(status, MPI_INTEGER, count, ierr)
print *, 'process ', rank, 'read ', count, 'integers'

call MPI_FILE_CLOSE(fh, ierr)
```

52


## Writing to a File

- Use `MPI_File_write` or `MPI_File_write_at`
- Use `MPI_MODE_WRONLY` or `MPI_MODE_RDWR` as the flags to `MPI_File_open`
- If the file doesn't exist previously, the flag `MPI_MODE_CREATE` must also be passed to `MPI_File_open`
- We can pass multiple flags by using bitwise-or '|' in C, or addition '+' in Fortran

53

## MPI Parallel I/O to Single File

- Processes write to shared file



- `MPI_File_set_view` assigns regions of the file to separate processes

54

## File Views

- Specified by a triplet (*displacement*, *etype*, and *filetype*) passed to **MPI\_File\_set\_view**
- *displacement* = number of bytes to be skipped from the start of the file
- *etype* = basic unit of data access (can be any basic or derived datatype)
- *filetype* = specifies which portion of the file is visible to the process

55

## MPI Parallel I/O to Single File

```
MPI_File thefile;

for (i=0; i<BUFSIZE; i++)
    buf[i] = myrank * BUFSIZE + i;
MPI_File_open(MPI_COMM_WORLD, "testfile",
               MPI_MODE_CREATE | MPI_MODE_WRONLY,
               MPI_INFO_NULL, &thefile);
MPI_File_set_view(thefile, myrank * BUFSIZE * sizeof(int),
                  MPI_INT, MPI_INT, "native",
                  MPI_INFO_NULL);
MPI_File_write(thefile, buf, BUFSIZE, MPI_INT,
               MPI_STATUS_IGNORE);
MPI_File_close(&thefile);
```

56

## MPI\_File\_set\_view

- Describes that part of the file accessed by a single MPI process.
- Arguments to `MPI_File_set_view`:
  - `MPI_File` file
  - `MP_Offset` disp
  - `MPI_Datatype` etype
  - `MPI_Datatype` filetype
  - `char *datarep`
  - `MPI_Info` info

57

## MPI I/O in Fortran

```
PROGRAM main

use mpi

integer ierr, i, myrank, BUFSIZE, thefile
parameter (BUFSIZE=100)
integer buf(BUFSIZE)
integer(kind=MPI_OFFSET_KIND) disp

call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierr)
do i = 0, BUFSIZE
    buf(i) = myrank * BUFSIZE + i
enddo

* in F77, see implementation notes (might be integer*8)
```

58

## MPI I/O in Fortran contd.

---

```
call MPI_FILE_OPEN(MPI_COMM_WORLD, 'testfile', &
                   MPI_MODE_WRONLY + MPI_MODE_CREATE, &
                   MPI_INFO_NULL, thefile, ierr)
call MPI_TYPE_SIZE(MPI_INTEGER, intsize)
disp = myrank * BUFSIZE * intsize
call MPI_FILE_SET_VIEW(thefile, disp, MPI_INTEGER, &
                      MPI_INTEGER, 'native', &
                      MPI_INFO_NULL, ierr)
call MPI_FILE_WRITE(thefile, buf, BUFSIZE, MPI_INTEGER, &
                   MPI_STATUS_IGNORE, ierr)
call MPI_FILE_CLOSE(thefile, ierr)
call MPI_FINALIZE(ierr)

END PROGRAM main
```

59

## C++ Version

---

```
// example of parallel MPI read from single file
#include <iostream.h>
#include "mpi.h"

int main(int argc, char *argv[])
{
    int bufsize, *buf, count;
    char filename[128];
    MPI::Status status;

    MPI::Init();
    int myrank = MPI::COMM_WORLD.Get_rank();
    int numprocs = MPI::COMM_WORLD.Get_size();
    MPI::File thefile = MPI::File::Open(MPI::COMM_WORLD,
                                         "testfile",
                                         MPI::MODE_RDONLY,
                                         MPI::INFO_NULL); 60
```

## C++ Version, Part 2

```
MPI::Offset filesize = thefile.Get_size();
filesize      = filesize / sizeof(int);
bufsize      = filesize / numprocs + 1;
buf = new int[bufsize];
thefile.Set_view(myrank * bufsize * sizeof(int),
                 MPI_INT, MPI_INT, "native",
                 MPI::INFO_NULL);
thefile.Read(buf, bufsize, MPI_INT, &status);
count = status.Get_count(MPI_INT);
cout << "process " << myrank << " read " << count
      << " ints" << endl;
thefile.Close();
delete [] buf;
MPI::Finalize();
return 0;
}
```

61


## Other Ways to Write to a Shared File

- `MPI_File_seek` } like Unix seek
- `MPI_File_read_at` }
- `MPI_File_write_at` } combine seek and I/O  
for thread safety
- `MPI_File_read_shared` }
- `MPI_File_write_shared` } use shared file pointer
- Collective operations

62

## Collective I/O in MPI

- A critical optimization in parallel I/O
- Allows communication of “big picture” to file system
- Framework for 2-phase I/O, in which communication precedes I/O (can use MPI machinery)
- Basic idea: build large blocks, so that reads/writes in I/O system will be large



63


## Noncontiguous Accesses

- Common in parallel applications
- Example: distributed arrays stored in files
- A big advantage of MPI I/O over Unix I/O is the ability to specify noncontiguous accesses in memory and file within a single function call by using derived datatypes
- Allows implementation to optimize the access
- Collective IO combined with noncontiguous accesses yields the highest performance.

64

## Example: Distributed Array Access


2D array distributed among four processes

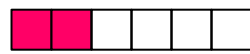



File containing the global array in row-major order

65

## A Simple File View Example

 etype = MPI\_INT

 filetype = two MPI\_INTs followed by  
a gap of four MPI\_INTs



66



## File View Code

```
MPI_Aint lb, extent;
MPI_Datatype etype, filetype, contig;
MPI_Offset disp;

MPI_Type_contiguous(2, MPI_INT, &contig);
lb = 0; extent = 6 * sizeof(int);
MPI_Type_create_resized(contig, lb, extent, &filetype);
MPI_Type_commit(&filetype);
disp = 5 * sizeof(int); etype = MPI_INT;

MPI_File_open(MPI_COMM_WORLD, "/pfs/datafile",
              MPI_MODE_CREATE | MPI_MODE_RDWR, MPI_INFO_NULL, &fh);
MPI_File_set_view(fh, disp, etype, filetype, "native",
                  MPI_INFO_NULL);
MPI_File_write(fh, buf, 1000, MPI_INT, MPI_STATUS_IGNORE);
```

67

## Collective I/O

- `MPI_File_read_all`,  
`MPI_File_read_at_all`, etc
- `_all` indicates that all processes in the group specified by the communicator passed to `MPI_File_open` will call this function
- Each process specifies only its own access information -- the argument list is the same as for the non-collective functions


68

## Collective I/O

- By calling the collective I/O functions, the user allows an implementation to optimize the request based on the combined request of all processes
- The implementation can merge the requests of different processes and service the merged request efficiently
- Particularly effective when the accesses of different processes are noncontiguous and interleaved

69

## Accessing Arrays Stored in Files



$nproc(1) = 2, nproc(2) = 3$

70

## Using the “Distributed Array” (Darray) Datatype

```
int gsizes[2], distribs[2], dargs[2], psize[2];

gsizes[0] = m;    /* no. of rows in global array */
gsizes[1] = n;    /* no. of columns in global array*/

distribs[0] = MPI_DISTRIBUTE_BLOCK;
distribs[1] = MPI_DISTRIBUTE_BLOCK;

dargs[0] = MPI_DISTRIBUTE_DFLT_DARG;
dargs[1] = MPI_DISTRIBUTE_DFLT_DARG;

psize[0] = 2; /* no. of processes in vertical dimension
              of process grid */
psize[1] = 3; /* no. of processes in horizontal dimension
              of process grid */
```

71

## Darray Continued

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Type_create_darray(6, rank, 2, gsizes, distribs, dargs,
                      psize, MPI_ORDER_C, MPI_FLOAT, &filetype);
MPI_Type_commit(&filetype);

MPI_File_open(MPI_COMM_WORLD, "/pfs/datafile",
              MPI_MODE_CREATE | MPI_MODE_WRONLY,
              MPI_INFO_NULL, &fh);

MPI_File_set_view(fh, 0, MPI_FLOAT, filetype, "native",
                  MPI_INFO_NULL);

local_array_size = num_local_rows * num_local_cols;
MPI_File_write_all(fh, local_array, local_array_size,
                  MPI_FLOAT, &status);

MPI_File_close(&fh);
```

72

## A Word of Warning about Darray

- The darray datatype assumes a very specific definition of data distribution -- the exact definition as in HPF
- For example, if the array size is not divisible by the number of processes, darray calculates the block size using a *ceiling* division ( $20 / 6 = 4$ )
- darray assumes a row-major ordering of processes in the logical grid, as assumed by cartesian process topologies in MPI-1
- If your application uses a different definition for data distribution or logical grid ordering, you cannot use darray. Use subarray instead.

73

## Using the Subarray Datatype

```
gsizes[0] = m; /* no. of rows in global array */
gsizes[1] = n; /* no. of columns in global array*/

psizes[0] = 2; /* no. of procs. in vertical dimension */
psizes[1] = 3; /* no. of procs. in horizontal dimension */

lsizes[0] = m/psizes[0]; /* no. of rows in local array */
lsizes[1] = n/psizes[1]; /* no. of columns in local array */

dims[0] = 2; dims[1] = 3;
periods[0] = periods[1] = 1;
MPI_Cart_create(MPI_COMM_WORLD, 2, dims, periods, 0, &comm);
MPI_Comm_rank(comm, &rank);
MPI_Cart_coords(comm, rank, 2, coords);
```

74

## Subarray Datatype contd.

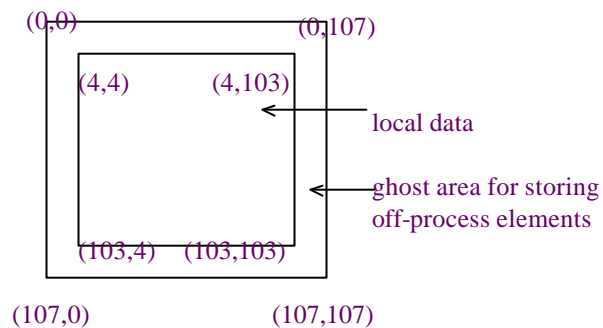
```
/* global indices of first element of local array */
start_indices[0] = coords[0] * lsizes[0];
start_indices[1] = coords[1] * lsizes[1];

MPI_Type_create_subarray(2, gsizes, lsizes, start_indices,
                        MPI_ORDER_C, MPI_FLOAT, &filetype);
MPI_Type_commit(&filetype);

MPI_File_open(MPI_COMM_WORLD, "/pfs/datafile",
              MPI_MODE_CREATE | MPI_MODE_WRONLY,
              MPI_INFO_NULL, &fh);
MPI_File_set_view(fh, 0, MPI_FLOAT, filetype, "native",
                  MPI_INFO_NULL);
local_array_size = lsizes[0] * lsizes[1];
MPI_File_write_all(fh, local_array, local_array_size,
                  MPI_FLOAT, &status);
```

75

## Local Array with Ghost Area in Memory



- Use a subarray datatype to describe the noncontiguous layout in memory
- Pass this datatype as argument to `MPI_File_write_all`

76

## Local Array with Ghost Area

```
memsizes[0] = lsizes[0] + 8;
/* no. of rows in allocated array */
memsizes[1] = lsizes[1] + 8;
/* no. of columns in allocated array */
start_indices[0] = start_indices[1] = 4;
/* indices of the first element of the local array
   in the allocated array */

MPI_Type_create_subarray(2, memsizes, lsizes,
                        start_indices, MPI_ORDER_C, MPI_FLOAT, &memtype);
MPI_Type_commit(&memtype);

/* create filetype and set file view exactly as in the
   subarray example */

MPI_File_write_all(fh, local_array, 1, memtype, &status);
```

77

## Accessing Irregularly Distributed Arrays

Process 0's data array




Process 1's data array




Process 2's data array




Process 0's map array



Process 1's map array



Process 2's map array



The map array describes the location of each element  
of the data array in the common file

78

## Accessing Irregularly Distributed Arrays

```
integer (kind=MPI_OFFSET_KIND) disp

call MPI_FILE_OPEN(MPI_COMM_WORLD, '/pfs/datafile', &
                   MPI_MODE_CREATE + MPI_MODE_RDWR, &
                   MPI_INFO_NULL, fh, ierr)

call MPI_TYPE_CREATE_INDEXED_BLOCK(bufsize, 1, map, &
                                   MPI_DOUBLE_PRECISION, filetype, ierr)
call MPI_TYPE_COMMIT(filetype, ierr)
disp = 0
call MPI_FILE_SET_VIEW(fh, disp, MPI_DOUBLE_PRECISION, &
                      filetype, 'native', MPI_INFO_NULL, ierr)

call MPI_FILE_WRITE_ALL(fh, buf, bufsize, &
                       MPI_DOUBLE_PRECISION, status, ierr)

call MPI_FILE_CLOSE(fh, ierr)
```

79

## Nonblocking I/O

```
MPI_Request request;
MPI_Status status;

MPI_File_iwrite_at(fh, offset, buf, count, datatype,
                  &request);

for (i=0; i<1000; i++) {
    /* perform computation */
}

MPI_Wait(&request, &status);
```

80

## Split Collective I/O

- A restricted form of nonblocking collective I/O
- Only one active nonblocking collective operation allowed at a time on a file handle
- Therefore, no request object necessary

```
MPI_File_write_all_begin(fh, buf, count, datatype);

for (i=0; i<1000; i++) {
    /* perform computation */
}

MPI_File_write_all_end(fh, buf, &status);
```

81

## Shared File Pointers

```
#include "mpi.h"
// C++ example
int main(int argc, char *argv[])
{
    int buf[1000];
    MPI::File fh;

    MPI::Init();

    MPI::File fh = MPI::File::Open(MPI::COMM_WORLD,
        "/pfs/datafile", MPI::MODE_RDONLY, MPI::INFO_NULL);
    fh.Write_shared(buf, 1000, MPI_INT);
    fh.Close();

    MPI::Finalize();
    return 0;
}
```

82



## Passing Hints to the Implementation

```
MPI_Info info;

MPI_Info_create(&info);

/* no. of I/O devices to be used for file striping */
MPI_Info_set(info, "striping_factor", "4");

/* the striping unit in bytes */
MPI_Info_set(info, "striping_unit", "65536");

MPI_File_open(MPI_COMM_WORLD, "/pfs/datafile",
              MPI_MODE_CREATE | MPI_MODE_RDWR, info, &fh);

MPI_Info_free(&info);
```

83

## Examples of Hints (used in ROMIO)

- **striping\_unit**
  - **striping\_factor**
  - **cb\_buffer\_size**
  - **cb\_nodes**
  - **ind\_rd\_buffer\_size**
  - **ind\_wr\_buffer\_size**
  - **start\_iodevice**
  - **pfs\_svr\_buf**
  - **direct\_read**
  - **direct\_write**
- } MPI-2 predefined hints
- } New Algorithm Parameters
- } Platform-specific hints

84

## I/O Consistency Semantics

- The consistency semantics specify the results when multiple processes access a common file and one or more processes write to the file
- MPI guarantees stronger consistency semantics if the communicator used to open the file accurately specifies all the processes that are accessing the file, and weaker semantics if not
- The user can take steps to ensure consistency when MPI does not automatically do so
- **Warning:** NFS (Network File System) does not support access from multiple processes. Clusters should use PVFS (Parallel Virtual File System) instead.

85

## File Interoperability

- Users can optionally create files with a portable binary data representation
- “datarep” parameter to `MPI_File_set_view`
- `native` - default, same as in memory, not portable
- `internal` - impl. defined representation providing an impl. defined level of portability
- `external32` - a specific representation defined in MPI, (basically 32-bit big-endian IEEE format), portable across machines and MPI implementations

86

## General Guidelines for Achieving High I/O Performance

---

- Buy sufficient I/O hardware for the machine
- Use fast file systems (such as PVFS), not NFS-mounted home directories
- Do not perform I/O from one process only
- Make large requests wherever possible
- For noncontiguous requests, use derived datatypes and a single collective I/O call

87

## Achieving High I/O Performance with MPI

---

- Any application as a particular “I/O access pattern” based on its I/O needs
- The same access pattern can be presented to the I/O system in different ways depending on what I/O functions are used and how
- See the SC98 paper in which we classify the different ways of expressing I/O access patterns in MPI-IO into four *levels*: level 0 — level 3  
([http://www.supercomp.org/sc98/TechPapers/sc98\\_FullAbstracts/Thakur447](http://www.supercomp.org/sc98/TechPapers/sc98_FullAbstracts/Thakur447))
- We demonstrate how the user’s choice of *level* affects performance

88


## Example: Distributed Array Access

Large array  
distributed  
among 16  
processes

P0	P1	P2	P3
P4	P5	P6	P7
P8	P9	P10	P11
P12	P13	P14	P15

Each square represents  
a subarray in the memory  
of a single process

Access Pattern in the file



89

## Level-0 Access (C)

- Each process makes one independent read request for each row in the local array (as in Unix)

```
MPI_File_open(..., file, ..., &fh)
for (i=0; i<n_local_rows; i++) {
    MPI_File_seek(fh, ...);
    MPI_File_read(fh, &(A[i][0]), ...);
}
MPI_File_close(&fh);
```

90

## Level-0 Access (Fortran)

---

- Each process makes one independent read request for each column in the local array (as in Unix)

```
call MPI_File_open(..., file, ..., fh, ierr)
do i=1, n_local_cols
    call MPI_File_seek(fh, ..., ierr)
    call MPI_File_read(fh, A(1,i), ..., ierr)
enddo
call MPI_File_close(fh, ierr)
```

91

## Level-1 Access (C)

---

- Similar to level 0, but each process uses collective I/O functions

```
MPI_File_open(MPI_COMM_WORLD, file, ...,
&fh);
for (i=0; i<n_local_rows; i++) {
    MPI_File_seek(fh, ...);
    MPI_File_read_all(fh, &(A[i][0]), ...);
}
MPI_File_close(&fh);
```

92

## Level-1 Access (Fortran)

- Similar to level 0, but each process uses collective I/O functions

```
call MPI_File_open(MPI_COMM_WORLD, file,  
..., fh, ierr)  
do i=1, n_local_cols  
    call MPI_File_seek(fh, ..., ierr)  
    call MPI_File_read_all(fh, A(1,i), ...,  
ierr)  
enddo  
call MPI_File_close(fh, ierr)
```

93

## Level-2 Access (C)

- Each process creates a derived datatype to describe the noncontiguous access pattern, defines a file view, and calls independent I/O functions

```
MPI_Type_create_subarray(..., &subarray,  
...);  
MPI_Type_commit(&subarray);  
MPI_File_open(..., file, ..., &fh);  
MPI_File_set_view(fh, ..., subarray, ...);  
MPI_File_read(fh, A, ...);  
MPI_File_close(&fh);
```

94

## Level-2 Access (Fortran)

- Each process creates a derived datatype to describe the noncontiguous access pattern, defines a file view, and calls independent I/O functions

```
call MPI_Type_create_subarray(..., subarray,  
..., ierr)  
call MPI_Type_commit(subarray, ierr)  
call MPI_File_open(..., file, ..., fh, ierr)  
call MPI_File_set_view(fh, ..., subarray,  
..., ierr)  
call MPI_File_read(fh, A, ..., ierr)  
call MPI_File_close(fh, ierr)
```

95

## Level-3 Access (C)

- Similar to level 2, except that each process uses collective I/O functions

```
MPI_Type_create_subarray(...,  
&subarray, ...);  
MPI_Type_commit(&subarray);  
MPI_File_open(MPI_COMM_WORLD, file,  
..., &fh);  
MPI_File_set_view(fh, ..., subarray,  
...);  
MPI_File_read_all(fh, A, ...);  
MPI_File_close(&fh);
```

96


## Level-3 Access (Fortran)

- Similar to level 2, except that each process uses collective I/O functions

```
call MPI_Type_create_subarray(...,  
subarray, ..., ierr)  
call MPI_Type_commit(subarray, ierr)  
call MPI_File_open(MPI_COMM_WORLD, file,  
..., fh, ierr)  
call MPI_File_set_view(fh, ..., subarray,  
..., ierr)  
call MPI_File_read_all(fh, A, ..., ierr)  
call MPI_File_close(fh, ierr)
```

97

## The Four Levels of Access



98



## Optimizations

---

- Given complete access information, an implementation can perform optimizations such as:
  - Data Sieving: Read large chunks and extract what is really needed
  - Collective I/O: Merge requests of different processes into larger requests
  - Improved prefetching and caching

99

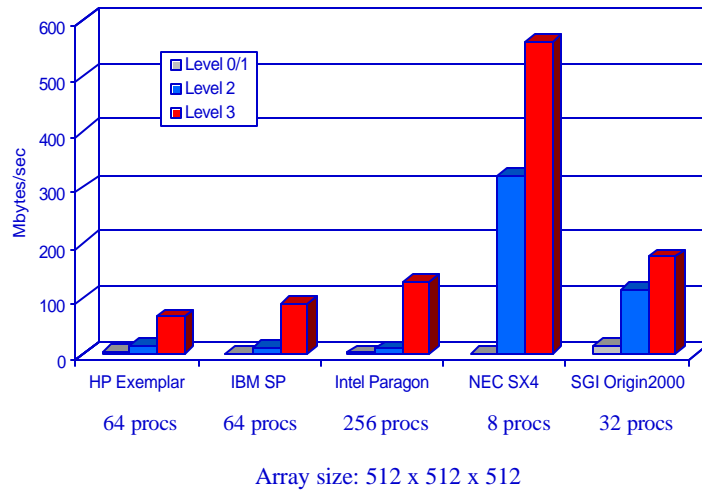
## Performance Results

---

- Distributed array access
- Unstructured code from Sandia
- On five different parallel machines:
  - HP Exemplar
  - IBM SP
  - Intel Paragon
  - NEC SX-4
  - SGI Origin2000

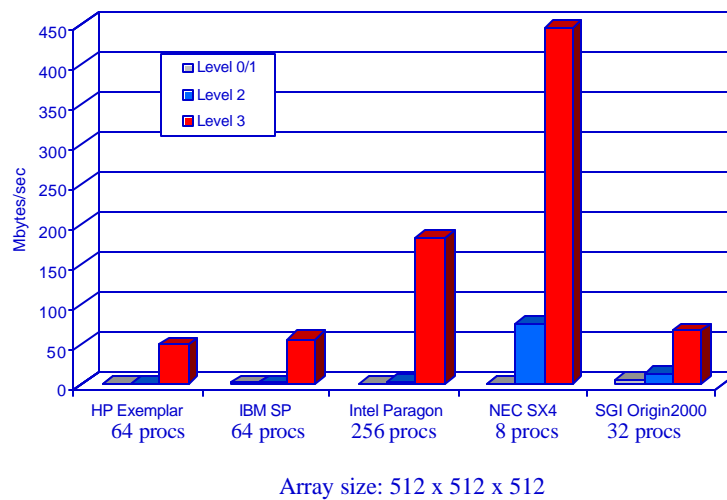
100

## Distributed Array Access: Read Bandwidth



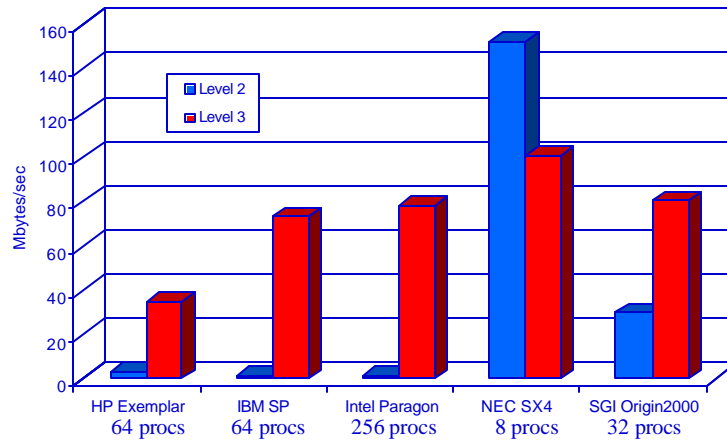
101

## Distributed Array Access: Write Bandwidth




102

## Unstructured Code: Read Bandwidth



103


## Unstructured Code: Write Bandwidth



104

## Independent Writes

- On Paragon
- Lots of seeks and small writes
- Time shown = 130 seconds




## Collective Write

- On Paragon
- Computation and communication precede seek and write
- Time shown = 2.75 seconds



## Independent Writes with Data Sieving

- On Paragon
- Access data in large “blocks” and extract needed data
- Requires lock, read, modify, write, unlock for writes
- 4 MB blocks
- Time = 16 sec.



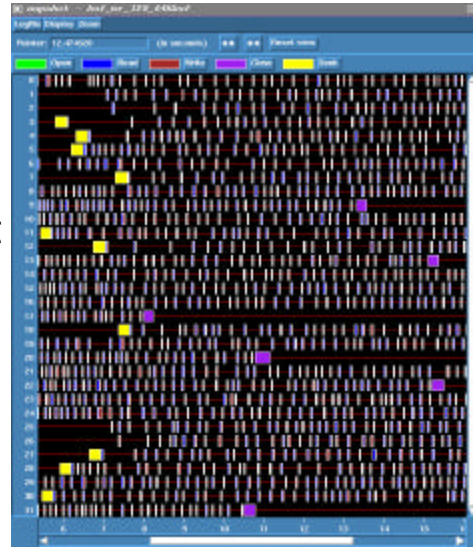
## Changing the Block Size

- Smaller blocks mean less contention, therefore more parallelism
- 512 KB blocks
- Time = 10.2 seconds



## Data Sieving with Small Blocks

- If the block size is too small, however, the increased parallelism doesn't make up for the many small writes
- 64 KB blocks
- Time = 21.5 seconds



## Common Errors

- Not defining file offsets as `MPI_Offset` in C and `integer (kind=MPI_OFFSET_KIND)` in Fortran (or perhaps `integer*8` in Fortran 77)
- In Fortran, passing the offset or displacement directly as a constant (e.g., 0) in the absence of function prototypes (F90 `mpi` module)
- Using `darray` datatype for a block distribution other than the one defined in `darray` (e.g., floor division)
- filetype defined using offsets that are not monotonically nondecreasing, e.g., 0, 3, 8, 4, 6. (happens in irregular applications)

110

## Summary

- MPI I/O has many features that can help users achieve high performance
- The most important of these features are the ability to specify noncontiguous accesses, the collective I/O functions, and the ability to pass hints to the implementation
- Users must use the above features!
- In particular, when accesses are noncontiguous, users must create derived datatypes, define file views, and use the collective I/O functions

111

## Dynamic Process Management in MPI

- Standard way of starting processes in PVM
- Not so necessary in MPI
- Useful in assembling complex distributed applications
- Issues
  - maintaining simplicity, flexibility, and correctness
  - interaction with operating system, resource manager, and process manager
  - connecting independently started processes

112

## Starting New MPI Processes

- MPI\_Comm\_spawn
  - Starts n new processes
  - Collective over communicator
    - Necessary for scalability
  - Returns an intercommunicator
    - Does *not* change MPI\_COMM\_WORLD
  - “SPMD” (Single Program Multiple Data)
- MPI\_Comm\_spawn\_multiple
  - Link MPI\_Comm\_spawn
  - “MPMD” (Multiple Program Multiple Data)

113

## The MPI Environment

- MPI does not specify a particular process management environment
  - You can use whatever you have (Condor, LoadLeveler, PBS, etc.)
  - No particular system is defined
    - MPI Forum attempted to define an abstract interface to third-party process managers, but could no consensus reached
- Still need to pass information to process management system
  - Led to MPI\_Info

114



## Communicating Data to (and through) MPI

- MPI-2 defines a new object, *MPI\_Info*
- Provides an extensible list of key=value pairs
- Used in I/O, One-sided, and Dynamic to package variable, optional types of arguments that may not be standard

115

## Using MPI\_Info in Fortran

- Example use for MPI\_Comm\_spawn:

```
numslaves = 10
call MPI_INFO_CREATE( spawninfo, ierr )
call MPI_INFO_SET( spawninfo, 'host', 'nome.mcs.anl.gov', ierr )
call MPI_INFO_SET( spawninfo, 'path', '/home/kosh/progs', ierr )
call MPI_INFO_SET( spawninfo, 'wdir', '/home/kosh/tmp', ierr )
call MPI_COMM_SPAWN( 'slave', MPI_ARGV_NULL, numslaves, &
                    spawninfo, 0, MPI_COMM_WORLD, &
                    slavecomm, MPI_ERRCODES_IGNORE, ierr )
call MPI_INFO_FREE( spawninfo, ierr )
```

116

## Using MPI\_Info in C

- Example use for MPI\_Comm\_spawn:

```
numslaves = 10;
MPI_Info_create( &spawninfo );
MPI_Info_set( spawninfo, "host", "nome.mcs.anl.gov" );
MPI_Info_set( spawninfo, "path", "/home/kosh/progs" );
MPI_Info_set( spawninfo, "wdir", "/home/kosh/tmp" );
MPI_Comm_spawn( "slave", MPI_ARGV_NULL, numslaves,
                spawninfo, 0, MPI_COMM_WORLD,
                &slavecomm, MPI_ERRCODES_IGNORE );
MPI_Info_free( &spawninfo );
```

117

## Other Process Management Features

- MPI\_Comm\_connect and MPI\_Comm\_accept allow two running MPI programs to connect and interoperate
  - Not intended for client/server applications
  - Useful in assembling complex distributed applications
- MPI\_Join allows the use of a TCP socket to connect two applications

118

## Dynamic Process Management

---

- Spawning new processes is *collective*, returning an intercommunicator.
  - Local group is group of spawning processes.
  - Remote group is group of new processes.
  - New processes have own `MPI_COMM_WORLD`.
  - `MPI_Comm_get_parent` lets new processes find parent communicator.

119


## Intercommunicators

---

- Contain a *local* group and a *remote* group
- Point-to-point communication is between a process in one group and a process in the other.
- Can be merged into a normal (intra) communicator.
- Created by `MPI_Intercomm_create` in MPI-1.
- Play a more important role in MPI-2, created in multiple ways.

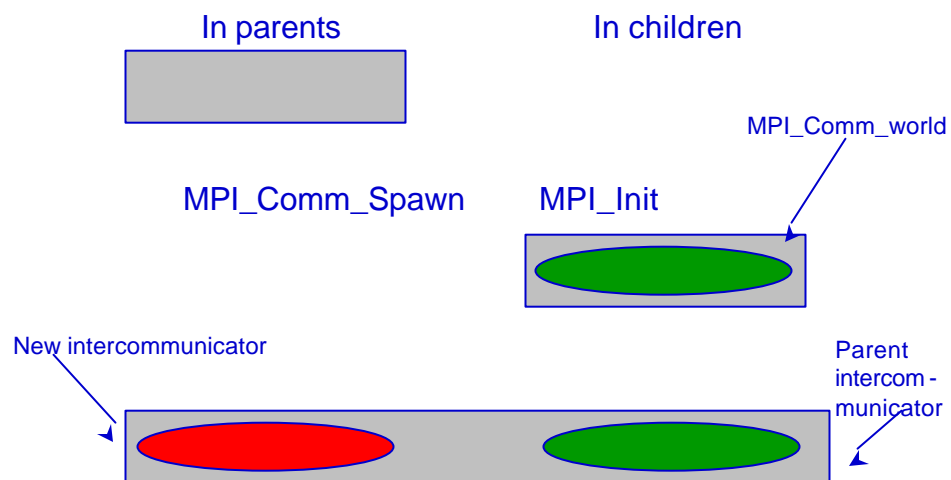
120

## Intercommunicators



121

## Spawning New Processes



122

## Spawning Processes

```
MPI_Comm_spawn(command, argv, numprocs, info,  
root, comm, intercomm, errcodes)
```

- Tries to start **numprocs** process running **command**, passing them command-line arguments **argv**.
- The operation is collective over **comm**.
- Spawnees are in remote group of **intercomm**.
- Errors are reported on a per-process basis in **errcodes**.
- **Info** used to optionally specify hostname, archname, wdir, path, file, softness.

123

## Spawning Multiple Executables

- **MPI\_Comm\_spawn\_multiple( ... )**
- Arguments **command**, **argv**, **numprocs**, **info** all become arrays.
- Still collective

124

## In the Children

- **MPI\_Init** (only MPI programs can be spawned)
- **MPI\_COMM\_WORLD** is processes spawned with one call to **MPI\_Comm\_spawn**.
- **MPI\_Comm\_get\_parent** obtains parent intercommunicator.
  - Same as intracommunicator returned by **MPI\_Comm\_spawn** in parents.
  - Remote group is spawners.
  - Local group is those spawned.


125

## Manager-Worker Example

- Single manager process decides how many workers to create and which executable they should run.
- Manager spawns  $n$  workers, and addresses them as 0, 1, 2, ...,  $n-1$  in new intercomm.
- Workers address each other as 0, 1, ...  $n-1$  in **MPI\_COMM\_WORLD**, address manager as 0 in parent intercomm.
- One can find out how many processes can usefully be spawned (**MPI\_UNIVERSE\_SIZE** attribute of **MPI\_COMM\_WORLD**).

126

## Parallel File Copy



127

## Parallelism in Distributing Files

- All processes writing in parallel
- Message-passing in parallel (assume scalable implementation of MPI\_Bcast)
- Pipeline parallelism with blocks from the file
- Use syntax adopted from cp:

```
pcp 0-63 /home/progs/cpi /tmp/cpi
```

128

## Code for pcp Master -1

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#define BUFSIZE      256*1024
#define CMDSIZE      80
int main( int argc, char *argv[] )
{
    int      num_hosts, mystatus, allstatus, done, numread;
    int      infd, outfd;
    char      outfilename[MAXPATHLEN], controlmsg[CMDSIZE];
    char      buf[BUFSIZE];
    char      soft_limit[20];
    MPI_Info hostinfo;
    MPI_Comm pcpslaves, all_procs;

    MPI_Init( &argc, &argv );
```

129

## pcp Master - 2

```
makehostlist( argv[1], "targets", &num_hosts );
MPI_Info_create( &hostinfo );
MPI_Info_set( hostinfo, "file", "targets" );
sprintf( soft_limit, "0:%d", num_hosts );
MPI_Info_set( hostinfo, "soft", soft_limit );
MPI_Comm_spawn( "pcp_slave", MPI_ARGV_NULL, num_hosts,
                hostinfo, 0, MPI_COMM_SELF, &pcpslaves,
                MPI_ERRCODES_IGNORE );
MPI_Info_free( &hostinfo );
MPI_Intercomm_merge( pcpslaves, 0, &all_procs );
```

130



## pcp master - 3

```
strcpy( outfilename, argv[3] );
if ( (infd = open( argv[2], O_RDONLY ) ) == -1 ) {
    fprintf( stderr, "%s doesn't exist\n", argv[2] );
    sprintf( controlmsg, "exit" );
    MPI_Bcast( controlmsg, CMDSIZE, MPI_CHAR, 0,
              all_procs );
    MPI_Finalize();
    return -1 ;
}
else {
    sprintf( controlmsg, "ready" );
    MPI_Bcast( controlmsg, CMDSIZE, MPI_CHAR, 0,
              all_procs );
}
```

131

## pcp Master - 4

```
MPI_Bcast( outfilename, MAXPATHLEN, MPI_CHAR, 0,
          all_procs );
if ( (outfd = open( outfilename,
                  O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU ) ) == -1 )
    mystatus = -1;
else
    mystatus = 0;
MPI_Allreduce( &mystatus, &allstatus, 1, MPI_INT,
              MPI_MIN, all_procs );
if ( allstatus == -1 ) {
    fprintf( stderr, "Output file %s not opened\n",
            outfilename );
    MPI_Finalize();
    return 1 ;
}
```

132

## pcp Master - 5

```
done = 0;
while (!done) {
    numread = read( infd, buf, BUFSIZE );
    MPI_Bcast( &numread, 1, MPI_INT, 0, all_procs );
    if ( numread > 0 ) {
        MPI_Bcast( buf, numread, MPI_BYTE, 0, all_procs);
        write( outfd, buf, numread );
    }
    else {
        close( outfd );
        done = 1;
    }
}
MPI_Comm_free( &pcpslaves );
MPI_Comm_free( &all_processes );
MPI_Finalize();
return 0;
```

133

## pcp Slave - 1

```
#include "mpi.h"
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#define BUFSIZE    256*1024
#define CMDSIZE    80
int main( int argc, char *argv[] )
{
    int      mystatus, allstatus, done, numread;
    char     outfilename[MAXPATHLEN], controlmsg[CMDSIZE];
    int      outfd;
    char     buf[BUFSIZE];
    MPI_Comm slavecomm, all_processes;

    MPI_Init( &argc, &argv );

    MPI_Comm_get_parent( &slavecomm );
```

134

## psp Slave - 2

```
MPI_Intercomm_merge( slavecomm, 1, &all_procs );
MPI_Bcast( controlmsg, CMDSIZE, MPI_CHAR, 0,
           all_procs );
if ( strcmp( controlmsg, "exit" ) == 0 ) {
    MPI_Finalize();
    return 1;
}
MPI_Bcast( outfilename, MAXPATHLEN, MPI_CHAR, 0,
           all_procs );
if ( (outfd = open( outfilename, O_CREAT|O_WRONLY|O_TRUNC,
                   S_IRWXU ) ) == -1 )
    mystatus = -1;
else
    mystatus = 0;
MPI_Allreduce( &mystatus, &allstatus, 1, MPI_INT, MPI_MIN,
               all_procs );
if ( allstatus == -1 ) {
    MPI_Finalize();
    return -1;
}
}
```

135

## pcp Slave - 3

```
/* at this point all files have been successfully opened */
done = 0;
while ( !done ) {
    MPI_Bcast( &numread, 1, MPI_INT, 0, all_processes );
    if ( numread > 0 ) {
        MPI_Bcast( buf, numread, MPI_BYTE, 0, all_procs );
        write( outfd, buf, numread );
    }
    else {
        close( outfd );
        done = 1;
    }
}
MPI_Comm_free( &slavecomm );
MPI_Comm_free( &all_processes );
MPI_Finalize();
return 0;
}
```


136

## Establishing Connections

- Two sets of MPI processes may wish to establish connections, e.g.,
  - Two parts of an application started separately.
  - A visualization tool wishes to attach to an application.
  - A server wishes to accept connections from multiple clients. Both server and client may be parallel programs.
- Establishing connections is collective but asymmetric (“Client”/“Server”).
- Connection results in an intercommunicator.

137

## Establishing Connections Between Parallel Programs



138

## Connecting Processes

- Server:
  - `MPI_Open_port( info, port_name )`
    - system supplies `port_name` (output argument)
    - might be `host:num`; might be low-level switch #, etc.
  - `MPI_Comm_accept( port_name, info, root, comm, &intercomm )`
    - collective over `comm`
    - returns `intercomm`; remote group is clients
- Client:
  - `MPI_Comm_connect( port_name, info, root, comm, &intercomm )`
    - remote group is server

139

## Optional Name Service

```
MPI_Publish_name( service_name, info,
                  port_name )
```

```
MPI_Lookup_name( service_name, info,
                 port_name )
```

- allow connection between `service_name` known to users and system-supplied `port_name`
- MPI implementations are allowed to ignore this service.
  - But most implementations could use openLDAP

140

## Bootstrapping

**MPI\_Join( fd, intercomm )**

- collective over two processes connected by a socket.
- The **fd** is a file descriptor for an open, quiescent socket.
- The **intercomm** is a new intercommunicator.
- Can be used to build up full MPI communication.
- The fd **fd** is *not* used for MPI communication.

141

## MPI and Threads

- MPI describes parallelism between *processes*
- *Thread* parallelism provides a shared-memory model within a process
- OpenMP and pthreads are common
  - OpenMP provides convenient features for loop-level parallelism

142

## Threads and MPI in MPI-2

---

- MPI-2 specifies four levels of thread safety
  - MPI\_THREAD\_SINGLE : only one thread
  - MPI\_THREAD\_FUNNELED : only one thread that makes MPI calls
  - MPI\_THREAD\_SERIALIZED : only one thread at a time makes MPI calls
  - MPI\_THREAD\_MULTIPLE : any thread can make MPI calls at any time
- MPI\_Init\_thread( ..., required, &provided) can be used instead of MPI\_Init

143

## MPI and OpenMP

---

- Loop-level parallelism
  - Requires only MPI\_THREAD\_FUNNELLED
- Task-parallelism
  - Requires MPI\_THREAD\_MULTIPLE

144

## Using OpenMP with MPI

- The famous cpi program, using OpenMP for loop-level
- Worker/manager

145

## cpi with Bcast/Reduce and OpenMP

```
<get n to use>
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
h  = 1.0 / (double) n;
sum = 0.0;
#pragma omp parallel for private(x) reduction(+:sum)
for (i = myid + 1; i <= n; i += numprocs) {
    x = h * ((double)i - 0.5);
    sum += f(x);
}
mypi = h * sum;
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
           MPI_COMM_WORLD);
if (myid == 0) {
    printf("pi is %.16f, Error is %.16f\n",
           pi, fabs(pi - PI25DT));
}
```

146



## Task Parallelism with MPI and OpenMP

```
#pragma omp sections
{
# pragma omp section
{
while (not_done) {
MPI_Recv( ... );
if (status.MPI_TAG == ALL_DONE) break;
<handle message>
}
}
# pragma omp section
{
<compute thread>
}
}
```

147

## MPI Thread Levels and OpenMP

- **MPI\_THREAD\_SINGLE**
  - OpenMP not permitted
- **MPI\_THREAD\_FUNNELED**
  - All MPI calls in the “main” thread
  - MPI calls outside of any OpenMP sections
  - Most if not all MPI implementations support this level
- **MPI\_THREAD\_SERIALIZED**
  - #pragma omp parallel
  - ...  
#pragma omp single
  - {  
MPI calls allowed here as well
  - }
- **MPI\_THREAD\_MULTIPLE**
  - Any MPI call anywhere
  - But be careful of races
  - Some MPI implementations support this level


148

## Conclusions

- MPI is a proven, effective, portable parallel programming model
- MPI has succeeded because
  - features are orthogonal (complexity is the product of the number of *features*, not routines)
  - programmer can control memory motion (critical in high-performance computing)
  - complex programs are no harder than easy ones
  - open process for defining MPI led to a solid design
- MPI I/O is widely available and efficient
- Dynamic process management becoming available

149

## Tutorial Material on MPI, MPI-2



<http://www.mcs.anl.gov/mpi/{usingmpi,usingmpi2}>

150