

Hierarchical cluster for scalable web servers

Presented at Cluster 2001, Newport Beach, CA., October 9, 2001

Dongseung Kim & Jonghuck Hong

Korea University

1

Parallel Computing & Architecture Lab. Korea Univ.

Motivation

- I Web servers & scalability
 - cluster: easy expansion
 - content based server switching
 - easy job migration
- I High efficiency and throughput of web servers
 - High throughput requirement for peak load time
 - Job dispatching bottleneck in the dispatcher when the arrival is burst


2

Parallel Computing & Architecture Lab. Korea Univ.

Expansion & performance scalability

I Expansion by front-end nodes

- multiple front-end node for dispatching
- load distribution among front-end nodes
 - 1. Bypass



3


- 2. Redirection

Parallel Computing & Architecture Lab. Korea Univ.

HTTP redirection-Front-end nodes

- I One virtual IP – transparent to outside users
- I HTTP redirection - expandability

- First contact thru virtual IP port no. 80
- second and later connection: port 8080
- connection consistency - designated servers
- Two connections needed



4

Parallel Computing & Architecture Lab. Korea Univ.

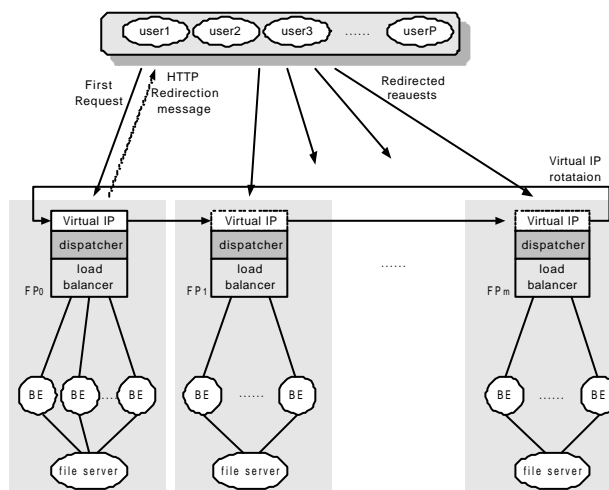
Load balancing/distribution

- I Imbalance among front-end nodes
 - weighted round-robin scheduling using HTTP redirection
- I Among backend nodes
 - Weighed round robin scheduling
 - load index: average file size transmitted
 - Normalized load index used for weight computation

5

Parallel Computing & Architecture Lab. Korea Univ.

Hierarchical Architecture



6

Parallel Computing & Architecture Lab. Korea Univ.

Experiments

I Hardware & software

- Cluster of 16 Linux PCs
- Pentium II CPU with 128 MB memory
- 100Mbps Ethernet switch
- Partition to both servers and clients
- Documents with various file sizes

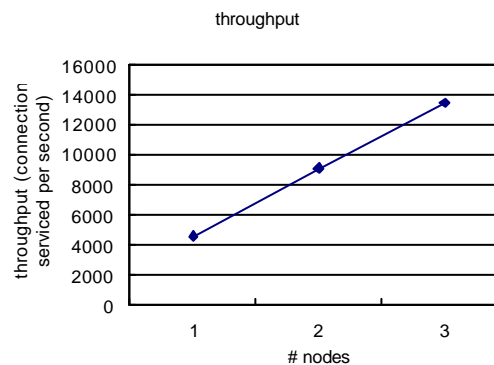
7

Parallel Computing & Architecture Lab. Korea Univ.

Scalability of front end nodes

I Configuration

- 1-3 front end nodes
- 6 backend
- 6 client no


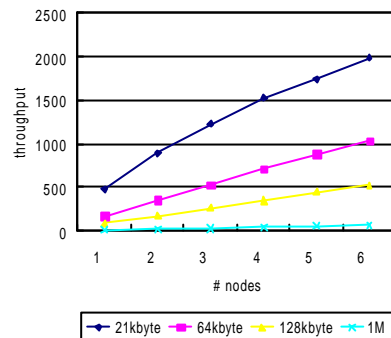


8

Parallel Computing & Architecture Lab. Korea Univ.

Scalability in backend nodes

Apache vs. khttpd



9

Parallel Comput

Load balancing in backend nodes


- I configuration
 - Four client nodes
 - One front end dispatcher
 - Four web server nodes
- I Load level
 - Full-loaded by client requests
- I Document sizes
 - 16k, 21k, 64k, 128k, 256k
- I Forced imbalance if Round Robin scheduling is used
 - 2:1:1:1
 - 8:1:1:1

10


Parallel Computing & Architecture Lab. Korea Univ.

Load distribution - 2:1:1:1

- RR scheduling




- Dynamic scheduling




Parallel Computing & Architecture Lab. Korea Univ.

Load distribution - 2:1:1:1

- RR



- Dynamic load balancing




12


Parallel Computing & Architecture Lab. Korea Univ.

Snapshot of load variation - 2:1:1:1 (load measured in every 0.3 second)

• RR



• Dynamic




13




Parallel Computing & Architecture Lab. Korea Univ.

8:1:1:1

• RR



• Dynamic scheduling



ws8 ws9 ws10 ws11


ws8 ws9 ws10 ws11

14


Parallel Computing & Architecture Lab. Korea Univ.

I Throughput with static load of 8:1:1:1

- RR



- Dynamic scheduling



15

Parallel Computing & Architecture Lab. Korea Univ.

Conclusions

I Multiple front-end nodes for concurrent job dispatching

- hierarchical architecture for content-based switching
- Bottleneck in dispatcher removed by using multiple front end nodes with HTTP redirection

I Redirection in back-end server nodes

- Weighted Round-robin scheduling
- Load index- number of average sized file served

16

Parallel Computing & Architecture Lab. Korea Univ.