



Dynamic Allocation of Nodes on a Large Space-shared Cluster

Lisa Kennicott Lee Ann Fisk

**Sandia National Laboratories
Albuquerque, New Mexico**

Sandia is a multiprogram laboratory operated by Sandia Corporation, a Lockheed Martin Company,
for the United States Department of Energy under contract DE-AC04-94AL85000.





Outline

- **Cplant™ runtime system overview**
- **Dynamic allocation motivation**
- **Simulation model**
- **Results**
- **Implementation decisions**
- **MPI-2 implementation**




Runtime Environment Components

- **Yod**
 - Parallel job launcher
 - Fans out executable and environment to compute nodes
- **Bebopd**
 - Compute node allocator
 - Allocates compute nodes to parallel jobs
- **PCT (Process Control Thread)**
 - Compute node daemon
 - Manages the compute and memory resources on a node
- **Pingd/Showmesh**
 - Compute node status tool
 - Displays the state of the compute nodes



Application Launch Performance





PBS on Cplant™

- **Added a “size” abstract resource to PBS**
- **PBS knows about total size and how much size is requested by jobs in the queue**
- **Bebopd notifies the PBS server when the number of compute nodes available to PBS changes**
- **Bebopd knows if a yod request is originating from a PBS job script and insures that a job does not exceed its size**
- **PBS insures that jobs do not exceed their allocated wall time**
- **PBS MOM daemons only run on each service node, not each compute node**



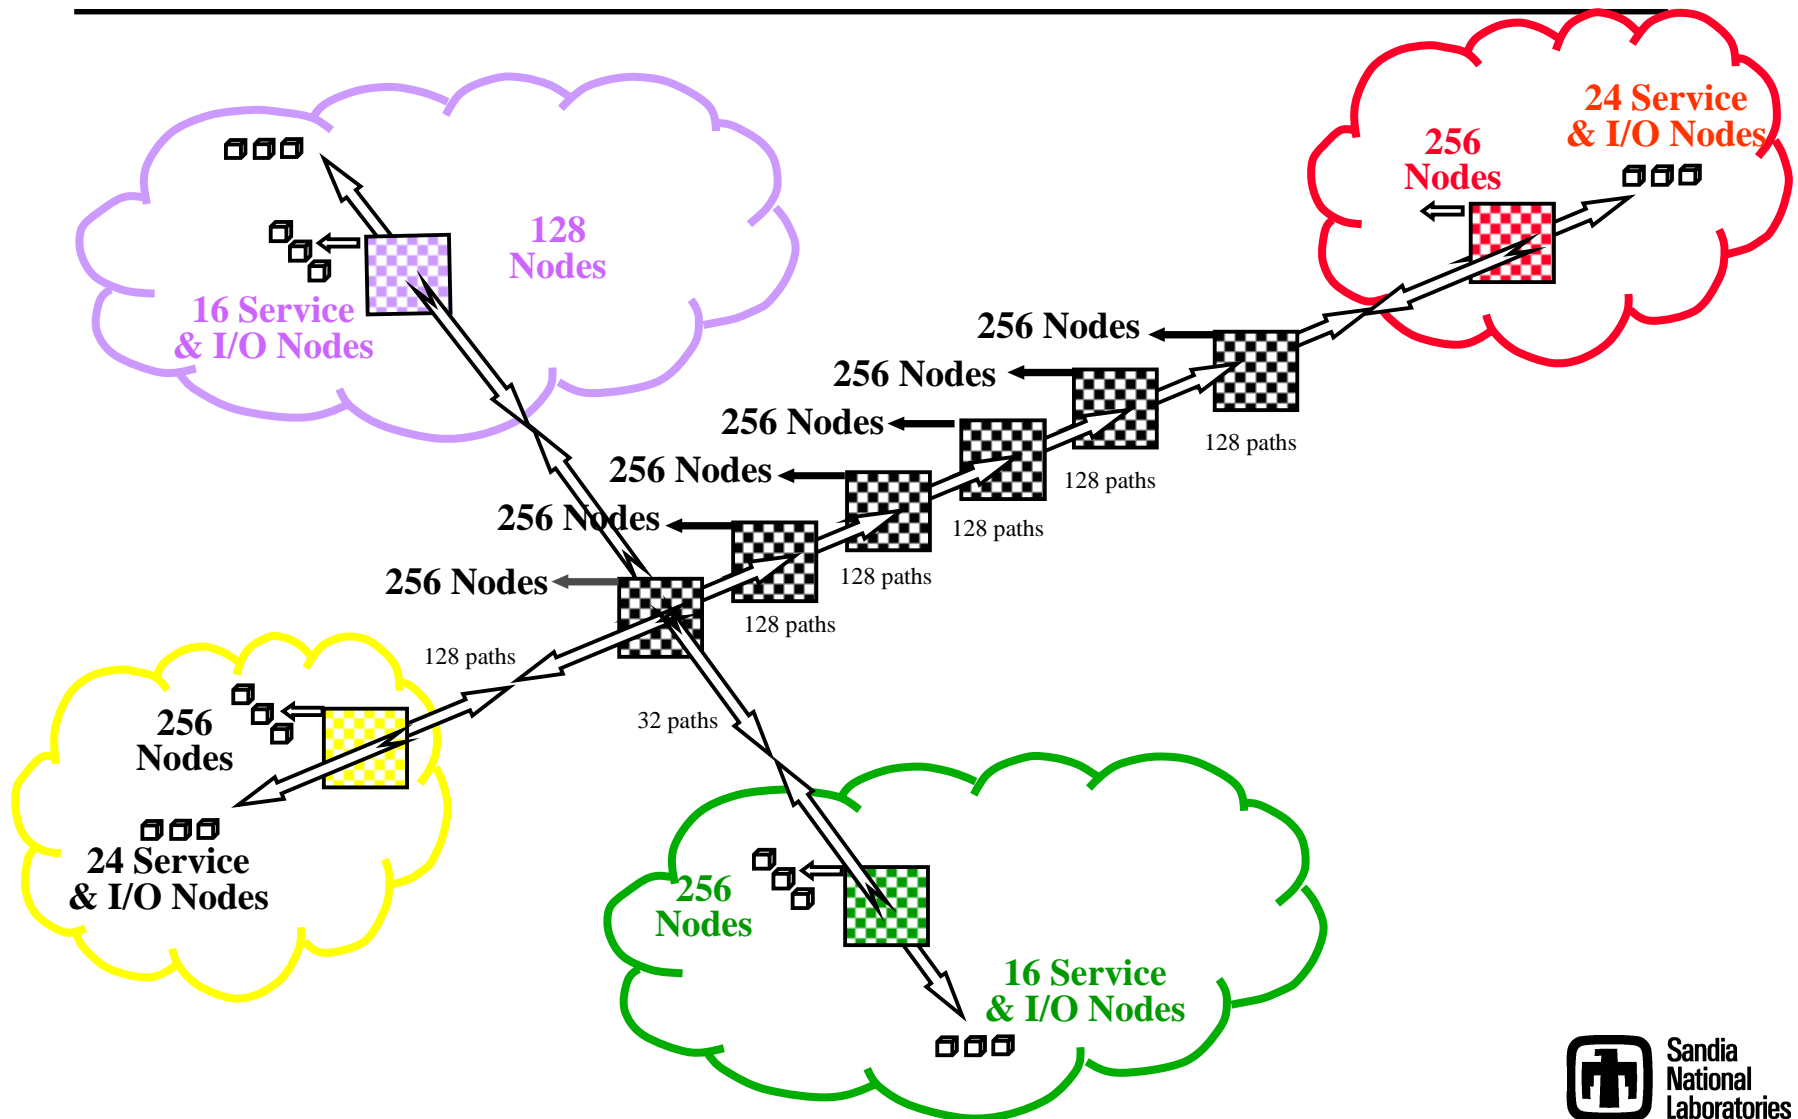
Cplant™ Alaska Cluster

- 400 Digital PWS 500a (Miata)
- 500 MHz Alpha 21164 CPU
- 2 MB L3 Cache
- 192 MB ECC SDRAM
- 16-port Myrinet SAN/LAN switch
- 32-bit, 33 MHz LANai-4 NIC
- 6 DEC AS1200, 12 RAID (.75 Tbyte) || file server
- 1 DEC AS4100 compile & user file server
- Integrated by Compaq





Cplant™ Antarctica Clusters





Dynamic Allocation Overview



- **Size (# nodes) of a running job can increase or decrease mid-run, new nodes are dynamically allocated to the job.**
- **System software detects conditions that are appropriate for DA and tells job(s) to reconfigure.**
- **Job reconfigures to more/less nodes at next convenient point (user-defined)**
- **Proposed to improve scheduling, system utilization, fragmentation, improved running time of some jobs**
- **Can be a very responsive scheduling policy to sudden changes in system load**



But.... It's a mixed bag

- Some studies find medium/large jobs do well - others find small jobs do better (until job arrival rates go up)
- Reconfiguration costs must be kept low - then it's great. But how easy is this?
- Low job arrival rate (more time to determine best reallocation) vs. high job arrival rate (less time available)
- High system utilization rates and high reconfiguration costs can cause thrashing
- Costs/benefits difficult to quantify and depend on implementation scheme
- So... Some groups are considering other scheduling policies such as backfilling and gang-scheduling



Different View of Dynamic Allocation

- Application decides when it needs more nodes vs. system detects conditions and orders job to change
- Application requests additional nodes. System fulfills it at earliest convenient time
- Goal: minimize interference of this new scheduling policy on OTHER jobs in the system (non-DA jobs)
- Metric is job queue time. In a space-sharing machine runtime of non-DA jobs will be minimally impacted
- We don't see dynamic allocation as a performance enhancement tool, but as a necessary and desirable scheduling feature for certain classes of jobs



Why are we interested in Dynamic Allocation?

User surveys and discussions revealed:

- Some jobs **NEED** it (calculations cannot continue unless additional nodes are available at some point mid-run). Otherwise you have to request many nodes and leave them sitting idle for potentially long times.
- Some jobs **WANT** it (load balancing is complex code; running time may improve)
- **NEED**: Finite element codes that employ adaptive mesh refinement
- **WANT**: Simulation codes (e.g. shock physics codes)



Why simulate?

- **Cost of developing it and being wrong too high**
- **Test out several different possible implementations in a short period of time**
- **Provide users a reasonable picture of costs/benefits to dynamic allocation**
- **But, it's important to validate the model to make sure we can fully trust our results**



Simulation Models

- **DA jobs can make a request for additional nodes at any time**
 - **“Need” request model**
 - Additional nodes requested all at once
 - Single request for 50% more nodes 60% of the way through original job run
 - **“Want” request model**
 - Smaller number of nodes is requested several times throughout the run
 - Requests 20 % of original node count at 25%, 50%, and 75% of the way through the original job run



Simulation Models (cont'd)

- DA “children” jobs must enter the job queue and are subject to various combinations of scheduling policies:
 - Priority - system gives scheduling priority to DA children jobs over other non-DA jobs in the queue
 - Non-priority - system gives equal priority to DA children and other non-DA jobs in the queue
 - Blocking vs. non-blocking - can the job do useful work until the additional nodes are allocated?
 - What are the costs of blocking? Should blocking be prohibited?

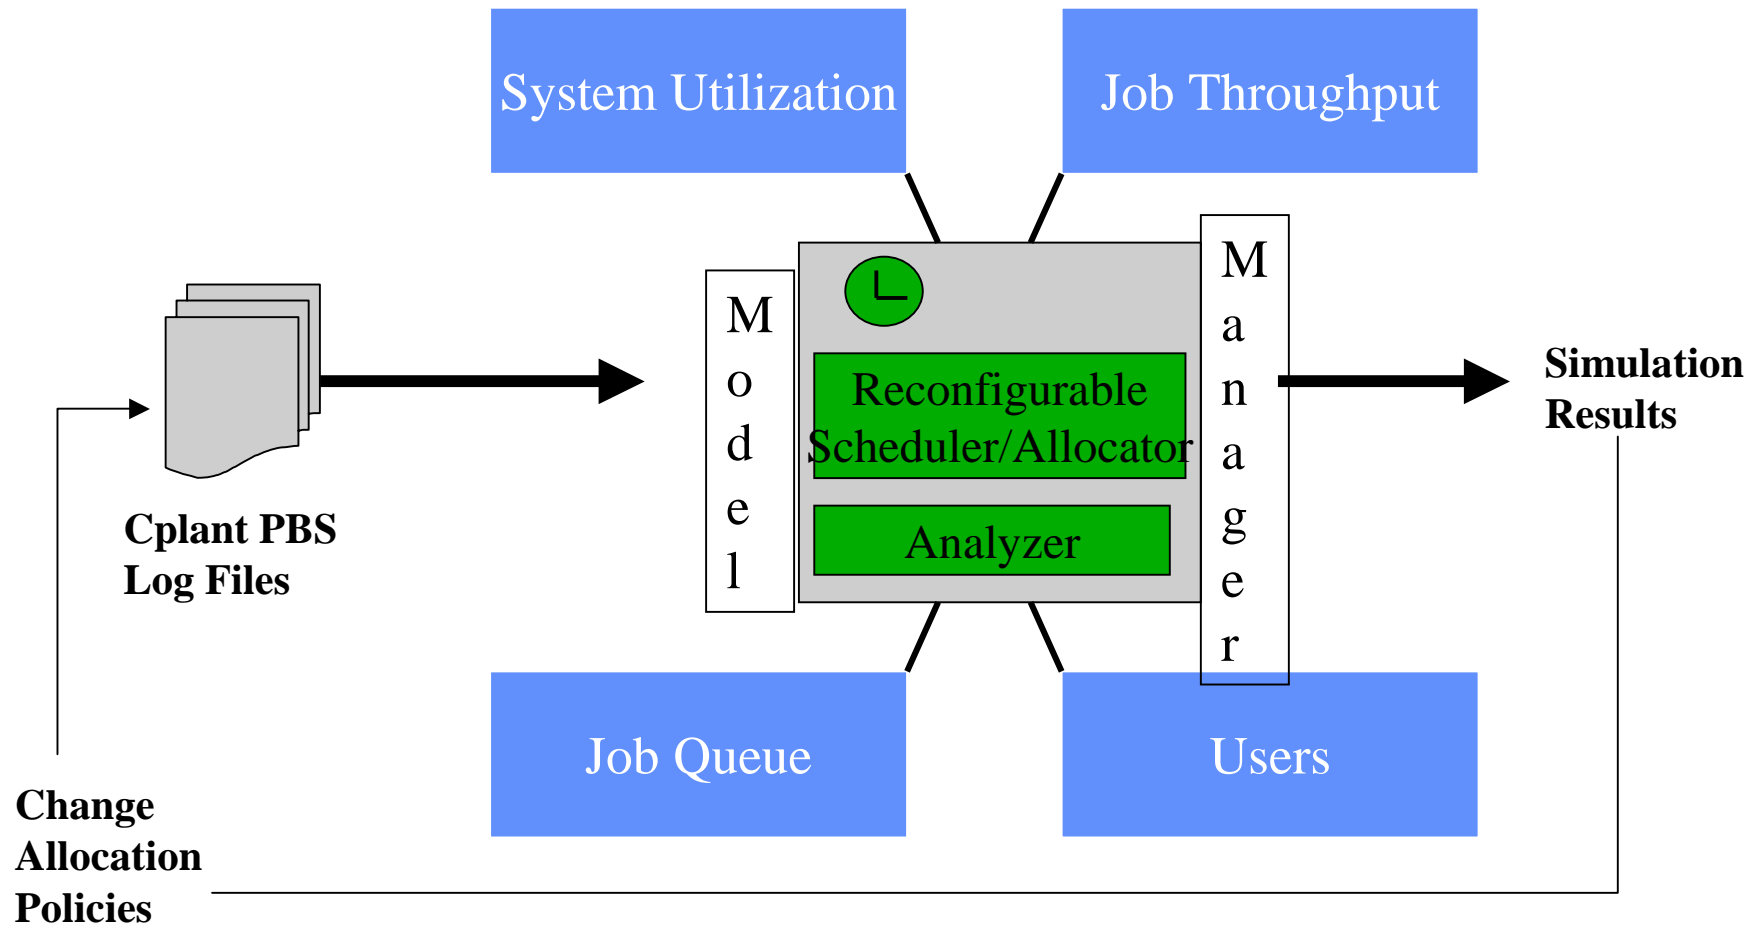


Simulator Description

- OO model of jobs, queues, processors, scheduling policies
- Scheduling policies can be quickly changed to play “what if” games
- Closely models Cplant™ scheduling (fair share, large vs. small job priorities, starvation policies, scheduled system time, etc.)
- Relied on past history of system usage rather than probabilistic predictions of job size, arrival rates, etc. We believe this gives a more accurate view of what’s really going on in the system
- Verified by taking historical job logs and matching results of simulation with history
 - Closely matched historical runs on Cplant with the exception of some human operator interference (minimal)
- All jobs - even those that immediately failed - were modeled. Even jobs that fail quickly affect job scheduling especially if they’re large



Simulation Model





Dynamic Allocation Results

Average wait times. Jobs requesting dynamic allocation are < 64 nodes in size.

SINGLE REQUEST	No Priority		Priority	
	<i>Blocking</i>	<i>Non-blocking</i>	<i>Blocking</i>	<i>Non-blocking</i>
Very small ($p < 16$)	2621	2615	2607	2604
Small ($16 \leq p < 64$)	3161	3161	3164	3164
Medium ($64 \leq p < 128$)	8999	8970	8945	8937
Large ($128 \leq p$)	103041	103041	103041	103041

TRICKLE REQUESTS	No Priority		Priority	
	<i>Blocking</i>	<i>Non-blocking</i>	<i>Blocking</i>	<i>Non-blocking</i>
Very small ($p < 16$)	2602	2586	2571	2572
Small ($16 \leq p < 64$)	3161	3161	3171	3163
Medium ($64 \leq p < 128$)	9098	8971	8954	8936
Large ($128 \leq p$)	103041	103041	103041	103041



Dynamic Allocation Results

Average wait times. Jobs requesting dynamic allocation are 64 -128 nodes in size.

One Request	No priority		Priority	
	Blocking	Non-blocking	Blocking	Non-blocking
Very Small ($p < 16$)	2413	2361	3937	3538
Small ($16 \leq p < 64$)	4186	3649	4171	4430
Medium ($64 \leq p < 128$)				
Large ($128 \leq p$)	103041	103041	134607	103041

Trickle requests	No priority		Priority	
	Blocking	Non-blocking	Blocking	Non-blocking
Very Small ($p < 16$)	5130	2360	5004	2955
Small ($16 \leq p < 64$)	7700	3646	21618	3846
Medium ($64 \leq p < 128$)				
Large ($128 \leq p$)	104350	103044	121908	103042



Dynamic Allocation Results

- With our current system usage and job distribution, it doesn't matter how we implement dynamic allocation.
- As the job size increases or the system utilization increases, this is no longer true.
- Cost of dynamic allocation on the rest of the system is heaviest for the “trickle” requests. Worse effect on small jobs than on large jobs.
- In general, it is better not to give scheduling priority to “dynamic children” for our implementation.
- Blocking until additional nodes are granted generally affects the rest of the system more than non-blocking.
- Small jobs do not receive additional nodes immediately 13-17% of the time. Large jobs 14-26% of the time.



Implementation Decisions

- Implementing DA in a no-priority, no-blocking (when possible) approach minimizes effects on other jobs in the queue
- Compromise: Go with no-priority but with blocking (ease of implementation) and give user the ability to call to layers below MPI for non-blocking
- Also found our scheduling policies could be improved
- Considering cycle-stealing jobs and backfilling



MPI-2 DA Support

- **Use MPI_Info object for**
 - **wdir** – working directory for new processes
 - **timeout** - timeout in seconds to wait for new processes to start
 - **soft** – extended specification of the number of new nodes that could be allocated (2:2:10 or 16,32)
- **Provide non-portable library routines for non-blocking spawning**



Conclusions

- Based on simulation data, implementing DA in a no-priority, non-blocking fashion is preferable
- Currently have an approach that allows for blocking, no-priority to be consistent with MPI-2
- Will also provide non-portable non-blocking spawning as well



Future Work

- Deploying backfilling policy to alleviate fragmentation as workload increases
- Experimenting with cycle-stealing jobs that can be killed when nodes are required for DA or queued jobs
- Evaluate these policies with regard to DA strategies
- Analyze log data for larger Antarctica clusters