

Approximation Algorithms for Data Distribution with Load Balancing of Web Servers

Li-Chuan Chen

lichen@mitre.org

The MITRE Corporation

Co-author: **Hyeong-Ah Choi**

(George Washington University)

Outline of the Talk

- Research Motivation and Goals
- Prior Literature
- Problem Formulation
- Research Results
- Conclusions and Future Directions

Research Motivation

With the increased popularity of World-Wide-Web (WWW or Web) there are a number of problems:

- Servers **overloaded**.
- Internet backbone **congestion**.
- **Slow** Web services access.

Background

Approaches to Reduce Server Load:

- **Mirror Web Sites**: Replicate web server contents throughout network. (User must select server.)
- **Web Caching**: Stores frequently requested Web documents closer to the users (Cache coherence).
- **Distributed Web Server**: Web documents are distributed among a cluster of servers acting as a single server (Load balancing).

Research Goals

To reduce Web server load and to increase efficiency and reliability of Web system performance via:

Load Balancing: Balancing the load among a set of distributed Web document servers.

Research Goals: (Continued)

We will consider the design of **optimization algorithms** for achieving these objectives.

- Note that most formulations of these problems are **NP-hard**.
- We consider **special cases** and **approximation algorithms** for Load Balancing.

Load Balancing for Web Servers

Assume:

- A **cluster** of back-end Web servers working together as a single Web server.
- A set of **Web documents** are to be allocated among these servers.
- A **mechanism** to redirect HyperText Transfer Protocol (HTTP) requests to one of the back-end servers.


Prior Literature

- **Client-based load balancing**: A list of replicated server's performance is maintained at the client's proxy server and then a URL is mapped onto one of the servers. [MDZ99, SBSV98, LM97]
- **Server-based load balancing**: Front-end server dispatches incoming Web requests to one of the back-end servers via round-robin Domain Name Service (DNS) or server load. [SNCC96, AYHI96 GGMP95]
- **Hybrid approaches**: Combination of DNS round-robin, HTTP redirection, and document's access rate to balance the load. [NRY97]

Problem Formulation

- Consider M servers and N documents. Server i is associated with memory size m_i and number of simultaneous HTTP connections l_i . Document j is associated with document size s_j and access rate r_j .
- Given an allocation of documents to servers, let R_i denote the total access rate for server i .
- Define load of server i to be R_i / l_i . The objective is to minimize the maximum load over all servers.
- **Input:** Quadruple $I = (r, l, s, m)$.
- **Output:** An allocation of documents to servers.

Example:



Research Results

- **Lower bound** on the optimal load is r/l .
- **Optimal allocation is NP-hard**: Reduction from bin-packing. This means optimal allocation problem probably cannot be solved in polynomial time.

Research Results (Continued)

We present approximation algorithms for various special cases.

- **No memory constraints:** $O(nm)$ -time factor-2 approximation of the optimal solution.
- **Memory and load constraints:** $O(n \log n)$ -time factor-4 approximation of the optimal solution.
- **Small documents:** If a server can hold at least k documents, a $2(1 + 1/k)$ -factor approximation.

Lower Bounds

Input: Quadruple $I = (r, l, s, m)$ with **no** memory constraint.

- Let r be the sum of access rates over all documents, and let l be the sum of connections over all servers.
- Let f^* be the optimal allocation cost. Then we have the following **lower bound** on f^* :


$$r_{\max} = \max_{1 \leq j \leq N} r_j, \quad l_{\max} = \max_{1 \leq i \leq M} l_i$$

$$f^* \geq \max \left\{ \frac{r_{\max}}{l_{\max}}, \frac{r}{l} \right\}$$

- Proof is based on averaging and the pigeon-hole principal.

NP-Completeness


- **0-1 Allocation** is NP-hard with **no memory** constraints by reduction from bin-packing, where r denotes size of the objects and bins are of size $l = l_1 = l_2 = \dots = l_M$.
- Similarly for **0-1 Allocation** with **memory** constraints, where objects are size of documents and bin sizes are size of memories.



Approximation Algorithm: No Memory

- Sort documents by decreasing order of access rate, r_j .
- For each document j
 - Place it in server i that minimizes $(R_i + r_j)/l_i$.
 - $R_i += r_j$.

Intuition: Put each document into the server with the greatest remaining load capacity.



Approximation Algorithm: No Memory (Continued)

- Approximation for no memory constraints is a $O(nm)$ -time, factor-2 approximation of optimal load.
- Why Factor-2? Each server will be utilized to at least half its capacity. Optimal cannot utilize to more than full capacity.


Approximation Algorithm: Memory and Load Constraints

- **Split** documents into 2 sets, D^1 , D^2 , where D^1 consists documents whose access cost is bigger than document size and D^2 consists documents whose document size is bigger than access cost.
- **Assign** as many documents which are in D^1 as possible and then assign the remaining documents which are in D^2 .
- If all documents have been assigned to some server then feasible allocation exist else no solution.

Approximation Algorithm: Memory and Load Constraints


Intuition:

- Access dominant and size dominant documents are allocated separately.
- Thus each server is not too poorly utilized according to either criteria.




Approximation Algorithm: Memory and Load Constraints

- For **equal memory** and **load** constraints, this is an $O(n \log n)$ -time, factor-4 approximation of optimal solution.
- **Why?** Intuitively each server is at least $\frac{1}{4}$ utilized in load or memory.



Approximation Algorithm: Small Documents

- Approximation for small documents is $2(1 + 1/k)$ time the optimal solution if each server can hold at least k documents.



Future Directions

Consider ways to strengthen our existing results either by improving efficiency of the algorithms or by eliminating some of the assumptions that are made.

- **Dynamic Load Balancing**: How to deal with server failures, access rate changes, and changes in server capacity.
- **On-line algorithm** for load balancing.