

# SOVIA: A User-level Sockets Layer Over Virtual Interface Architecture

Jin-Soo Kim

jinsoo@computer.org

Computer System Research Department Electronics and Telecommunications Research Institute (ETRI)

Korea

#### **Outline**

- **Introduction**
- **Example 2** Performance Issues
- Compatibility Issues
- Results on Real Applications
- **Concluding Remarks**



## Introduction

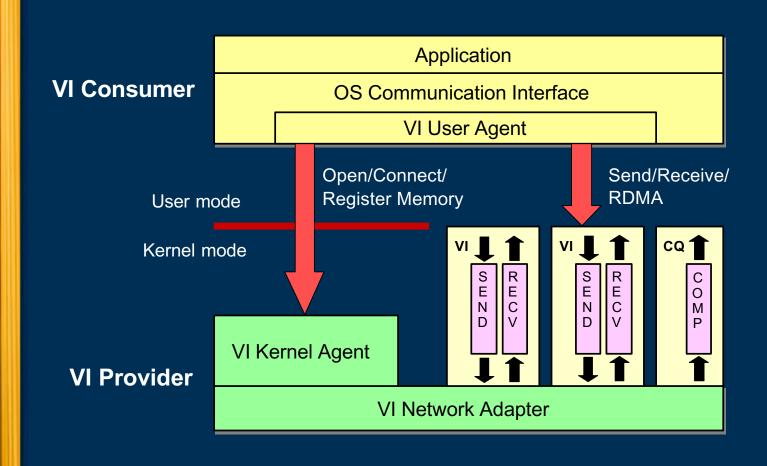


# Overheads in traditional communication architecture

- Protocol overhead in TCP/IP
- Data copying
- Context switching

#### Virtual Interface Architecture

- Led by Compaq, Intel, and Microsoft
- A standard software interface for userlevel access to a network hardware



#### VI Provider Library (VIPL)

#### **Hardware Connection**

VipOpenNic () VipCloseNic ()

#### Endpoint Creation and Destruction

VipCreateVi () VipDestoryVi ()

#### **Connection Management**

VipConnectWait ()
VipConnectAccept ()
VipConnectReject ()
VipConnectRequest ()
VipDisconnect ()

#### Memory Protection and Registration

VipCreatePtag ()
VipDestroyPtag ()
VipRegisterMem ()
VipDeregisterMem ()

#### Querying

VipQueryNic ()
VipQueryVi()
VipSetViAttributes ()
VipSetMemAttributes ()
VipQueryMem ()
VipQuerySystemManagementInfo ()

#### **Error**

VipErrorCallback ()

#### **Data Transfer**

VipPostSend ()
VipSendDone()
VipSendWait ()
VipSendNotify ()
VipPostRecv ()
VipRecvDone ()
VipRecvWait ()
VipRecvNotify ()

#### **Completion Queues**

VipCreateCQ ()
VipDestroyCQ ()
VipResizeCQ ()
VipCQDone ()
VipCQWait ()
VipCQNotify ()



#### **∠** Problem in VIPL

 "VIA is considered by many systems designers to be too low a level for application programming."

-- TFCC Cluster Computing White Paper

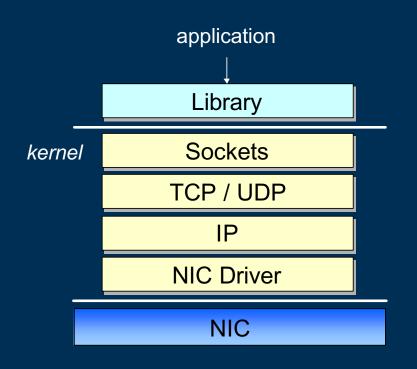
#### Requirements for a new layer

- High Performance
- Portability



## **Berkeley Sockets API**

#### **Traditional** architecture

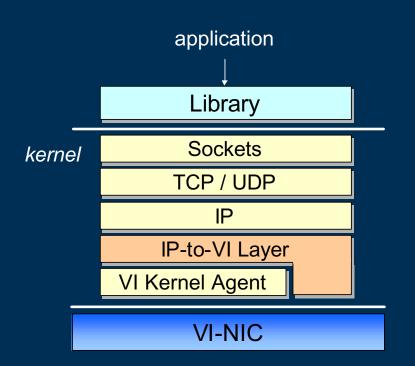


- Pros
  - ∠ Full compatibility
- Cons
  - ∠ TCP/IP overhead
  - ∠ Context switching
  - ∠ Data copy

#### Sockets over VIA (1)

#### ∠ Using an IP-to-VI Layer

• Giganet's LANE (LAN Emulation) Driver

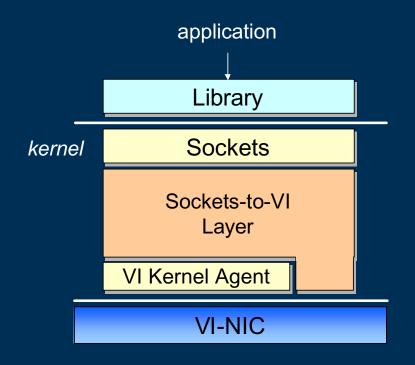


- Pros
  - ∠ Full compatibility
- Cons
  - ∠ TCP/IP overhead
  - ∠ Context switching
  - ∠ Data copy
  - Emulating Connectionless IP over connectionoriented VIA

#### Sockets over VIA (2)

#### ∠ Using a Sockets-to-VI Layer

VIsocket on Solaris

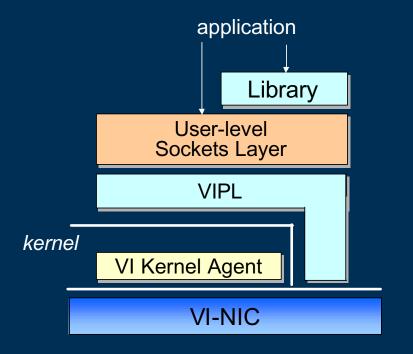


- Pros
  - TCP/IP overhead
- Cons
  - ∠ Context switching
  - ∠ Data copy

#### Sockets over VIA (3)

#### ∠ Using a User-level Sockets Layer

• MS WinSock Direct Path, SOVIA



- Pros
  - TCP/IP overhead
  - Context switching
  - Data copy
- Cons
  - ∠ Less compatibility



- Provide a simple, yet versatile communication service
- ∠ Accelerate the existing Sockets-based applications with a reasonable effort
  - Parallel and cluster file systems
  - User-level software DSMs, ...
- **∠** Target for another upper layers
  - RPC (Remote Procedure Calls)
  - MPI (Message Passing Interface), ...



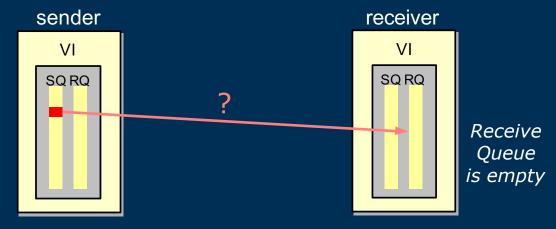
### Performance Issues

- Minimizing latency
- Maximizing bandwidth
- Evaluation

#### Synchronization Protocol (1)

#### VIA's pre-posting constraint

- The receiver should pre-post a descriptor before the sender initiates a data transfer.
- A high-level synchronization protocol needs to be implemented between the sender and receiver.



#### **Synchronization Protocol (2)**

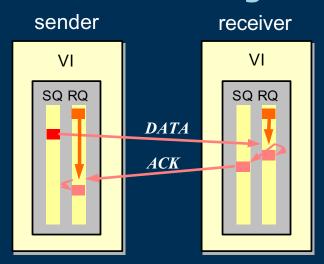
#### Satisfying pre-posting constraint

 Guarantee that at least one descriptor is available on the RQ for each send.

# Three-way Handshaking

# Sender receiver VI VI SQ RQ REQ ACK DATA

# Two-way Handshaking



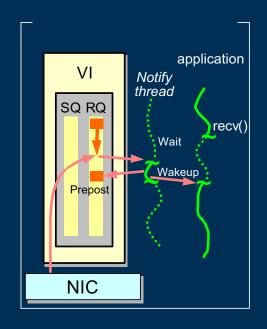


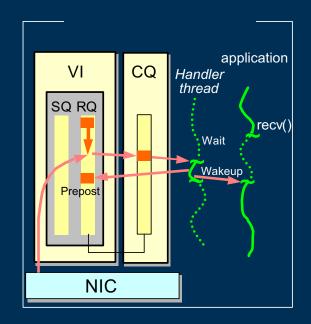
#### Message Handling Strategies

#### **∠** Incoming message handling

- The completed descriptors should be extracted from a queue manually.
- Incoming messages are delivered asynchronously.
- When and by whom?

# Multi-threaded Handling

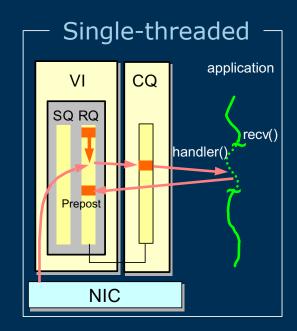




?

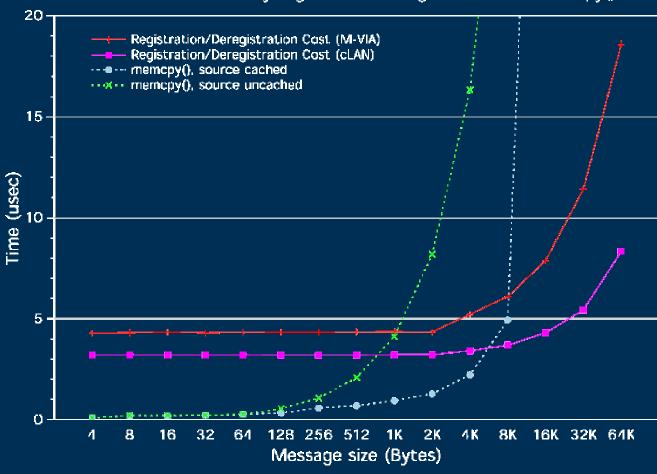
#### Single-threaded Handling

- The application thread handles incoming messages when it calls communicationinvolved functions.
- Can overlap the communication with the computation if multiple descriptors are pre-posted.



#### **Conditional Registration**

Unit cost for memory registration/deregistration and memopy()





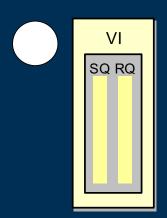
### Performance Issues

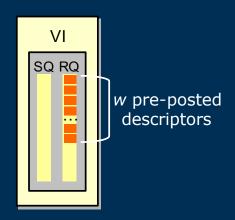
- Minimizing latency
- Maximizing bandwidth
- Evaluation

#### Flow Control

#### A sliding window protocol

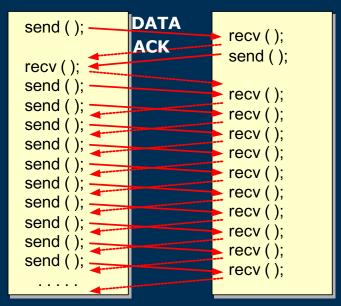
- Receiver pre-posts w descriptors
- send() decreases w by one
- Sender is blocked if w = 0
- Window size w is increased by ACK

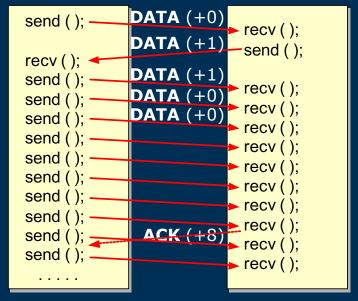




#### **Delayed ACKs and Piggybacking**

∠ Delay up to t ACKs
∠ Piggyback ACKs to DATA





**Before** 

After (t = 8)



- **∠** TCP uses Nagle algorithm
- ∠ Append outgoing data into a buffer (< 2KB), and start a timer</p>
- The buffer is flushed either
  - when the timer expires
  - when there is no enough room
  - when the size is > 2KB
  - when the application calls recv() or close()
- Dynamically turned off for latency-sensitive applications.



### Performance Issues

- Minimizing latency
- Maximizing bandwidth
- Evaluation



#### **Linux servers**

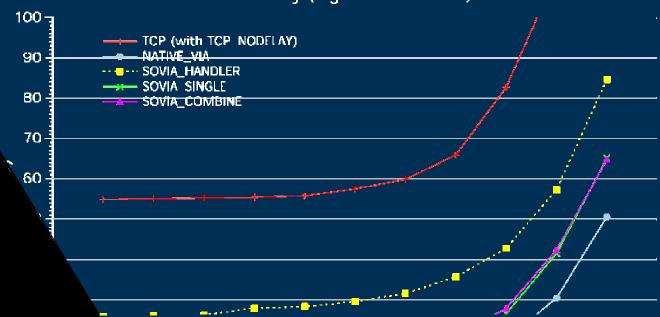
- Intel L440GX+ motherboard
- Intel Pentium III-500MHz
- 256MB main memory
- Linux kernel 2.2.16

#### VIA implementations

• cLAN v1.1.1 on Giganet cLAN1000

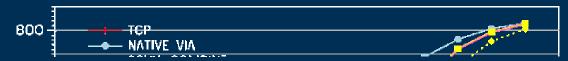
## Latency (cLAN)

Latency (Giganet cLAN1000)



## Bandwidth (cLAN)

Bandwidth (Giganet cLAN1000)





#### Minimizing latency

- Two-way handshaking
- Single-threaded implementation
- Conditional memory registration

#### Maximizing bandwidth

- A sliding window protocol
- Delayed acknowledgments and piggybacking
- Combining small messages

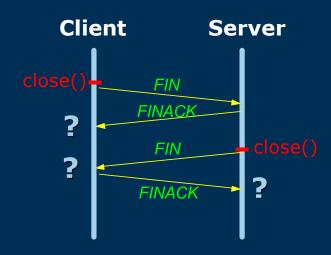


### **Compatibility Issues**

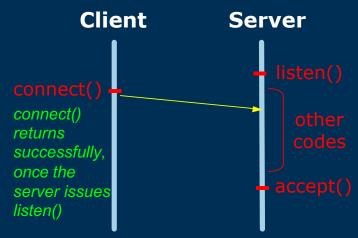
- Connection management
- Enhancing portability

#### **Connection Management**

# Problems in the single-threaded implementation

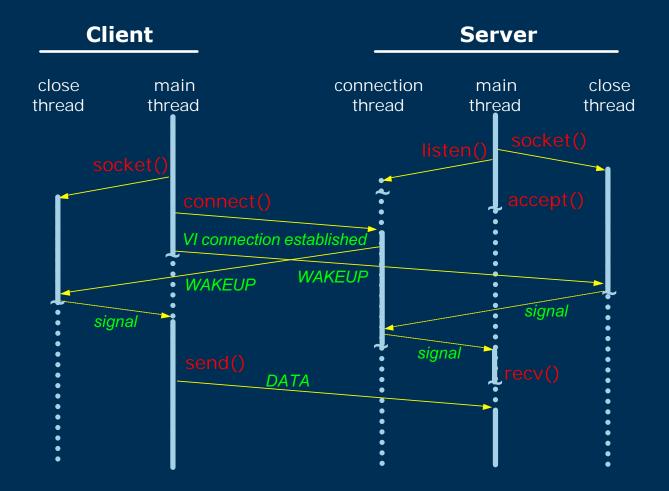


No chance to handle incoming messages after last close()

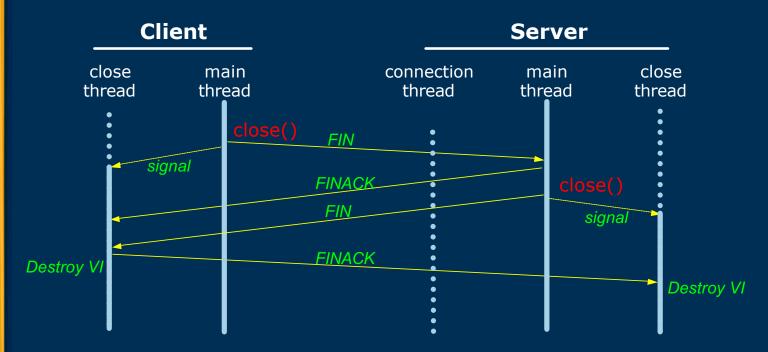


VipConnectWait() can't be called in listen(), as it would block the thread.

#### **Establishing a Connection**



#### Closing a Connection



- The close thread is not activated if there is an open connection.
- Hence, the presence of the close/connection thread dose not affect the application's performance.



### **Compatibility Issues**

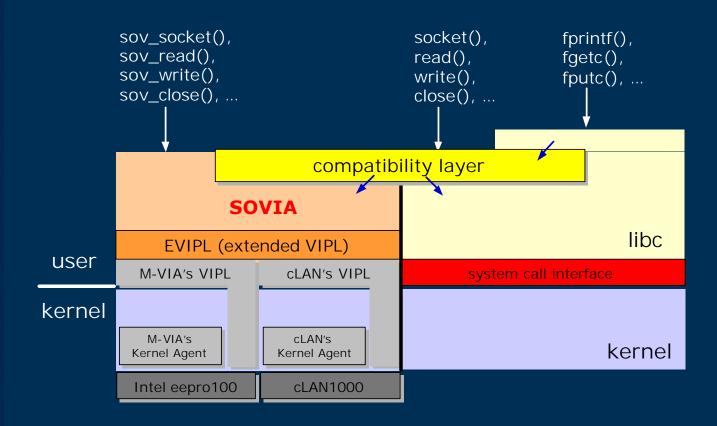
- Connection management
- Enhancing portability

#### **Issues in Porting Applications**

- **∠** Sockets-specific interface
  - socket(), connect(), accept(), ...
- **∠** File system interface
  - read(), write(), close(), ...
- Standard I/O library

```
int s;
FILE *fp;
...
s = socket (AF_INET, SOCK_STREAM, 0);
connect (s, (struct sockaddr *) &server, sizeof(server));
fp = fdopen (s, "w");
fprintf (fp, "Hello, world...\n");
```

#### **Enhancing Portability**





#### Other Compatibility Issues

#### Fork()

- OK if the child process is not involved in any communication – Had to fix the VIPL due to the copy-on-write problem.
- Sockets can not be shared with child processes.

#### Exec()

Exec() will destroy any user-level data.

#### Normal termination

Register a cleanup function using atfinalize().

#### Abnormal termination

• Catch the abnormal termination in signal handlers and call a cleanup function.



# Results on Real Applications



#### ∠ Linux NetKit 0.16

• linux-ftpd-0.16 & netkit-ftp-0.16

	File 1	File 2
File size (bytes)	19,090,223	145,864,380
TCP/IP on Fast Ethernet	90Mbps (1.63 sec)	90Mbps (12.7 sec)
TCP/IP on cLAN	262Mbps (0.59 sec)	254Mbps (4.61 sec)
SOVIA on cLAN	573Mbps (0.27 sec)	532Mbps (2.20 sec)
Local copy (in ramdisks)	611Mbps (0.25 sec)	538Mbps (2.17 sec)





#### 

- High performance from VIA
- Portability based on Sockets API

#### ∠ On-going & Future Work

- Accelerate other applications
  - ∠ Parallel file systems (e.g. PVFS)
  - Network Block Device (e.g. GNBD)
  - ∠ User-level software DSMs
  - Message-passing library, etc.
- Compare with kernel-level SOVIA