

Parallel Jobs Scheduling on Multi-cluster Computing Systems

J. H. Abawajy and S. P. Dandamudi


School of Computer Science

Carleton University, Ottawa, Ontario, Canada

{abawjem, sivarama}@scs.carleton.ca

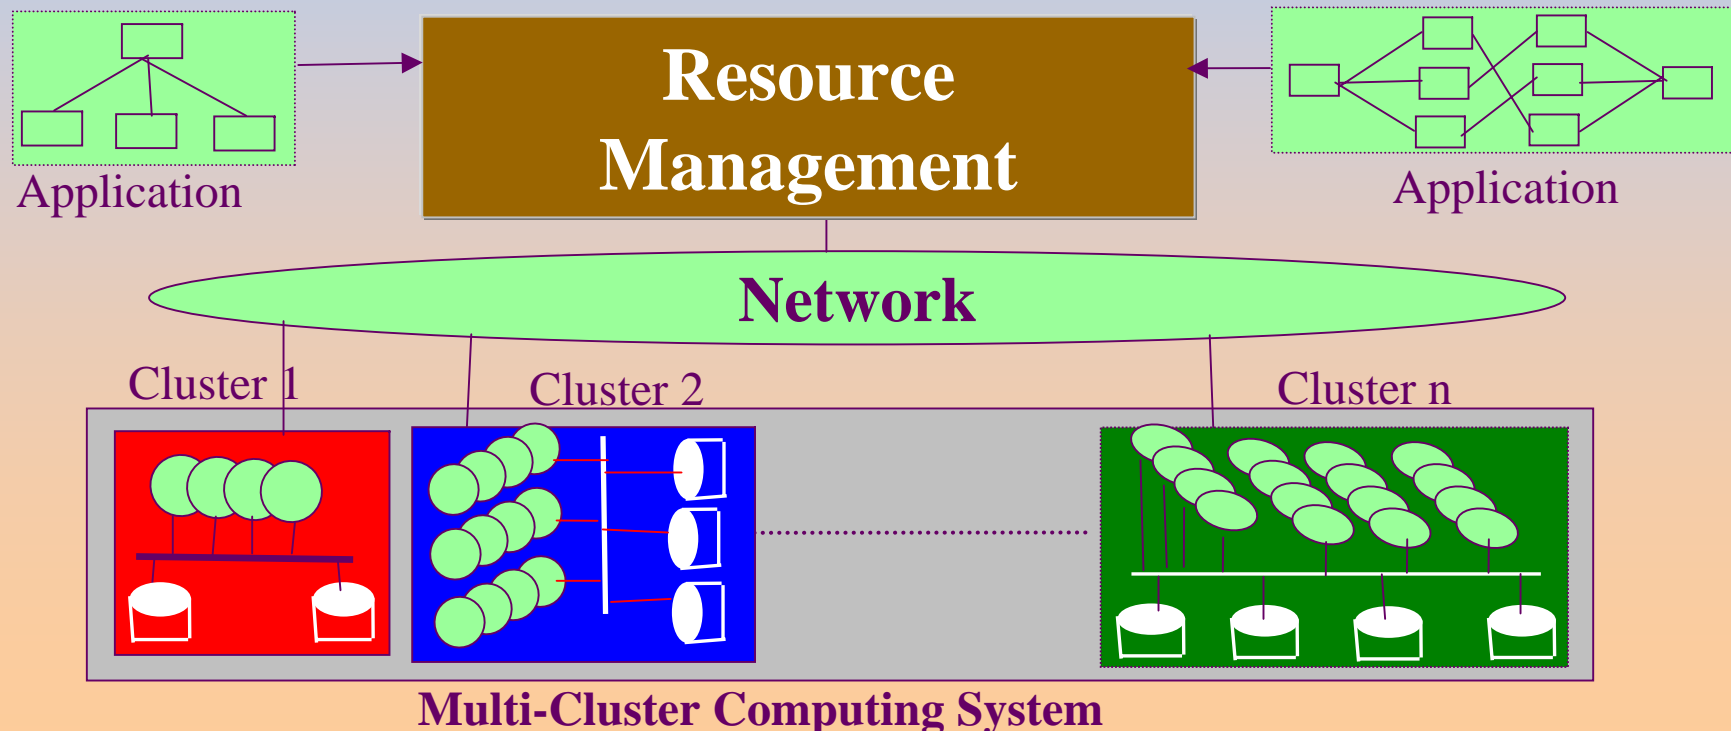
Background: A cluster

- A cluster is a collection of n independent workstations
 - ☞ Interconnected by LAN
 - ☞ Community-based (i.e., shareable)
 - ☞ May be homogeneous or heterogeneous.



Background: Multiple-clusters

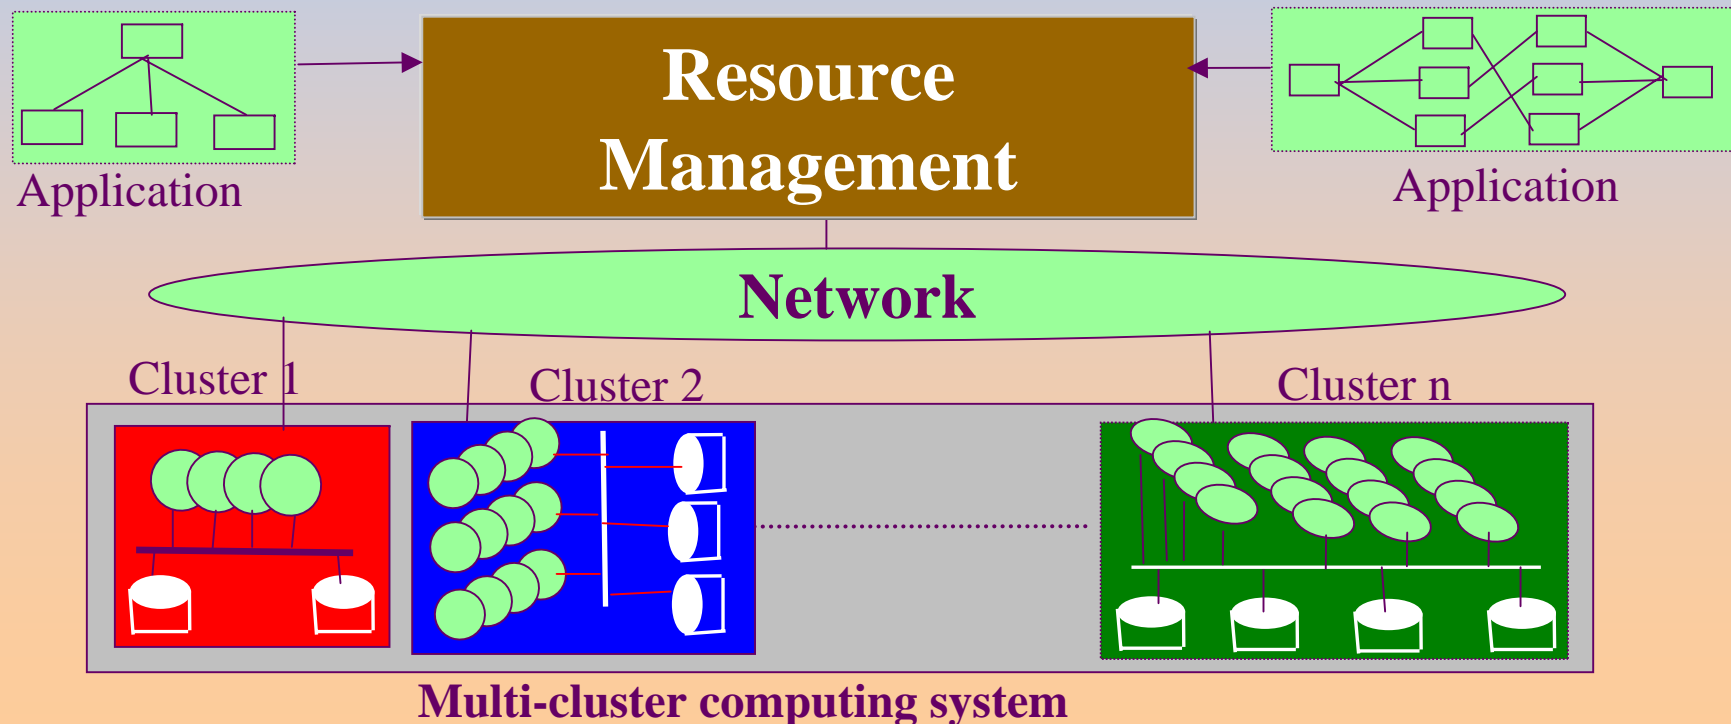
- Multi-cluster Computing System is formed by interconnecting multiple clusters via WAN.
- Each cluster can have different number of workstations



Background: Applications

■ Applications that can benefit from multi-cluster systems include:

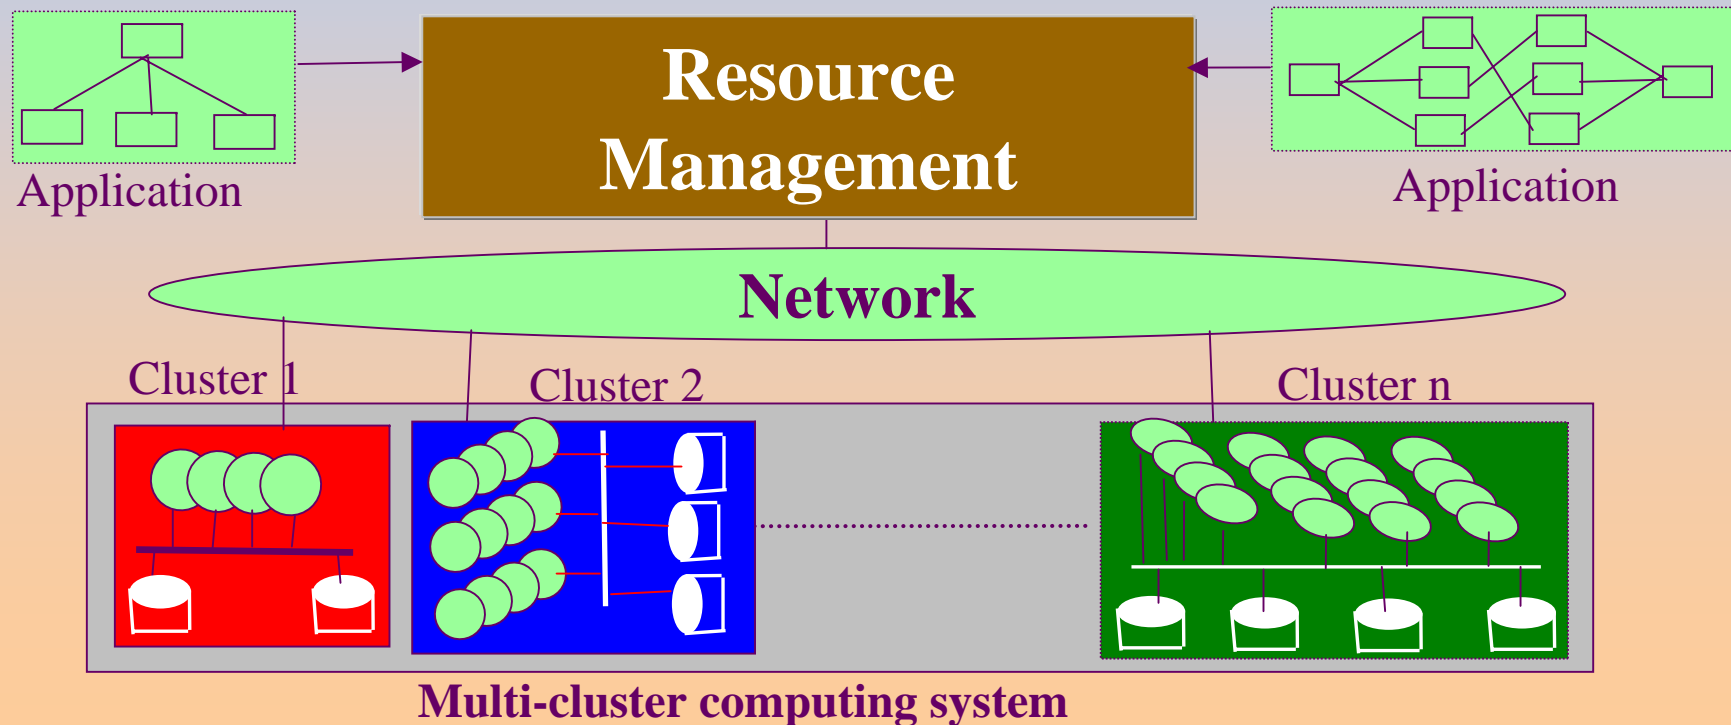
- ☞ parameter studies
- ☞ probabilistic analysis



Background: Resource Management

■ Proper resource management is a critical issue for multi-cluster computing systems:

- ☞ Competing users
- ☞ Diverse set of workloads coexist
- ☞ Distributed nature of the resources



Background: Resource Management

■ Current Focus is

- ➡ On single-cluster system
- ➡ Rarely address challenges posed by CC systems.
 - 📄 Scale of the system
 - 📄 Dynamically changing resource availability
 - 📄 Resource and workload heterogeneity.
 - 📄 Susceptibility to Failure

■ To address these challenges we have to dynamically

- ➡ Allocate resources to competing jobs
 - ➡ Re-allocate resources based on system states
- ## ■ We developed and experimentally validated
- ➡ Adaptive and scalable job scheduling policy
 - ➡ Manage multiple job streams across multiple clusters
 - ➡ Good response-time and system utilization

Problem Statement

- **Given:** A set of computationally intensive jobs, $J = \{J_1, \dots, J_r\}$, that arrive stochastically
- A set of clusters, $S = \{C_1, \dots, C_n\}$, each cluster with
 - $P = \{P_1, \dots, P_m\}$ community-based workstations
 - $D = \{D_1, \dots, D_k\}$ sharable disks.
- **Problem:** Schedule the J jobs onto the S clusters with objectives of good
 - Mean response time, and
 - System utilization
- **Constraints:**
 - No advance knowledge of resource availability, arrival and service times of the jobs/tasks.
 - A task can execute only after all input data received
 - At most one task can access any data storage at any given time.

Multi-cluster Computing Infrastructure

Scheduling Policy

- All incoming jobs are submitted to the system scheduler
- The scheduler is composed of three core components
 - Self-Scheduling
 - Pull Algorithm
 - Push Algorithm

Self-Scheduling

- We associate with each node in the cluster tree a parameter called base load level.
- Whenever the current load level $<$ base load level:
 - ☞ Request for a set of computation from parent
 - ☞ A request that cannot be satisfied by a parent is backlogged and processed when jobs become available
 - ☞ A request may recursively ascend the cluster tree
- Note that a node can have only one pending request at any given time

Pull Algorithm

- It is a form of space-sharing policy
- It assigns a partition size to jobs based on
 - ☞ Parent-child relationship; and
 - ☞ Negotiations between the nodes in the cluster.
- When it is invoked, the algorithm performs the following steps:
 - ☞ Determines an ideal number of jobs that can be assigned to a child
 - ☞ Adjust the number of jobs transferred by taking the request from child into account
 - ☞ Dispatch the jobs to the child

Push Algorithm

- Intrinsically allows jobs/tasks to be relocated
 - ☞ from overloaded clusters to under loaded clusters
- When the push algorithm is invoked, it performs the following steps
 - ☞ First it determines the average load in the entire system
 - ☞ It then identifies those clusters that are above or below the average workload
 - ☞ Instruct over-loaded clusters to send a set of jobs to under-loaded clusters.

Performance Evaluation

- We used discrete event simulation to study the performance of the proposed policy and compare it with two baseline policies:
 - ☞ A derivative of the job-based time-sharing policy
 - ☞ An adaptive space-sharing policy originally introduced in Rosti et. al. and subsequently modified in Thanalapati and Dandamudi
- Metrics used
 - ☞ Mean response time - the sum of all job response time divided by the number of completed jobs
 - ☞ Average utilization - the job arrival rate times the mean service demand of the jobs divided by the number of processors in the system.

Performance Evaluation

- **We used** synthetic workload with different characteristics:

- ☞ **W1 – Matrix multiplication**
- ☞ **W2 – Parameter Sweep**
- ☞ **W3 – Burns Hat**

- **System Environment**

- ☞ **Shared Homogeneous Environment:** all processors run at the same speed, but workstations are shared by parallel and local workloads.
- ☞ **Dedicated Heterogeneous Environment:** No background workload, but processors with different speed.
- ☞ **Shared Heterogeneous Environment:** processors differ in speed and workstations may receive background workload at any time.

Simulation Results and Discussions

Simulation Results and Discussions

Simulation Results and Discussions

Thank You ...

