

# JavaSplit

*A Portable  
Distributed Runtime  
for Java*

**Michael Factor**

**IBM**

**Assaf Schuster**

**Technion**

**Konstantin Shagin**

**Technion**

# Outline

---

- The goal & features
- Related work
- Implementation overview
- Performance evaluation
- Conclusion & future work

# The Goal



***Create a runtime environment that executes a multithreaded Java application on a set of interconnected machines***

- Why Java?
  - Built-in multithreading and synchronization
  - Shared memory abstraction
  - Popular

# JavaSplit Features

- **Transparency**
  - Executes standard, possibly pre-existing, Java applications
  - The programmer is completely unaware of the distributed nature of the runtime
- **Portability**
  - Any machine can participate without regard of its operating system and hardware architecture
- **Scalability**
  - Designed to efficiently support a large number of nodes

# Related Work




- Distributed (non standard) JVMs
  - Java/DSM (1997), Cluster VM for Java (former cJVM) (1999), JESSICA(1999)
  - Sacrifice portability
- Translation to native code combined with a DSM
  - Hyperion (2001), Jackal (2001)
  - Also not portable
  - Perform compiler optimizations
- Systems built on top of standard JVMs
  - JavaParty(1997), ProActive (1998), JSDM (2001)
  - Introduce unorthodox programming constructs and style
- Unlike previous works, JavaSplit combines portability with transparency

# Java Basics

## Source code

```
class A {  
    ....  
    ...  
}
```



## Bytecode

```
0a0b0c0d  
0c626243  
b6d68816  
2a0b0c0d
```

## Java Virtual Machine

```
bacb0c0d  
0c623431  
c1d61836  
ab8ce321
```

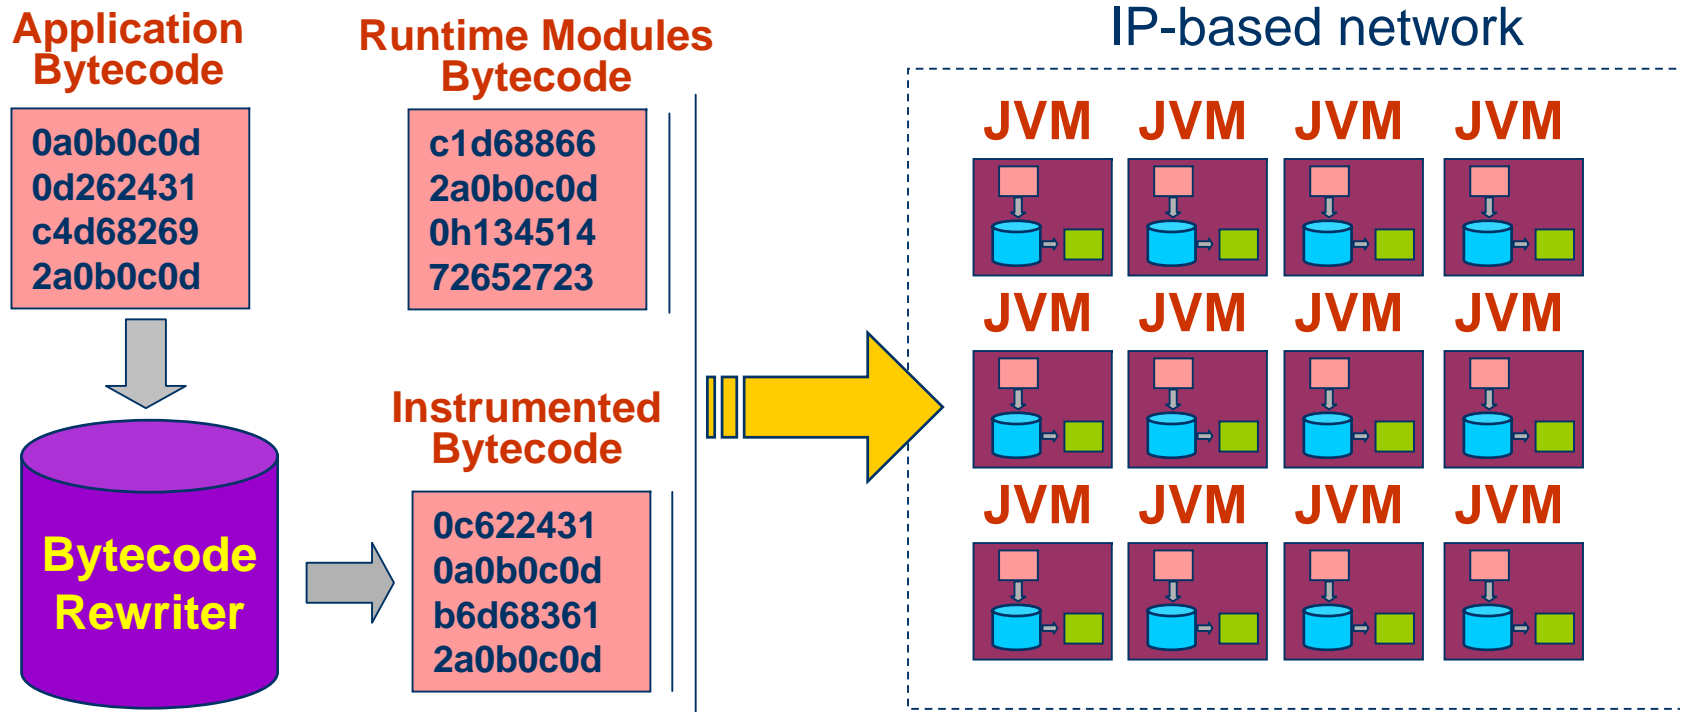
## Bootstrap classes



## Machine code

```
000110010  
010101011  
011100010  
100111010  
111110101  
011011110
```

# JavaSplit Overview



- Rewriting intercepts synchronization, accesses to shared data etc.
- Threads and objects are distributed among the machines
- Each node uses a **standard** Java Virtual Machine (JVM)

# Advantages of using standard JVMs

- Any machine with a JVM can be utilized
- Nodes can join the system using a Java-enabled browser
- No need to install any software
- Can use **just-in-time compiler (JIT)**
  - Unlike most distributed JVMs
- Can employ a standard **garbage collector**



# Distributed Shared Memory (DSM)

- Object-based
  - More suitable for Java
  - Few false-sharing
- Designed to be scalable
  - Never requires global cooperation
  - Allows multiple simultaneous writers
- Implements **Lazy Release Consistency**
  - Consistent with the proposed revisions to the **Java Memory Model (JMM)**
- Only objects accessed by more than one thread are managed by the DSM
  - Detected at runtime

# Instrumentation Details

- Classes are augmented with utility fields and methods
- The fields indicate the state and of an object
  - Inserted at the topmost hierarchy classes
  - The state data can be quickly accessed and easily disposed.
- Implementation of the utility methods is class-specific
  - Perform the same operation on each field of the class

```
class A extends C {  
    ...  
    // inherited utility fields  
    public byte  __JS__state;  
    public int   __JS__version;  
    public long  __JS__globalID;  
  
    public void  _JS_pack  
        (OutputStream out)  
        {...}  
    public void  _JS_unpack  
        (InputStream in)  
        {...}  
    public _JS_Diff _JS_compare  
        (JS.A other)  
        {...}  
}
```

# Read Access Check Example

...	
ALOAD 1	// load instance of A
<b>DUP</b>	// duplicate instance of A on stack
<b>GETFIELD</b>	A::byte __JS__state__
<b>IFNE</b>	// jump if the state allows reading
<b>DUP</b>	// duplicate instance of A on stack
<b>INVOKESTATIC</b>	JS.Handler::readMiss
<b>GETFIELD</b>	A::intField
...	

# Access Check Elimination

```
A aObject = new A();  
  
...  
<WRITE ACCESS CHECK of aObject>  
aObject.intField = 2003;  
    ... // no lock acquires  
<READ ACCESS CHECK of aObject>  
for(int k=0; k<N; k++){  
    ... // no lock acquires  
    <READ ACCESS CHECK of aObject>  
    System.out.println(k+aObject.intField);  
    ... // no lock acquires  
}
```

# Access Checks Overhead

[nanosec.]	Original	Rewritten	Slowdown
Field read	0.84	1.82	2.17
Field write	0.97	2.48	2.56
Static write	0.84	1.84	2.2
Static read	0.97	2.97	3.1
Array read	0.98	5.45	5.57
Array write	1.23	5.05	4.1

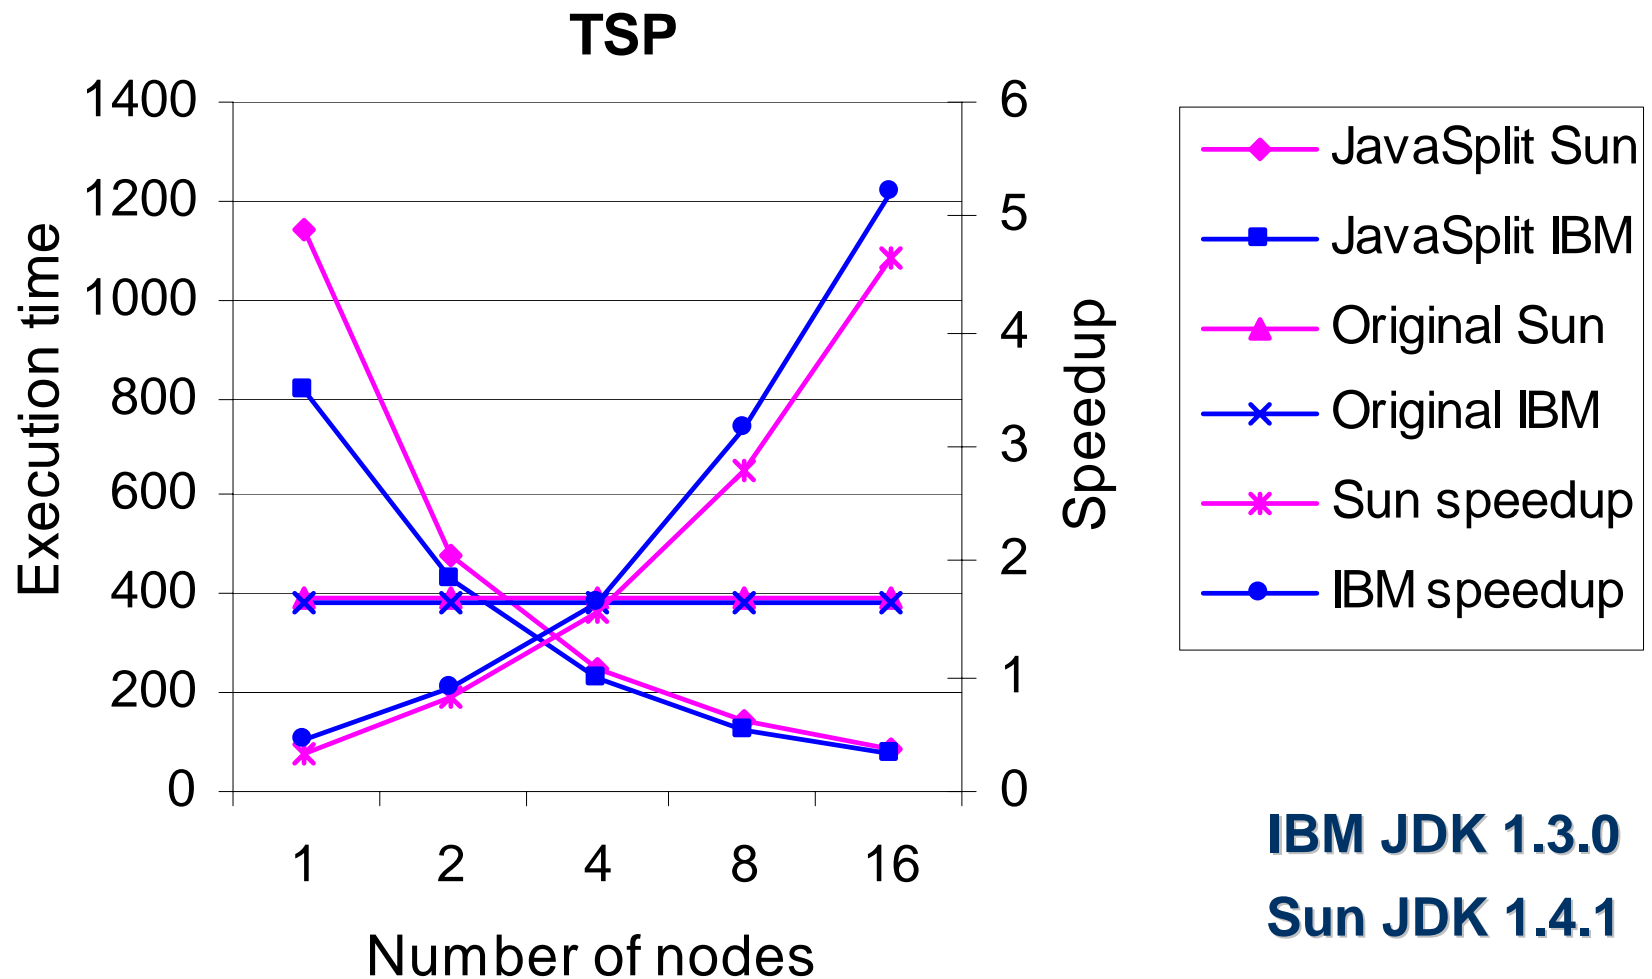
Sun JDK 1.4.1

# Efficient Synchronization

- Java applications contain a great amount of unneeded synchronization [Choi et. al., OOPSLA'99]
  - May cause significant performance degradation in instrumented classes
- We distinguish between synchronization operations on local and shared objects
  - Lightweight synchronization for local objects (similar to [Bacon et. al., PLDI'98])
  - Synchronization of local objects is cheaper than in original Java

[nanosec.]	Original	Local	Shared
Sun 1.4.1	90.6	19.6	281
IBM 1.3.0	93.4	54.7	327

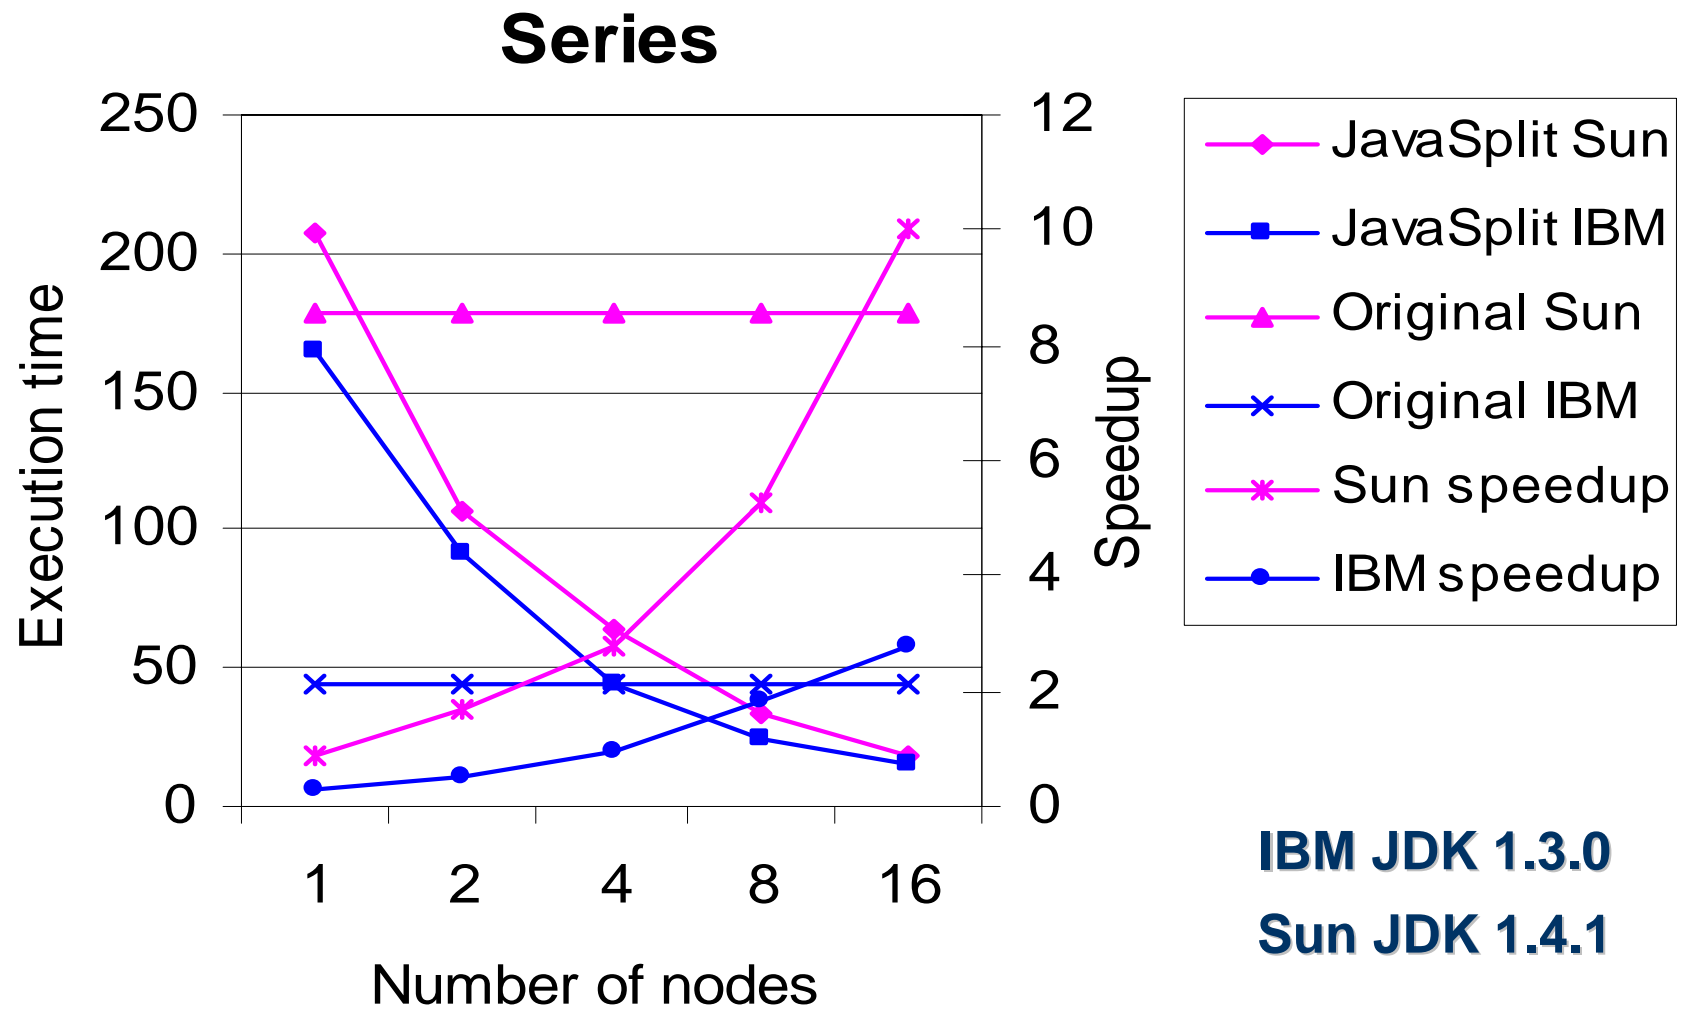
# Performance







# Performance (3)



# Conclusion

- JavaSplit is a first step towards creating a **convenient** and **portable** infrastructure for distributed computing
  - Provides shared memory abstraction
  - Can be used by any Java developer
  - Any platform with a JVM can participate
- Achieves scalable speedups executing computation-intensive applications
  - Despite relatively slow access to the network
  - With few simple optimizations

# Future Work

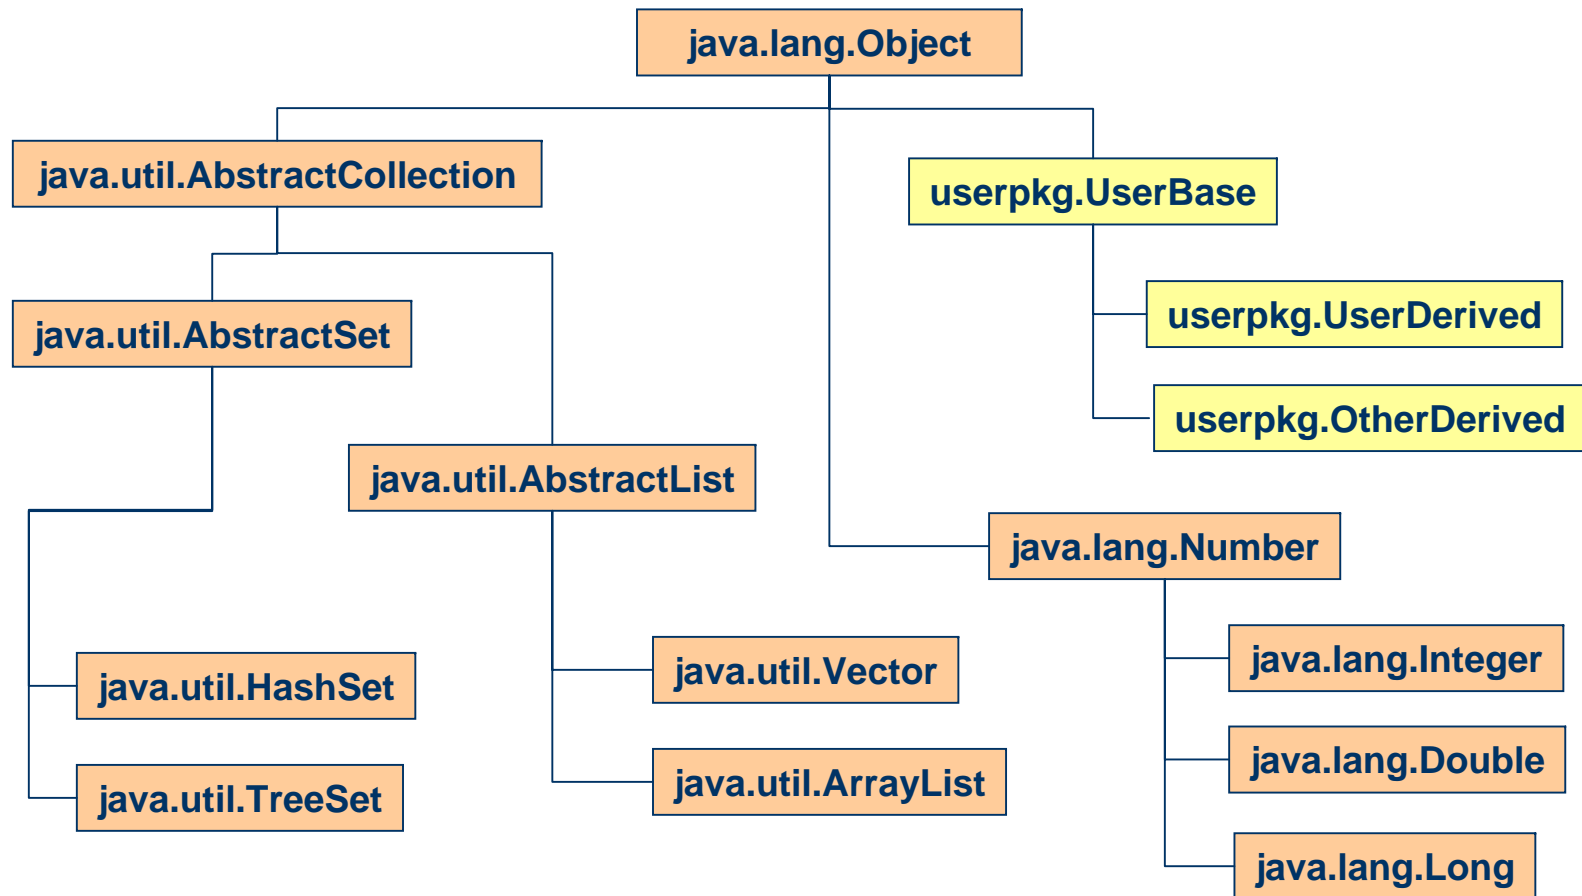
- DSM optimizations
  - Load balancing
  - Data distribution
- Compiler optimizations
  - Inter-procedural access check elimination
  - Static detection of local objects and classes
- Fault tolerance (High availability)
  - For wide-area cycle stealing
  - An algorithm is ready, need only to implement

**The End**

---

Questions?

# Original Class Hierarchy



# Twin Class Hierarchy

