

# Compiler Optimized Remote Method Invocation

Ronald Veldema

University of Erlangen-Nuremberg

Computer Science 2

Programming Languages

# Remote Method Invocation (RMI)

- Object-Oriented Remote Procedure Call
  - Polymorphism
  - Inheritance
- Heterogeneity
- Portable


# Problem Statement

- Java RMI for parallel programming too slow ?
  - Millisecond level
  - Many improvements already made
    - 30-200 microsecond level
      - PPopP'99 Manta-RMI
      - JavaGrande 99 KaRMI
      - etc

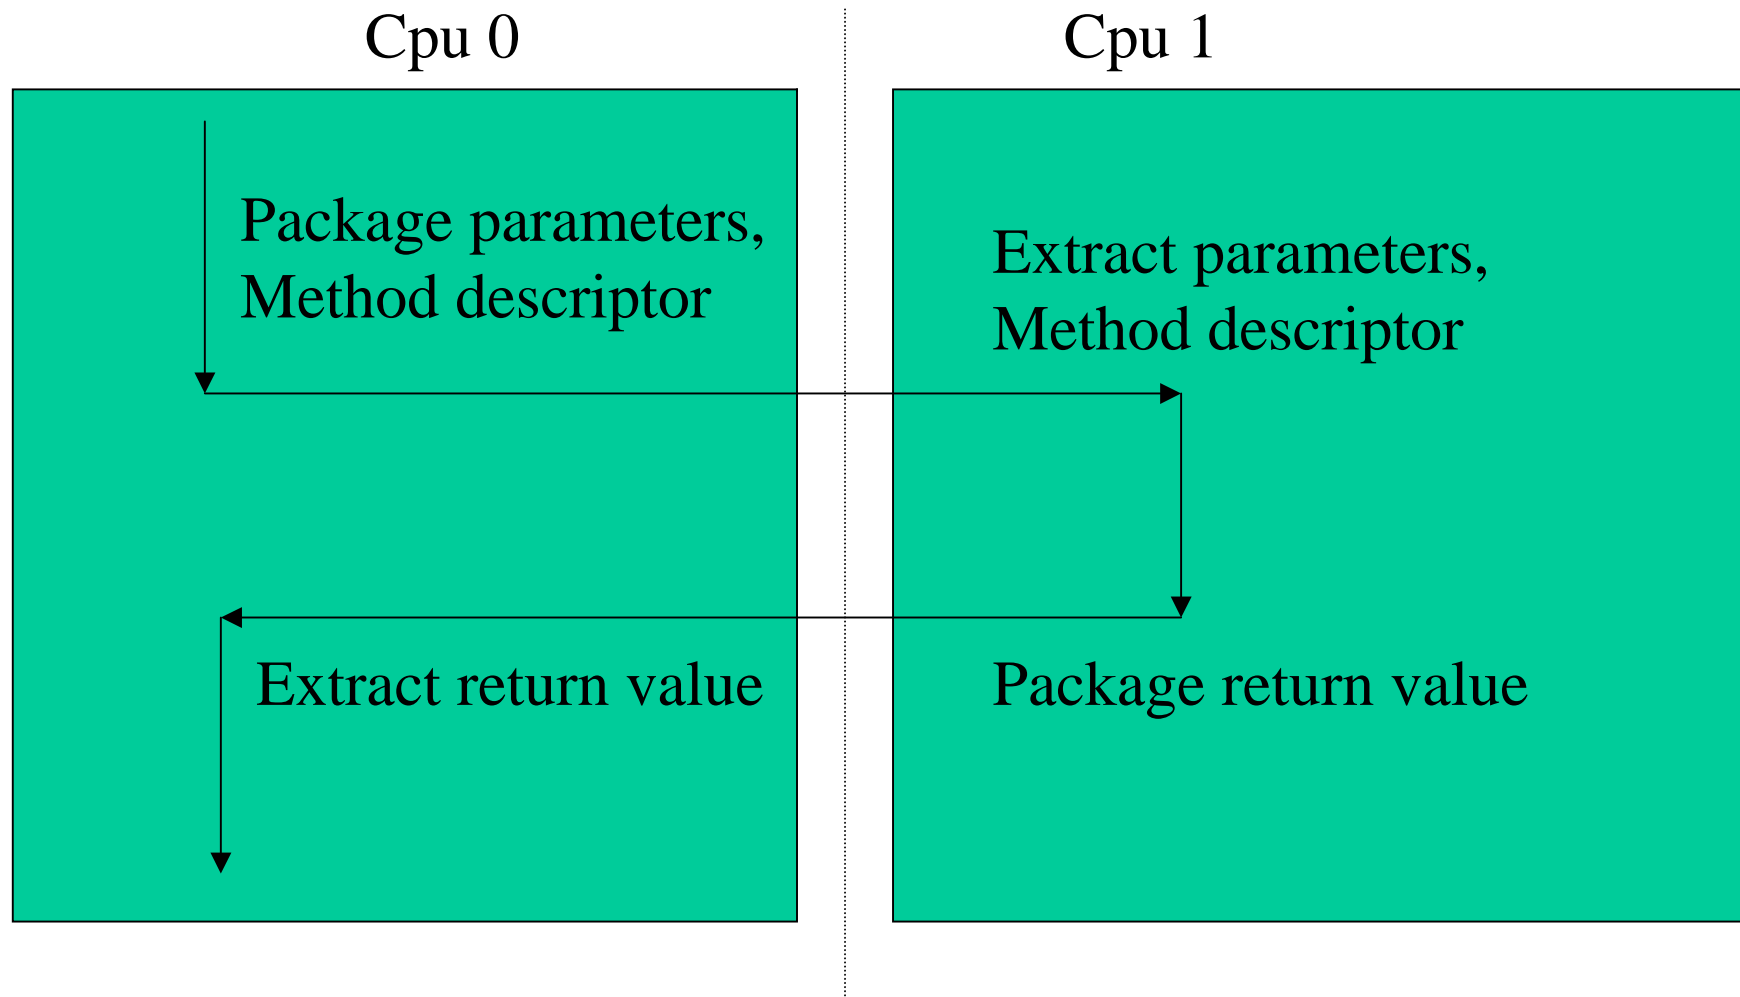
# JavaParty

- Layer on top of RMI
- JavaParty programs compile to RMI programs
- Advantages
  - No remote interfaces
  - No registry
  - Network exceptions masked by JavaParty
  - Object Distribution is handled by JavaParty RTS
  - Simpler programming model = code easier to analyze by a compiler

# RMI mechanism (1)



# RMI mechanism (2)



# Package Parameters: Earlier Work (1)

- Traverse the reachable object graph and **serialize** it to a byte array

```
class Data {  
    int temp;  
    String str;  
}
```

```
Data param1 = new Data();  
object.foo(param1);
```

```
Serialize_Data(Stream s, Data this) {  
    write_class_descriptor(s, Data.class);  
    Serialize_int(s, this.temp);  
    Serialize_Object(s, this.str);  
}
```

# Extract Parameters: Earlier Work (2)

- Read byte stream and reconstruct parameters by creating *new* objects

```
class Data {  
    int temp;  
    String str;  
}
```

```
Data param1 = new Data();  
object.foo(param1);
```

```
Object deserialize(Stream s) {  
    ObjectDescriptor d =  
        read_object_descriptor();  
    return d.deserializer(s);  
}
```

// for each class generate:

```
Object Data.deserializer(Stream s) {  
    Data this = new Data();  
    this.temp = read_int(s);  
    this.str = deserialize(s);  
    return this;  
}
```



# Problems With Naïve Approaches (1)

- Cannot handle cycles
  - The runtime system cannot look into the future
  - A compiler can look into the future
- For each deserialized object, a new object is created
  - The runtime system does not know if after an RMI the old object is garbage

# Problems With Naïve Approaches (2)

- Each reference causes an indirect call to locate the (de)serializer for that type
  - If the compiler knows for each reference what it points to, this can be avoided
- Search needs to be performed to locate the meta object for the deserialized object
  - If the compiler knows for each reference what it points to, this can be avoided

# Problems With Naïve Approaches (3)

- Cycle table lookups
  - Even for objects where programmer knows no cycles exist !

```
Serialize_Data(Stream s, Data this) {  
    if already_serialized(this) {  
        write_stream_backreference(s);  
    } else {  
        write_class_descriptor(s, Data.class);  
        Serialize_int(s, this.temp);  
        Serialize_Object(s, this.str);  
    }  
}
```

# Intermezzo: Heap Analysis

- Statically analyze which types of objects allocated from where in a program a reference can point to
- Algorithm
  - Convert program to SSA form
  - Assign to each ‘new’ a unique number
  - Propagate numbers throughout the program
  - Continue until a stable state is reached

# Heap Approximation (1)

```
class B {}

class A {
    B b;

    A() {
        B b = new B();
        this.b = b;
    }

    void foo() {
        // what do we know of "this.b" ?
    }

    public static void main(String args[]) {
        A a = new A();
        a.foo();
    }
}
```


# Heap Approximation (2)

```
class B {}
```

```
class A {  
    B b;
```

```
    A() {  
        B b = new B();  
        this.b = b;  
    }
```


// b = {1}



```
    void foo() {  
        // what do we know of "this.b" ?  
    }
```

```
    public static void main(String args[]) {  
        A a = new A();  
        a.foo();  
    }  
}
```

// a = {2}




# Heap Approximation (3)

```
class B {}
```

```
class A {  
    B b;
```

```
    A() {  
        B b = new B();  
        this.b = b;  
    }
```

```
// this = {2}  
// b = {1}
```




```
    void foo() {  
        // what do we know of "this.b" ?  
    }
```

```
// this = {2}
```

```
    public static void main(String args[]) {  
        A a = new A();  
        a.foo();  
    }  
}
```

```
// a = {2}
```



# Heap Approximation (4)

```
class B {}
```

```
class A {  
    B b;
```

```
    A() {  
        B b = new B();  
        this.b = b;    // {2}.b = {1}  
    }
```


```
    void foo() {  
        // what do we know of "this.b" ?  
    }
```

```
    public static void main(String args[]) {  
        A a = new A();  
        a.foo();  
    }  
}
```

// this = {2}  
// b = {1}


// this = {2}

// a = {2}





# Heap Approximation (5)






# Heap Approximation Results (1)

- Whole program analysis:
  - For each variable we know its possible origins
    - We know that variable *V* was the result of *new* from locations *X*, *Y*, *Z* in source code
    - For each reference we know where it may point to
    - For each object field we know where it may point to

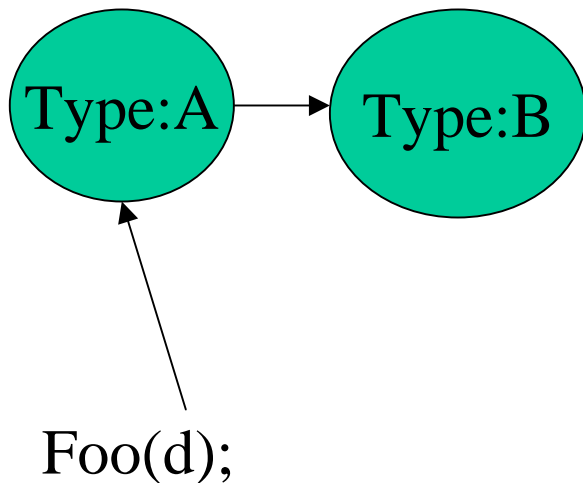
# Heap Approximation Results (2)

- Given a heap graph, we can see if the graph rooted in the RMI call's parameter contains cycles



# Heap Approximation Results (3)

- If we know the specific shape/types of nodes in the heap graph: generate specialized code *for that callsite*



ISO calling “foo” call:

```
call_foo_remotely(A d) {  
    write_method_descriptor(“foo”);  
    serialize_fields_of_A(d);  
    serialize_fields_of_B(d.b);  
}
```

# RMI Escape Analysis (1)

- If after an RMI the entire parameter tree is garbage, reuse the tree
  - Avoid garbage collection cost
  - Avoid costs of allocating new parameter objects
- Note: normal escape analysis is defined on single objects

# RMI Escape Analysis (2)

No !

```
// modified
static Data temp;

void foo(Data d) {
    temp = d;
}

void zoo(Data d) {
    temp = d.x.y.z.field;
}
```

Yes !

```
// read only:
static int temp;

void foo(Data d) {
    print (d.field);
}

int zoo(Data d) {
    return d.field * temp;
}
```

# Results

- 19 % performance gain for simple programs
- Most gain by generating specialized code for each callsite
  - Removes type lookups from code
  - Cycle elimination for duplicate reference tests
- Escape analysis doesn't help much
  - Garbage collector already efficient