

Optimizing Mechanisms for Latency Tolerance in Remote Memory Access Communication on Clusters

Jarek Nieplocha

Vinod Tipparaju

Manoj Krishnan

Pacific Northwest National Laboratory

Gopalakrishnan Santhanaraman

Dhabaleswar Panda

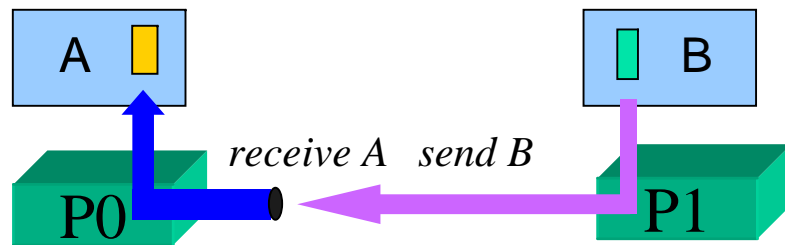
Ohio State University




Overview

- Background & Motivation
- Implementation of Mechanisms for Latency Tolerance
 - Nonblocking operations
 - Aggregation of small messages
- Performance Studies
 - Microbenchmarks
 - Application kernels
- Conclusions


Remote Memory Access and other Communication Models



message passing





remote memory access (RMA)



shared memory load/stores

Motivation: Processor-Network Speed Gap

1990



64-CPU
board

NCUBE/2
2MFLOP/s CPU
100 μ s latency
(35 μ s AM layer)
2.5MB/s B/W

12/02/2003

Today



Itanium-2/Myrinet cluster
6GFLOP/s CPU
10 μ s latency, 230MB/s (MPI)



Cray X1
12.8GFLOP/s CPU
9 μ s latency 10GB/s (MPI)


Improvement over 12 years

- CPU speed 10^4
- B/W 10^{2-4}
- Latency 10^1


What it means for efficient communication?

- Application level
 - increase problem granularity to scale
 - avoid small messages
 - explore algorithmic opportunities for latency hiding
- Comm. system level
 - effective nonblocking communication
 - message aggregation

Motivation: Processor-Memory Speed Gap



W. Wulf S. McKee, "Hitting the memory wall: Implications of the obvious", ACM Computer Architecture News, 1995



What it implies for efficient communication?

- Avoid processor touching data unnecessarily e.g.:
 - copying between user and special (DMA) buffers
 - data packing (sparse, strided)
 - format conversions
- Instead, we want
 - zero-copy communication
 - offload all/most data processing to the network adapter
 - including noncontiguous data
 - simplicity of RMA model attractive

Aggregate Remote Memory Copy Interface (ARMCI)

- Remote Memory Access Functionality
 - *put*, *get*, *accumulate* (also handles noncontiguous data transfers)
 - nonblocking communication interfaces
 - atomic *read-modify-write*, *mutexes* and *locks*
 - *memory allocation*
 - *fence* operations for managing memory consistency
- Characteristics
 - simpler & less synchronous model of RMA than MPI-2 1-sided
 - widespread portability (uses native protocols)
 - works with MPI
- Run-time component of the DoE Programming Models project
 - Co-Array Fortran compiler (Rice)
 - portable SHMEM (Ames lab)
 - Global Arrays library (PNNL)
 - Applications e.g., MPQC (Sandia)

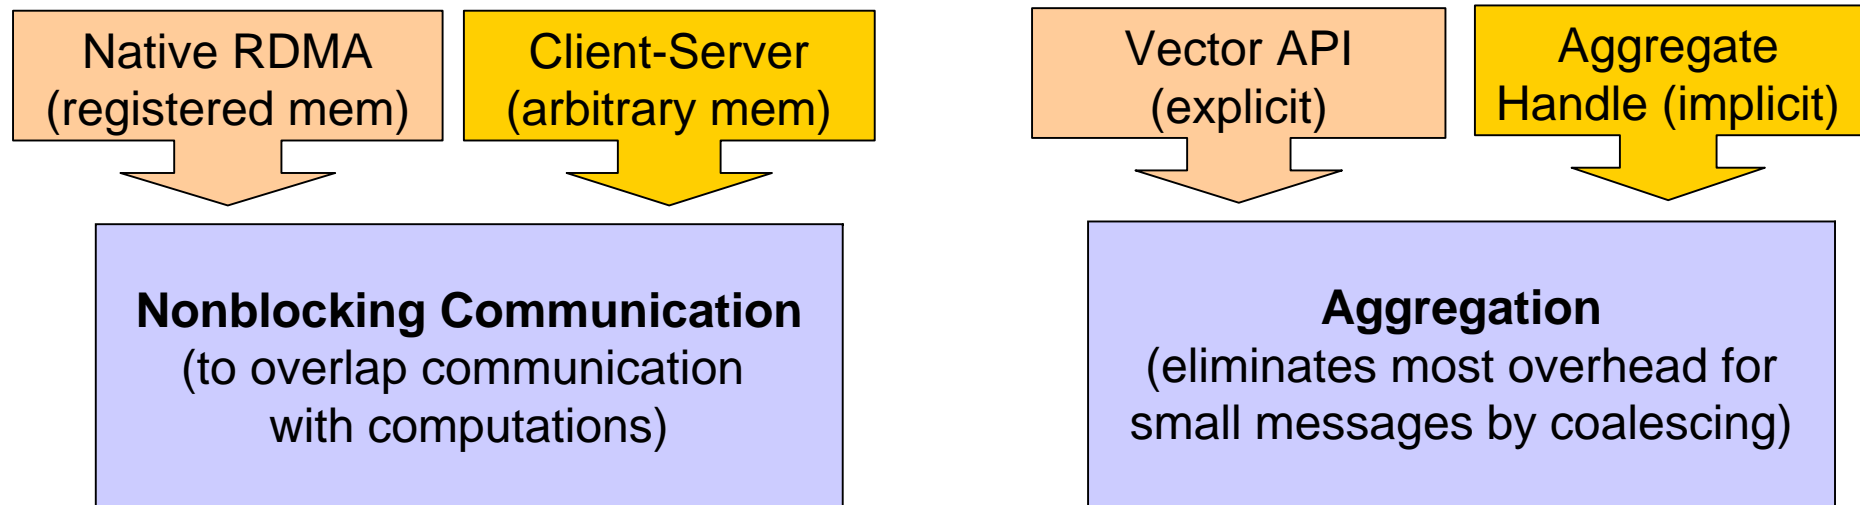
Memory Management Challenge

- User Communication Buffers
 - On systems with virtual memory, type of memory used for application communication buffers is *critical to performance*
 - *Good memory* (registered/pinned) that network adapter can communicate to/from directly (Myrinet, Infiniband, VIA, Hitachi SR8K)
 - *Good memory* on SMP nodes is shared memory (shmget/mmap)
 - Communication libraries designed to handle worst case scenario
 - on demand memory registration or data copying
 - this does not apply to Cray X1 and Quadrics (Elan-3, Elan-4)
- Our approach to memory management
 - Try to give user “good memory” suitable for fast communication
 - ARMCI_Malloc (registered + shared memory)
 - ARMCI_Malloc_local (registered)
 - If not, handle normal memory with some performance degradation
 - Want to provide effective nonblocking operations for both!

Two Mechanisms for Latency Tolerance

We implement portably two well known ideas:

- nonblocking coms: *Strumpen & Casavant 1994, Culler et al 1993*
- aggregation: *Pham & Albrecht 1999*



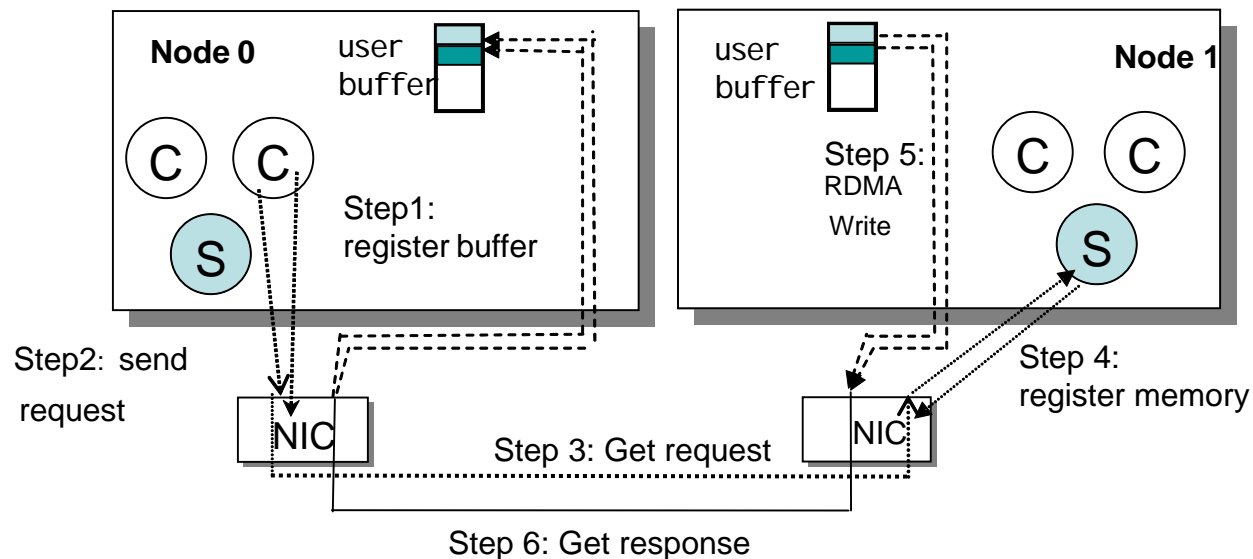
Details in the Paper:
descriptor and buffer management layers

Design Approach for Nonblocking Operations

- Aim for widespread portability
 - Implemented over GM, VIA, IBA, LAPI
 - Offer RMA (mostly in limited scope)
 - Papers on blocking RMA in CAC'02 and Cluster Computing 6, 2003
- Use multi-protocols to deliver best performance possible
 - switching between shared memory, zero-copy, buffered with and w/o pipelining
 - exploit and augment native vendor protocols
 - e.g., some not do well within the SMP node or miss noncontiguous data interfaces
- Manage multi-protocols transparently to the user
 - Track registered memory buffers
 - Request handle stores info on protocol used
- Pending requests
 - Either *active* or *completed* state
 - Older requests that use buffered protocols buffers might get internally completed in order to initiate a new request
 - A nonblocking call does not return w/o initiating network communication
 - Rather to be queued for future processing
 - Tries to provide effective overlapping

Example of Multiprotocols

- Get operation using dynamic memory registration
 - Server thread
 - RDMA put (write)
- Applicability
 - User memory not registered
 - For small messages it is less competitive than buffered protocols (registration cost)
 - Memory registered but network does not offer RDMA Read (GM 1.X, cLAN VIA)

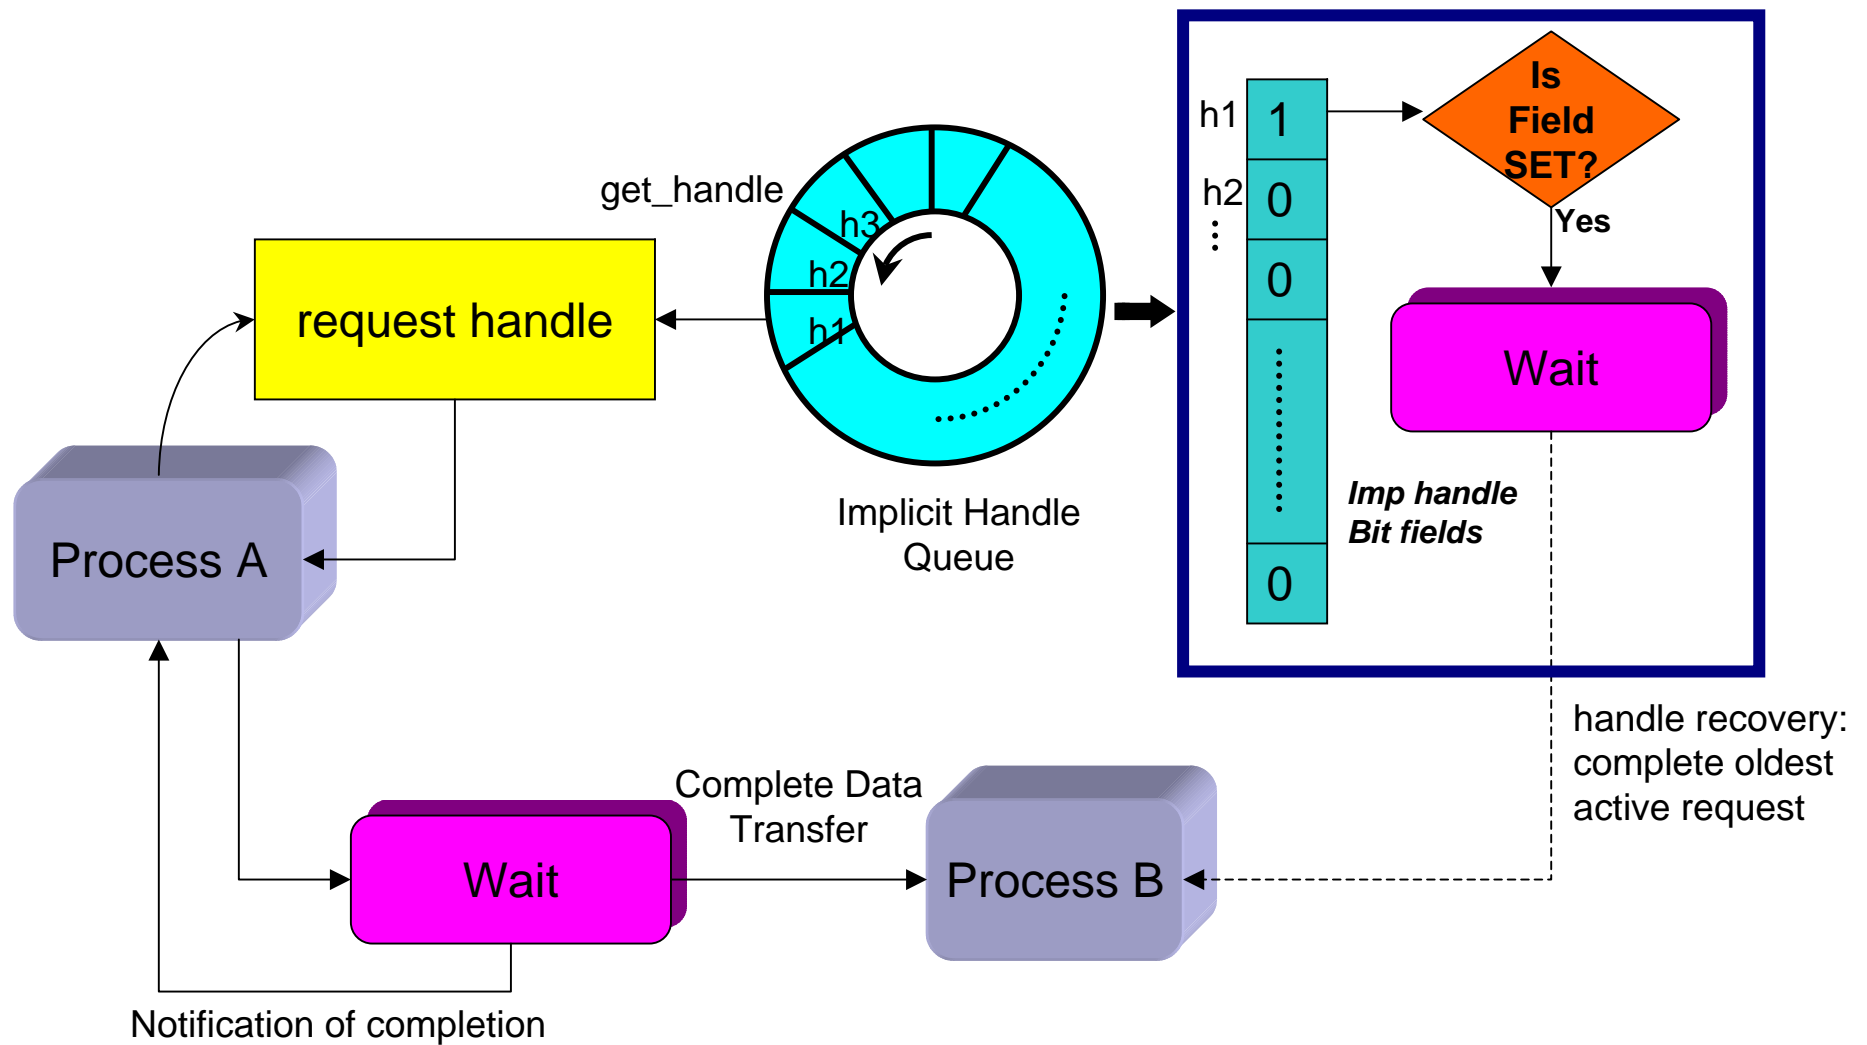


Request Handles

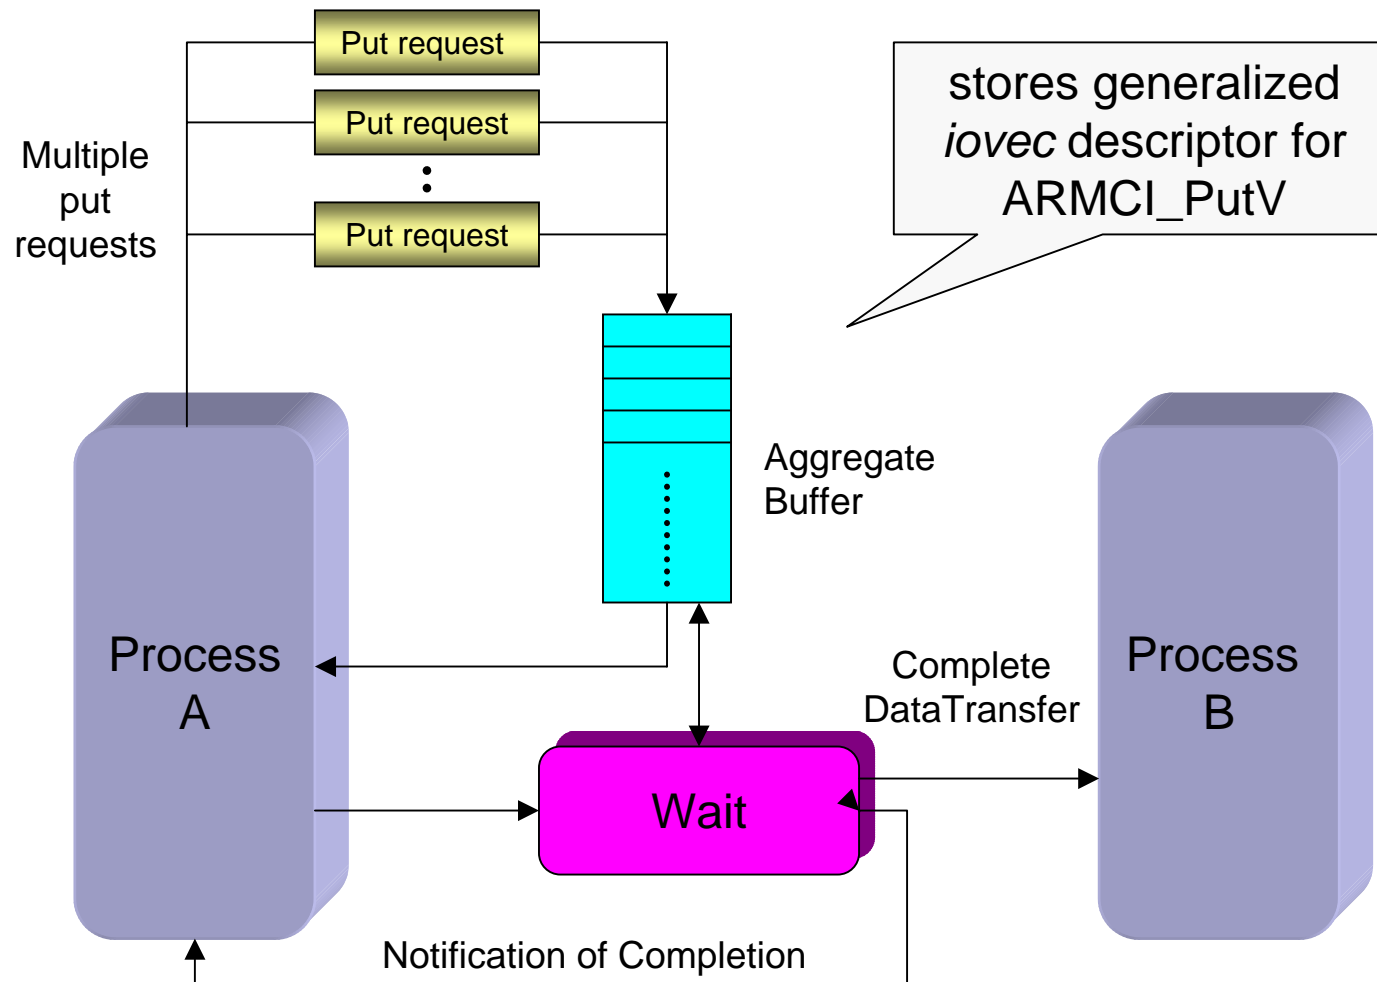
```
int ARMCI_NbPut(void* src, void *dst, int bytes,int proc, armci_hdl_t* handle)
```

- explicit handle
 - Used to wait for completion
 - Opaque data structure
 - RDMA descriptors
 - Buffer(s) info
 - other
 - User initialized
 - Can be used to aggregate messages
 - ARMCI_SET_AGEGATE macro
- implicit handle
 - handle=NULL
 - convenience to the user
 - No need to manage handles

Implicit Handle Non-blocking Transfer



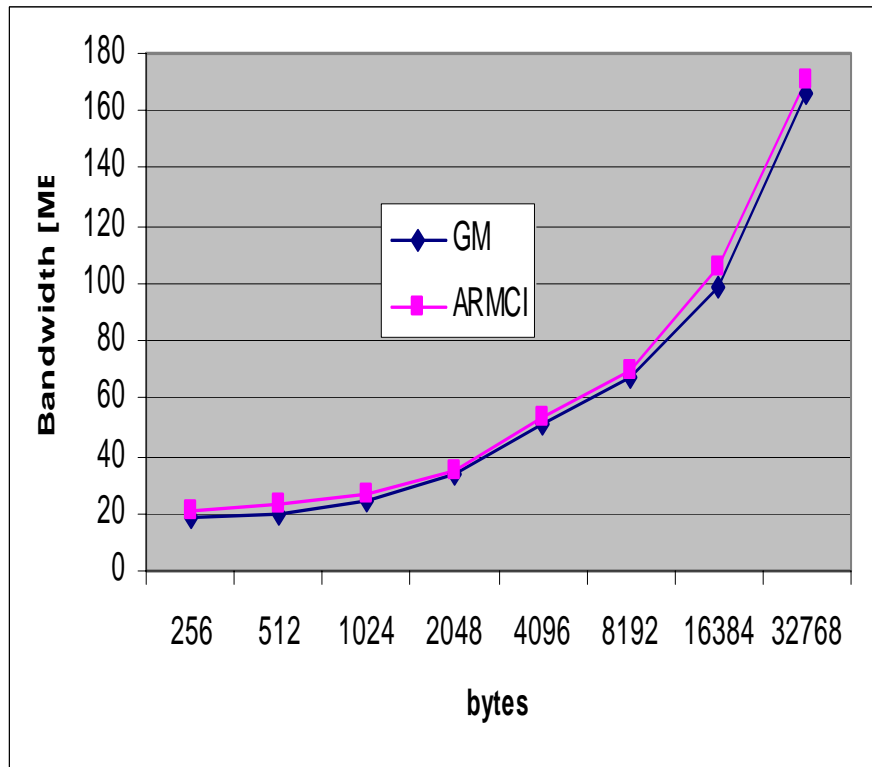
Message Aggregation: Put Example



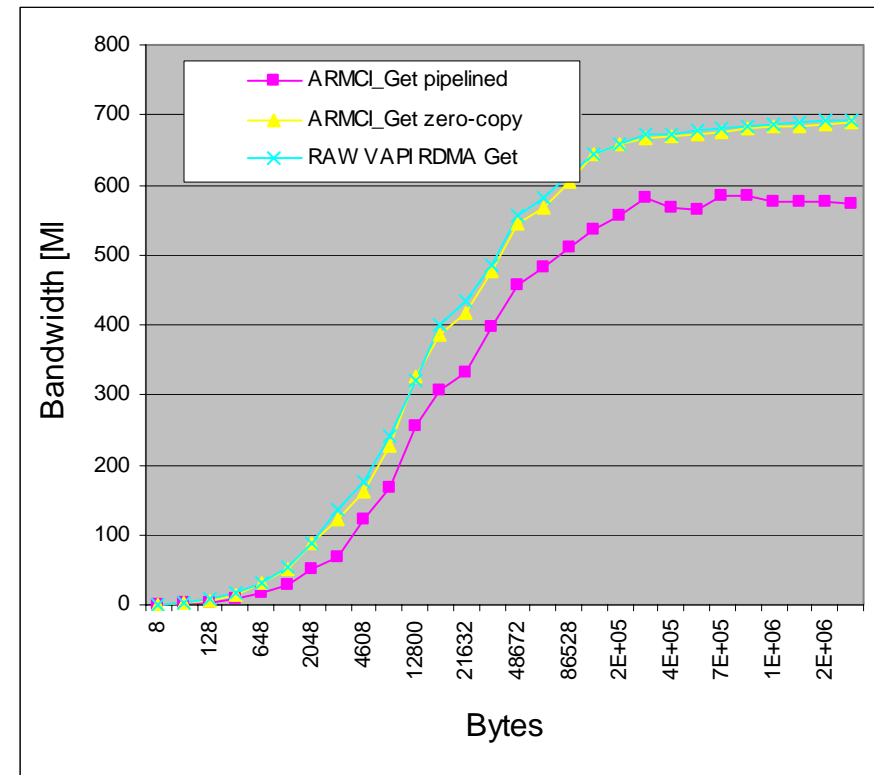
Experimental Results

- Platforms
 - Primary: dual CPU 2.4GHz P4 with Myrinet 2000
 - Secondary: Dual 1GHz Itanium-2 with Mellanox 4X IBA HCA
- Microbenchmarks
 - ARMCI overhead over native RDMA calls
 - Registered memory
 - nonregistered memory (pipelined protocols, see CAC'02 paper)
 - Potential for latency hiding in nonblocking operations
- Application Kernel Results
 - NAS CG & MG
 - Dense Matrix multiplication
 - Sparse Matrix-Vector multiplication

GET Overhead over Native RDMA Read



Myrinet/Pentium-4 GM 1.6

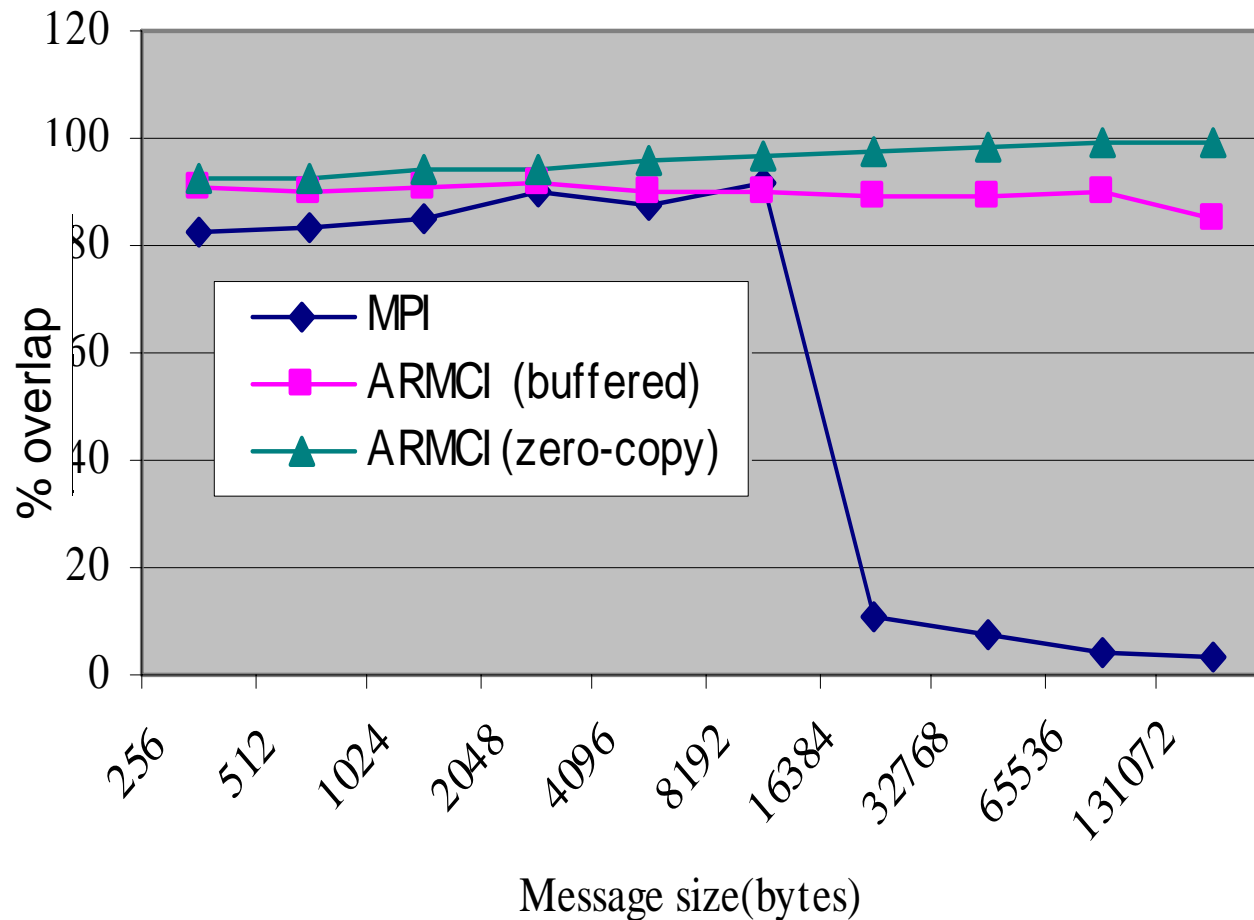


Infiniband/Itanium-2

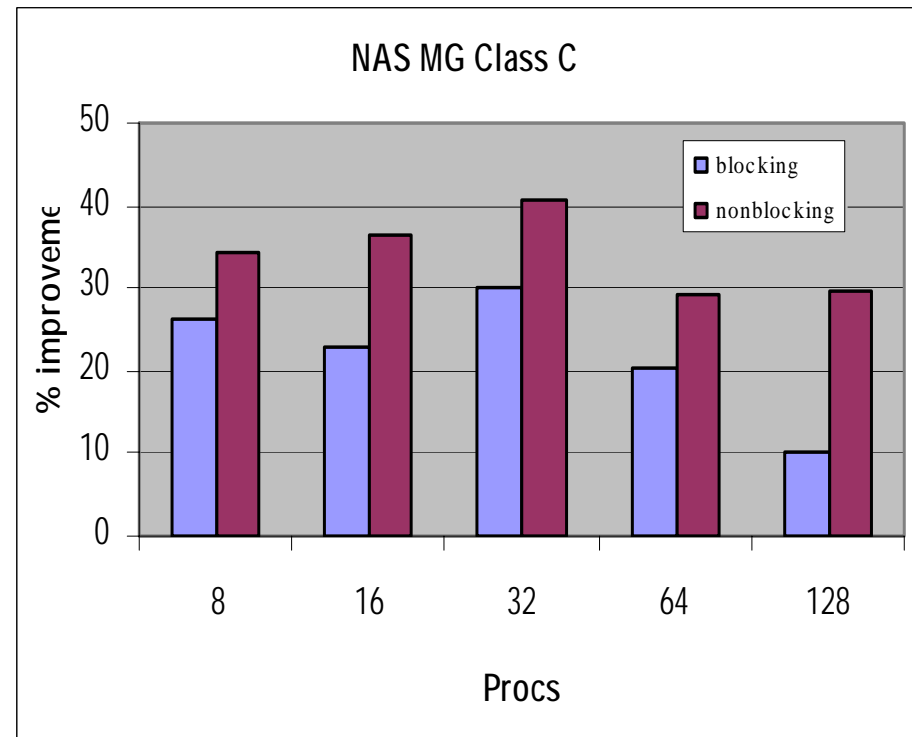

Overlapping Communication with Computations

Benchmark

- Issues a nonblocking call
- inserts an increasing amount of computation
- calls wait


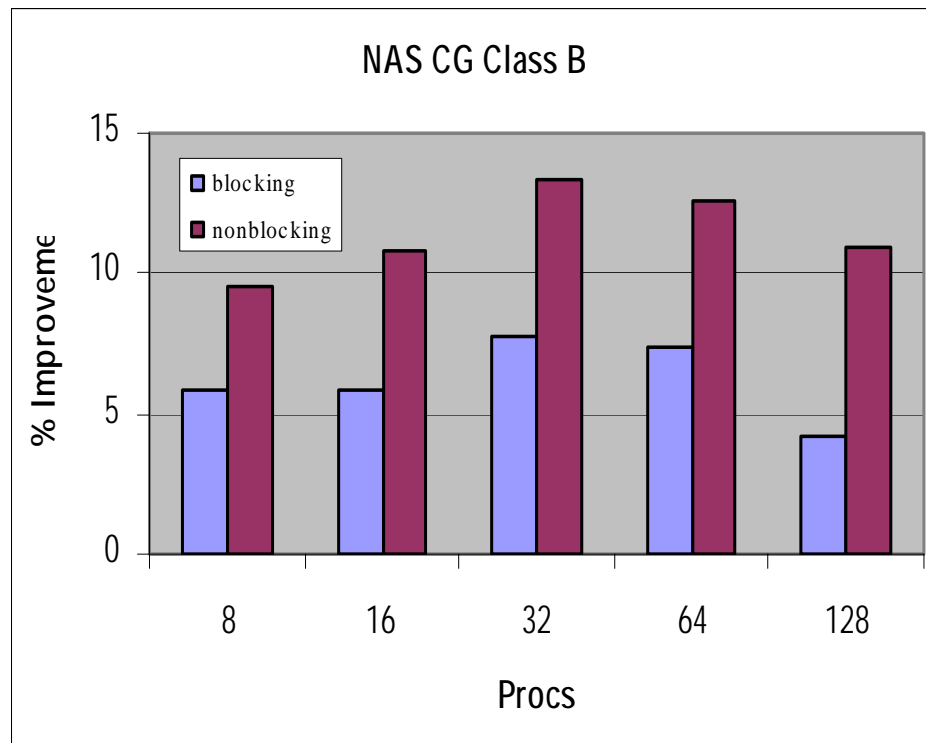


Improvement over MPI: NAS MG



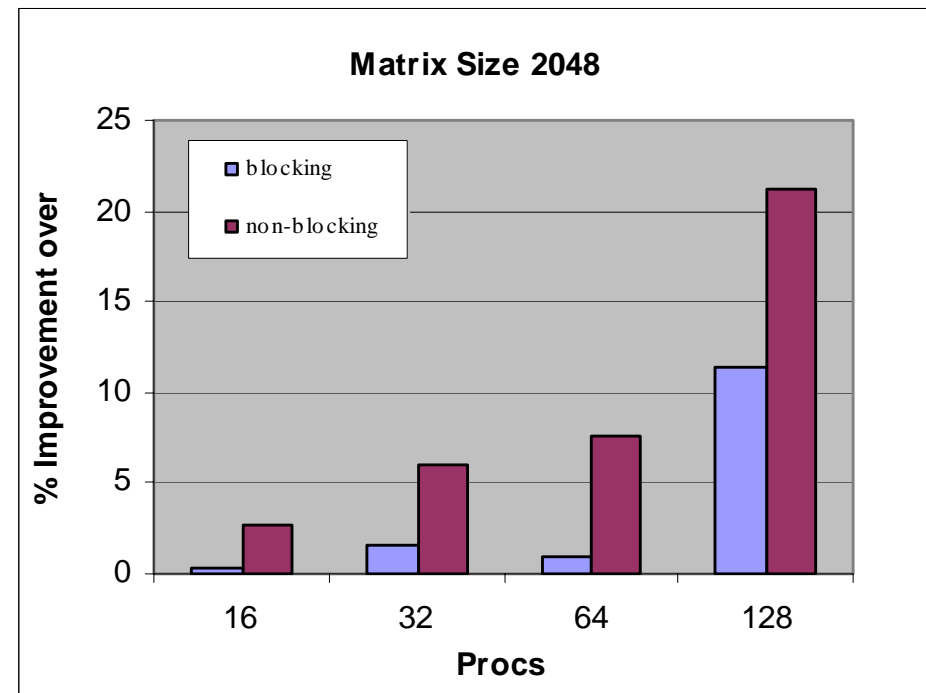
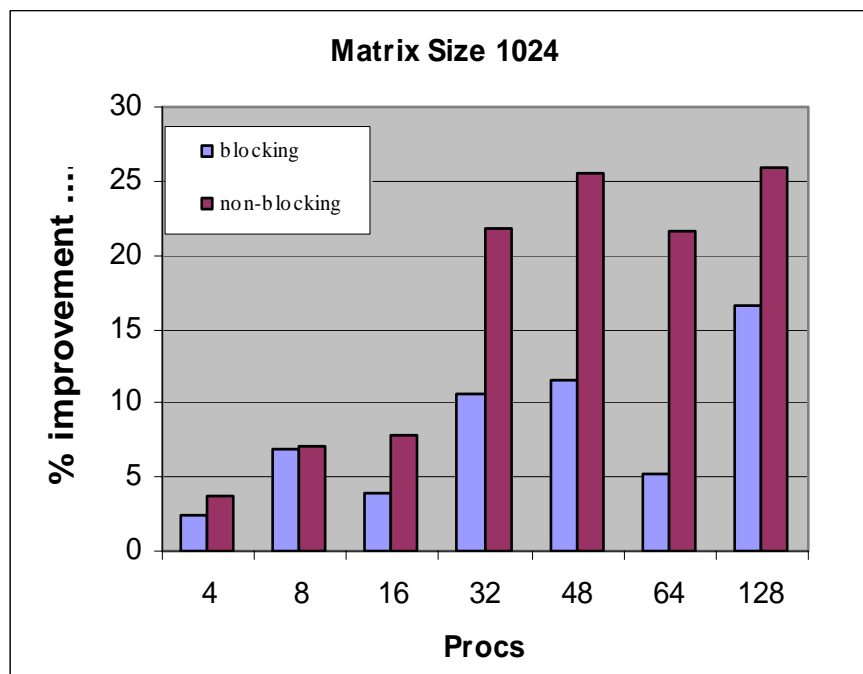
Linux, Pentium-4, Myrinet

Improvement over MPI: NAS CG



Linux, Pentium-4, Myrinet

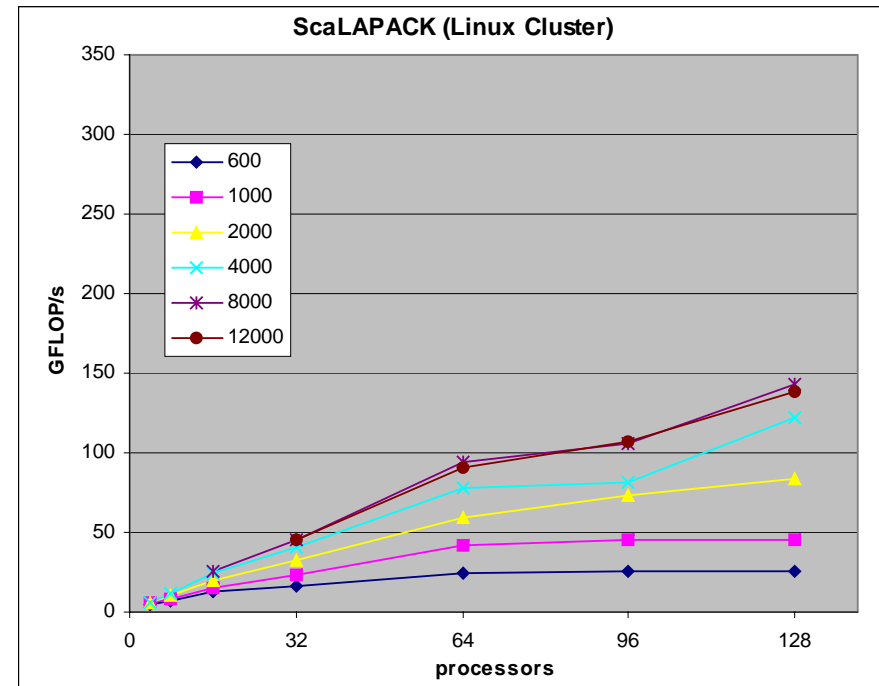
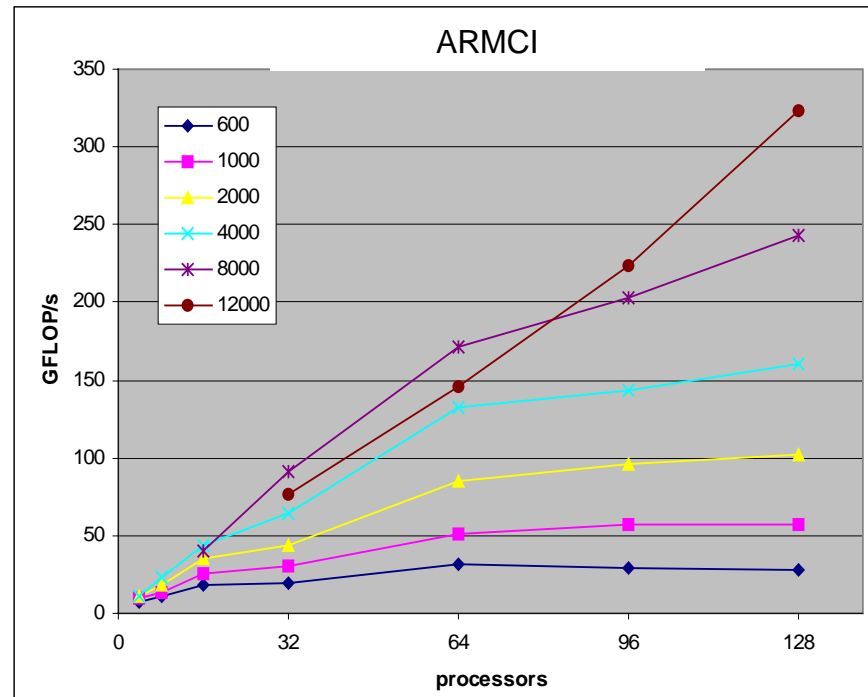
Matrix Multiplication SUMMA (MPI) and ARMCI



Linux, Pentium-4, Myrinet


Aggregate Performance

ARMCI and Scalapack/PBLAS pdgemm



Sparse Matrix Vector Multiplication


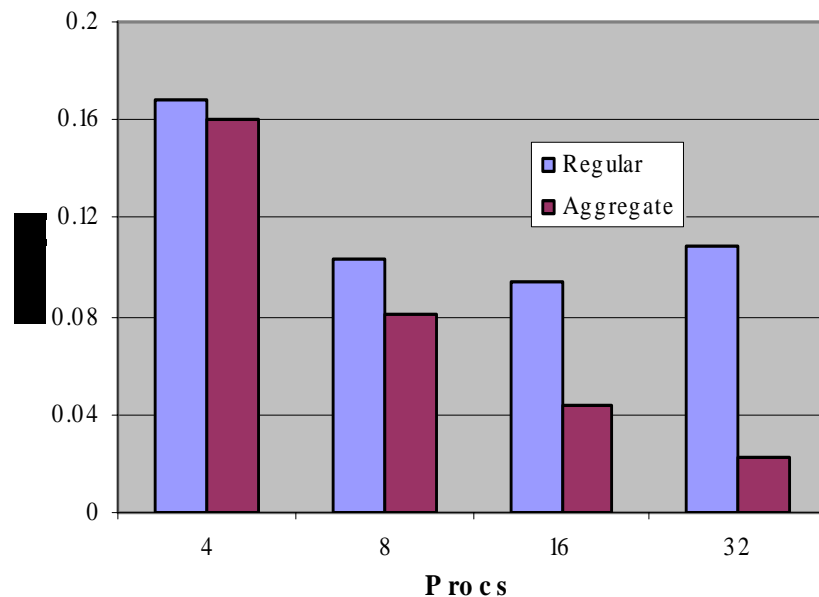
message aggregation



Harwell-Boeing Sparse Matrix collection

- a finite element problem
- $N = 41092$ Non-zero elements: 1683902 ($\sim 0.1\%$)

“aggregate handle” nonblocking ARMCI Get



Summary and Conclusions


- Mechanisms for latency tolerance increasingly important
 - Processor-Memory and Processor-Network performance gaps growing
- Paper focused on aggregation and nonblocking operations in context of RMA model
 - Efficient implementation on clusters
 - Handles underlying protocols transparently to the user through consistent and portable (GM, VIA, IBA, LAPI) interface
 - Zero-copy and buffered
 - Contiguous and noncontiguous data
 - Nonblocking operations provide real overlap
 - Effective implementation of small message aggregation
- Benchmarks show that these mechanisms help RMA model to deliver competitive performance to the send-receive model on clusters

Backup

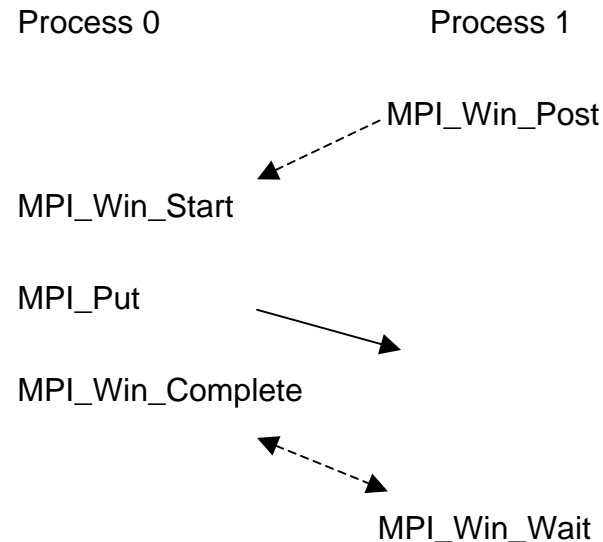
MPI-2 and other RMA models

Cray SHMEM


(IBM LAPI, GM, Elan, IBA similar)



MPI-2 1-Sided “active target”



MPI-2 1-Sided “passive target”



(Note: lock and put can be combined in networks that support active messages like IBM LAPI or sophisticated, user programmable adapters like Quadrics)

- MPI-2 1-sided is more synchronous than native RMA protocols
- Other RMA models decouple synchronization from data transfer