



Implications of a PIM Architectural Model for MPI

Keith Underwood
Sandia National Laboratories

Sandia is a multiprogram laboratory operated by Sandia Corporation, a Lockheed Martin Company,
for the United States Department of Energy under contract DE-AC04-94AL85000.






Overview

- **Growing use of COTS components in HPC (price pressure)**
- **Increasing gap between processor & memory speed**
- **Novel architectures on the horizon, but a huge investment exists in current parallel codes**
- **Next-next-Generation must...**
 - ... be high volume (commodity)
 - ... break the memory wall
 - ... support current applications



The Memory Wall

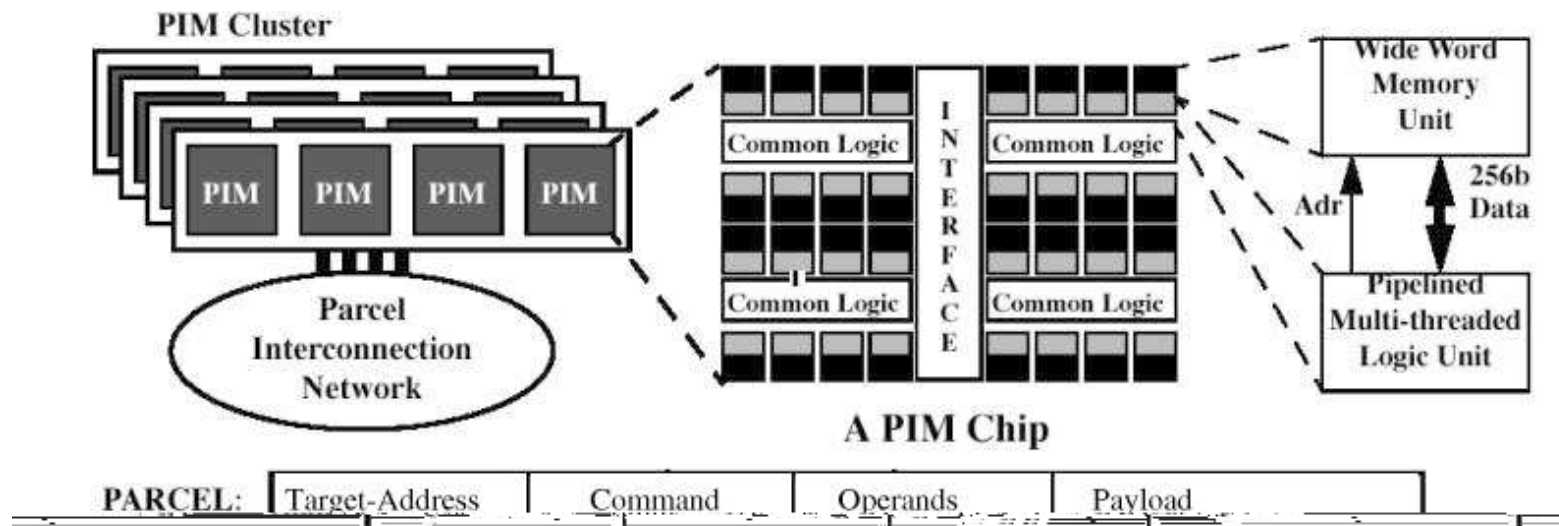
- Processor clock speed has grown dramatically
- Main memory access (DRAM) speed has grown much more slowly
- Caches and prefetching help, but problem is fundamental





PIM Systems

- Merge processing and memory on same die
- Very low latency to memory (< 10 ns)
- Simple, small, cheap multithreaded processor
- PIMs + interconnect = All-PIM supercomputer





PIM Programming Model

- **Pervasive Multithreading**
 - Hardware support for thread creation
 - Low-cost synchronization
 - Makes threading CHEAP!
- **Traveling threads**
 - Threads can migrate from PIM to PIM as required to access memory
 - Hardware support for thread state packaging and continuation on remote node
 - Send the thread to the data, not the data to the thread



MPI For PIM Implementation

- **Subset of MPI 1.2**
 - No collectives (other than MPI_Barrier)
 - MPI_COMM_WORLD only
 - Basic datatypes only
- **Explore PIM-specific issues**
 - Multi-threading
 - Thread migration



PIM Effects

- **Avoid “juggling”**
 - Each Request assigned a thread
 - No central thread which must iterate through each active and pending request
- **Hide latency**
 - High latency tasks (i.e. `memcpy()`) divided between many threads
- **Simplified implementation**
 - Traveling thread sends data + execution
 - Receiving node does not have to interpret incoming message. Message can “look after itself”



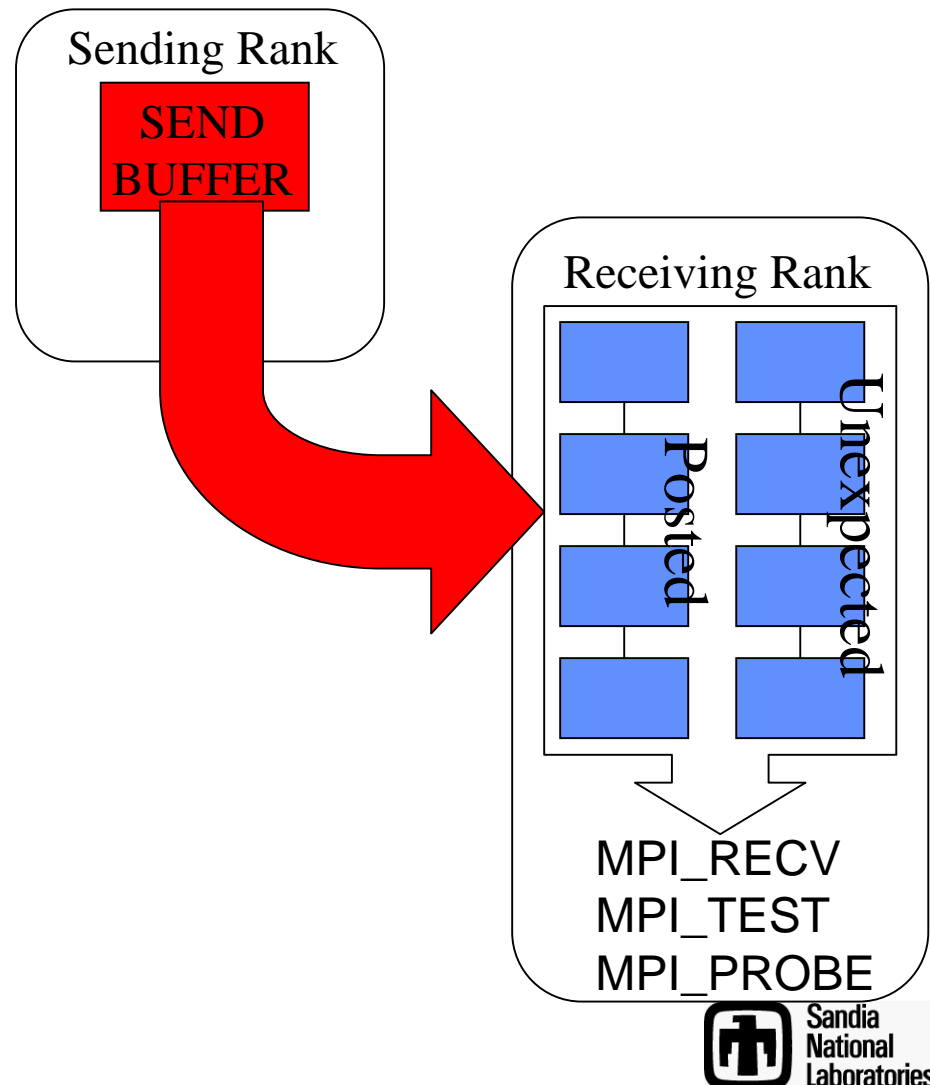
Related Work

- **Similar to MPI on Active messages**
 - MPI-FM, MPI-LAPI, MPI-AM
- **But, MPI-PIM does not require polling**
 - Saves processor cycles
 - Maintains MPI progress rule
- **Similar to RDMA MPIs**
 - Infiniband, MVICH, T3D
- **Incoming messages do not require processing by MPI library**
- **Traveling threads can “process themselves”**



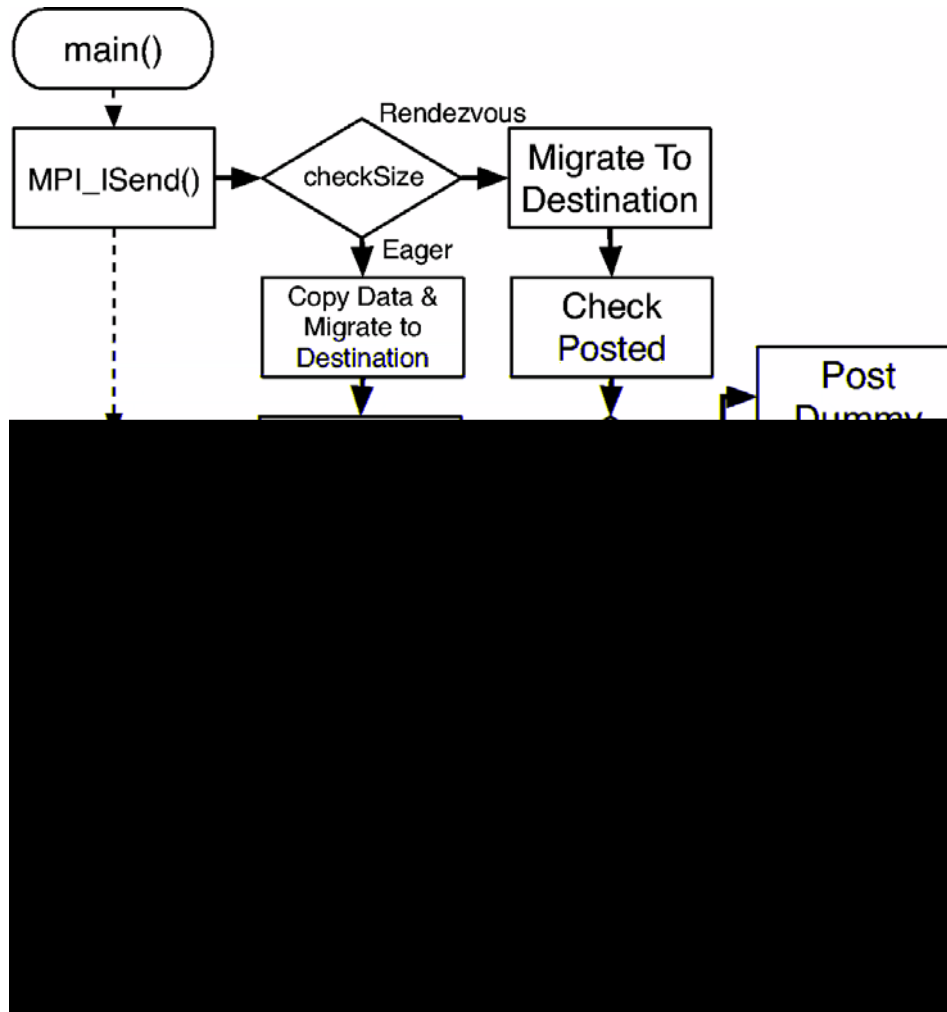
Basic Operation

- A thread is spawned to perform the MPI function
 - A “send thread” migrates to deliver data
 - Recv/Test threads check local state
- Posted and Unexpected Queues coordinate between threads





Isend() Implementation

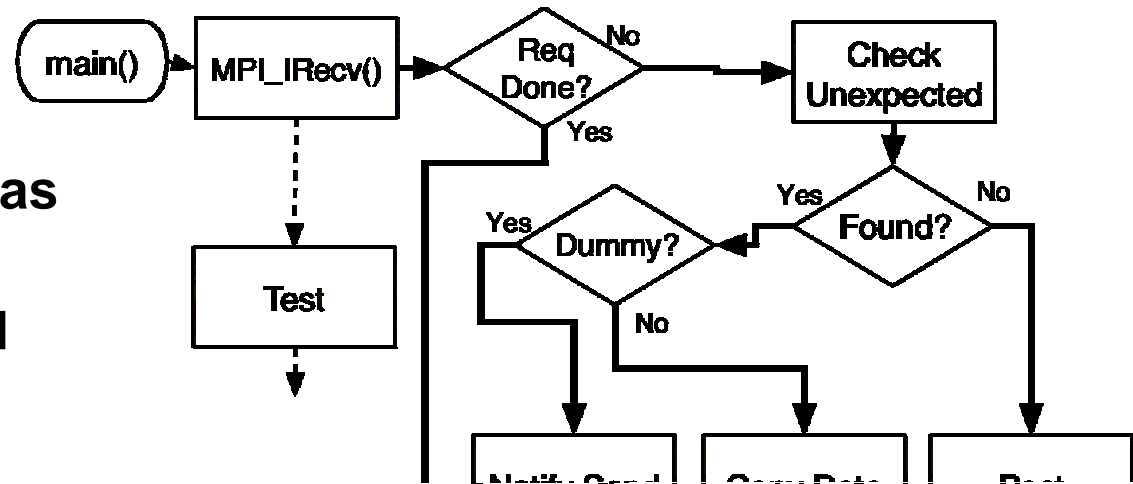


- Thread spawned to perform send
- Rendezvous
 - Migrate to destination
 - Find posted buffer, or wait for one
 - Retrieve data & deliver
- Eager
 - Copy data, migrate
 - Deliver to posted, or post to unexpected



IRecv()/Probe()

- IRecv()
 - Check if request has been completed
 - Check unexpected queue
- Probe()
 - Check unexpected Queue
- Ordering semantics preserved by locking





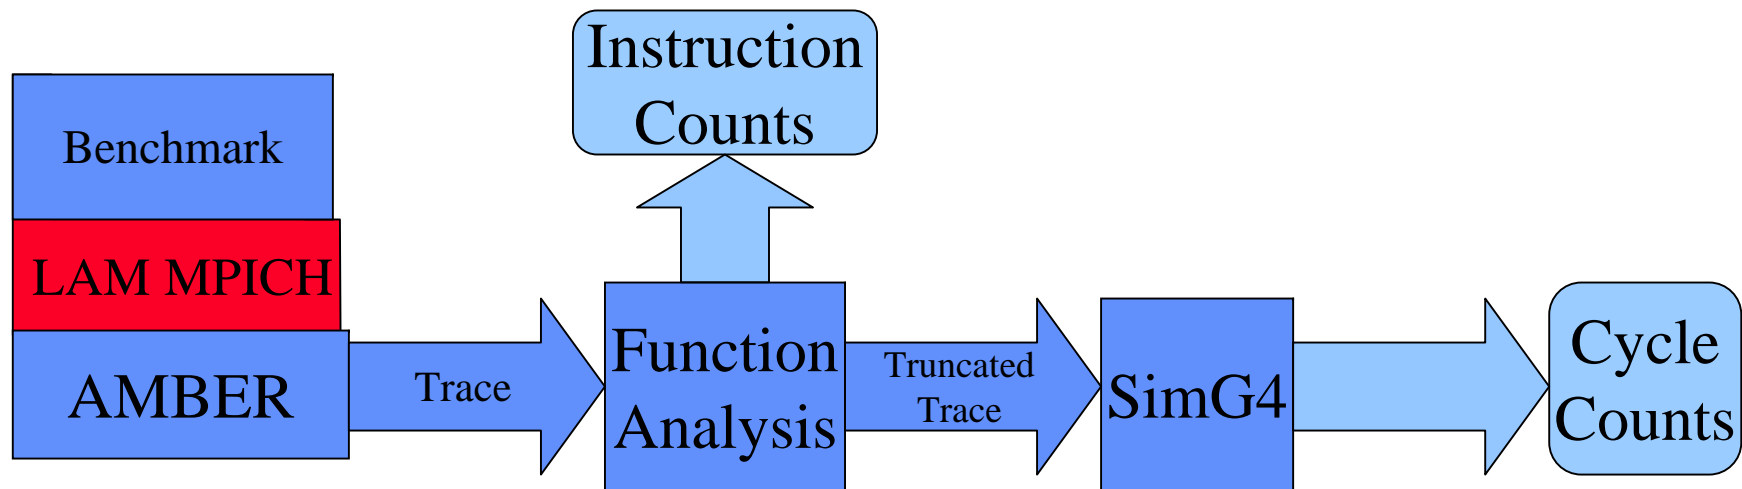
Test Methodology

- **Compare MPI for PIM against LAM MPI and MPICH**
 - LAM/MPICH traced on PowerPC, simulated on cycle accurate PowerPC simulator
 - MPI for PIM executed on PIM Simulators
- **Focus on MPI overhead**
 - OS & Network effects removed
 - Features not supported by MPI for PIM removed (error checking, user data types, etc.)



LAM MPICH

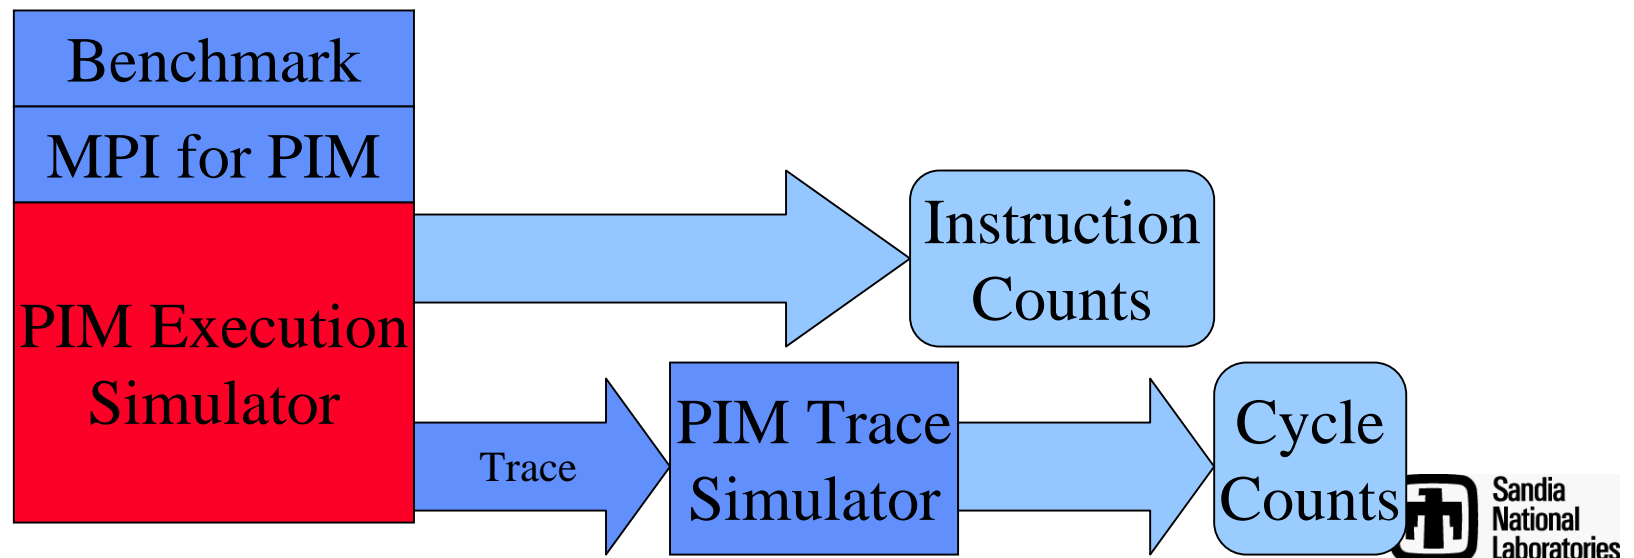
- **Amber captures trace**
- **Function Analysis**
 - Functions which concern features not present in MPI for PIM are removed (user data types, communicator lookup)
 - Functions categorized (setup, queue handling, etc...)
- **SimG4: Cycle accurate G4 simulator from Motorola**





MPI for PIM

- **PIM Execution Simulator**
 - MPI for PIM code annotated to categorize code
 - SimpleScalar with PIM-specific extensions for threading and synchronization
- **PIM Trace Simulator**
 - More detailed simulation





Processor Comparison

- **G4 (MPC7400)**
 - Superscalar OOO processor with vector unit
 - Traditional memory hierarchy
- **PIM**
 - Interwoven 1-wide multithreaded processor
 - Traveling thread execution model

	G4	PIM
Memory Latency, Open Page	20	4
Memory Latency, Closed Page	44	11
L2 Latency	6	N/A
Pipelines	7	1
Pipeline Depth	4 (Integer)	4 (Interwoven)



mpi-bw Benchmark

- **Mpi-bw** written to consider the effect of posted vs. unexpected messages
- Exercises a common subset of MPI


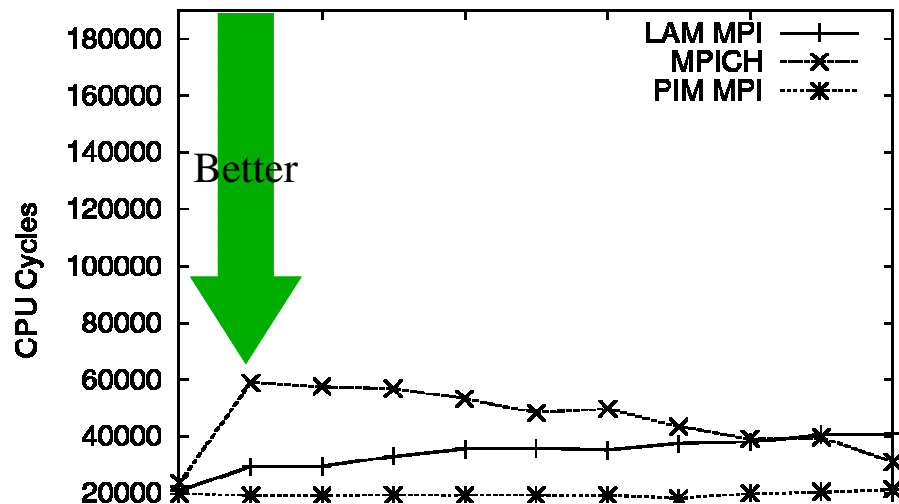
```
for (post=0; post<=10; post++) {  
    if (rank == 0) {  
        for (i = 0; i < post; i++)  
            MPI_Irecv(rmsg_buf[i], ...  
MPI_Barrier(MPI_COMM_WORLD);  
        for (i = 0; i < 10; i++)  
            MPI_Send(smsg_buf[i], ...  
        for (i = post; i < 10; i++) {  
            MPI_Probe(1, 0, ...  
            MPI_Recv(rmsg_buf[i], ...  
        }  
    }  
}
```

Abbreviated mpi-bw



MPI Performance: Overhead

- PIM MPI requires fewer cycles for MPI overhead
- More consistent over % posted



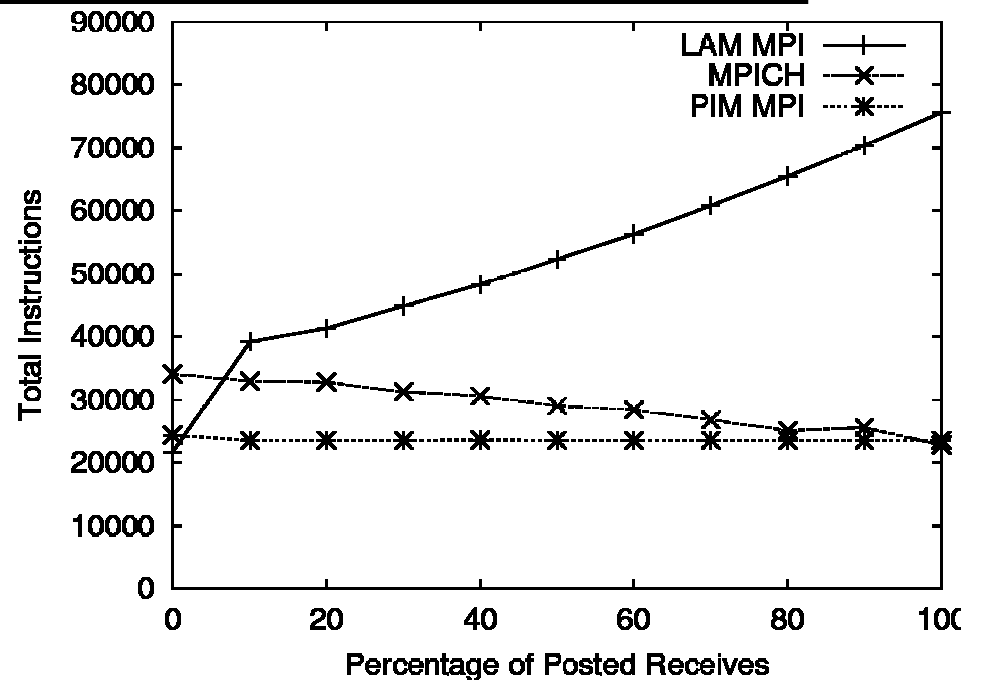
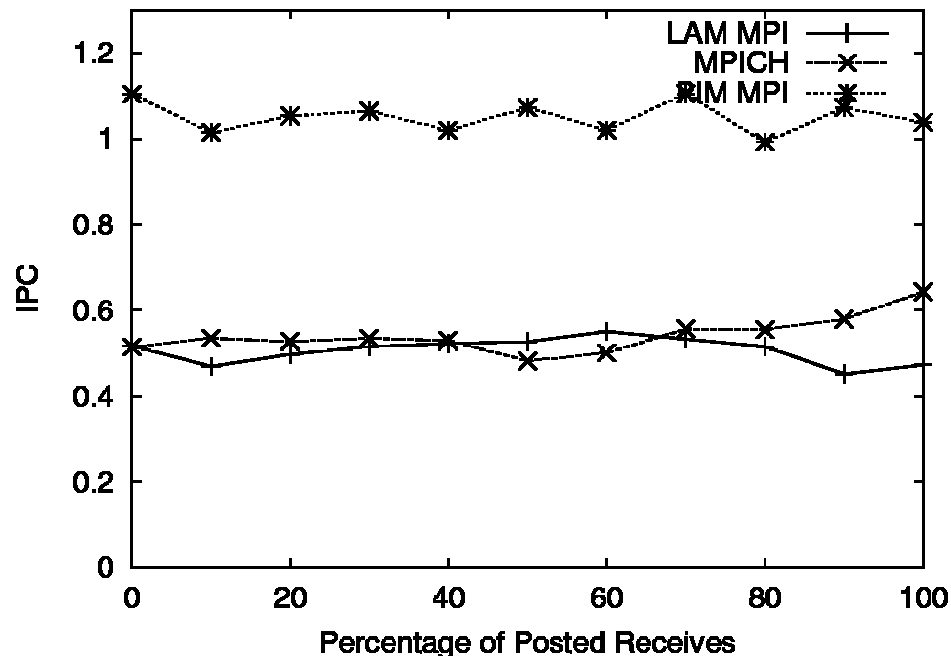
← Eager

↑ Rendezvous



Performance: Overhead

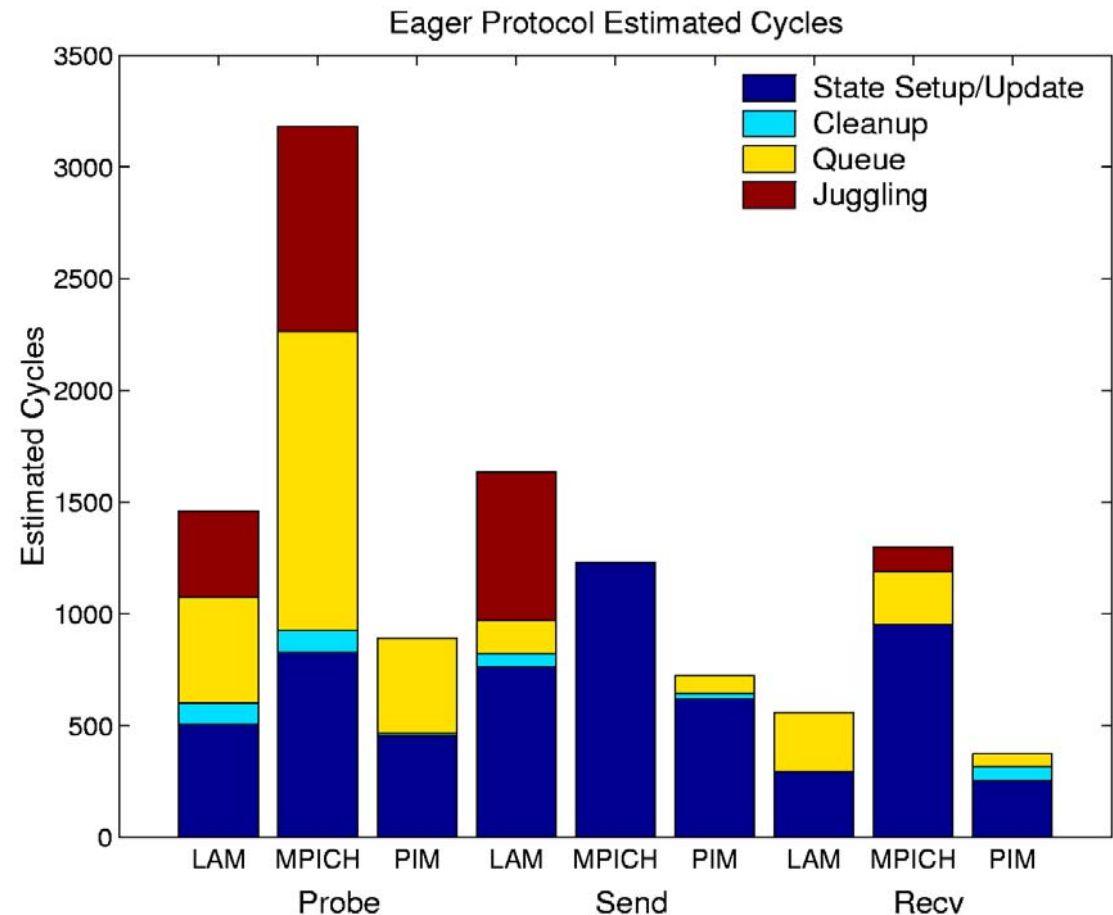
- MPI-PIM requires fewer instructions
- MPI-PIM IPC much higher






Function Analysis

- **Function Categories**
 - **Setup: State initialization and Update**
 - **Cleanup: Resource deallocation**
 - **Queue: inserting, interacting through lists**
 - **Juggling: MPI context switches**
- **PIM removes “Juggling”**





Overhead: PIM Benefits

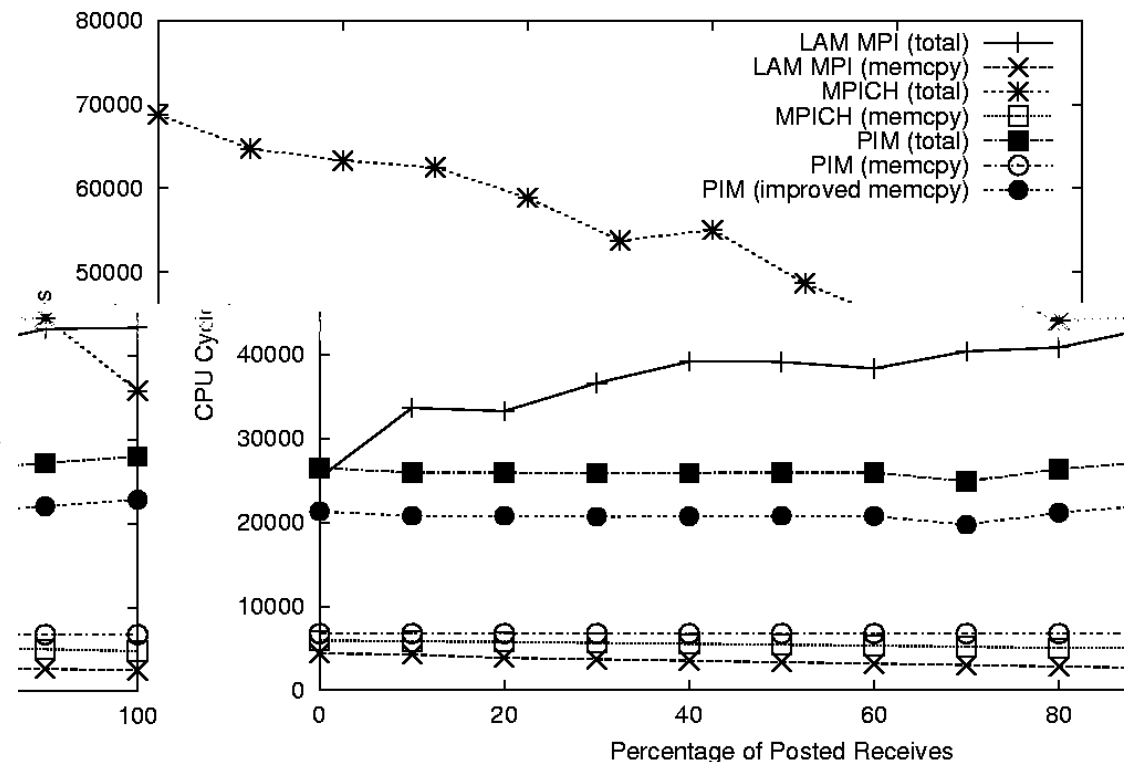


- PIM avoids juggling
- Higher IPC
 - Faster memory access
 - Fewer branch penalties
- Traveling thread decreases state setup



Other Benefits

- **Massive on-chip bandwidth allows fast memcpy()**
 - Closer to memory
 - Can copy entire open row at a time
- **Reduce negative effects of unexpected messages**





Conclusion

- **MPI maps onto PIM effectively**
- **Performance improved by**
 - Multithreading
 - Traveling threads
 - Fast memory copy
- **Implementation simplified**



Future Work

- **Expand to full MPI implementation**
- **More detailed simulation**
- **Optimized collectives**
- **One-sided communication**



Acknowledgments

- **Arun Rodriques and Rich Murphy**
 - University of Notre Dame students
 - The guys who did the work
- **Peter Kogge and Jay Brockman**
 - University of Notre Dame faculty
 - The guys who provided the students