



Comparing Map-Reduce and FREERIDE for Data-Intensive Applications

Wei Jiang, Vignesh T. Ravi and Gagan Agrawal

Outline

- ▶ Introduction
- ▶ Hadoop/MapReduce
- ▶ FREERIDE
- ▶ Case Studies
- ▶ Experimental Results
- ▶ Conclusion

Motivation

- ▶ Growing need for analysis of large scale data
 - ▶ Scientific
 - ▶ Commercial
- ▶ Data-intensive Supercomputing (DISC)
- ▶ Map-Reduce has received a lot of attention
 - ▶ Database and Datamining communities
 - ▶ High performance computing community
 - ▶ E.g. this conference !!

Map-Reduce: Positives and Questions

- ▶ Positives:
 - ▶ Simple API
 - ▶ Functional language based
 - ▶ Very easy to learn
 - ▶ Support for fault-tolerance
 - ▶ Important for very large-scale clusters
- ▶ Questions
 - ▶ Performance?
 - ▶ Comparison with other approaches
 - ▶ Suitability for different class of applications?

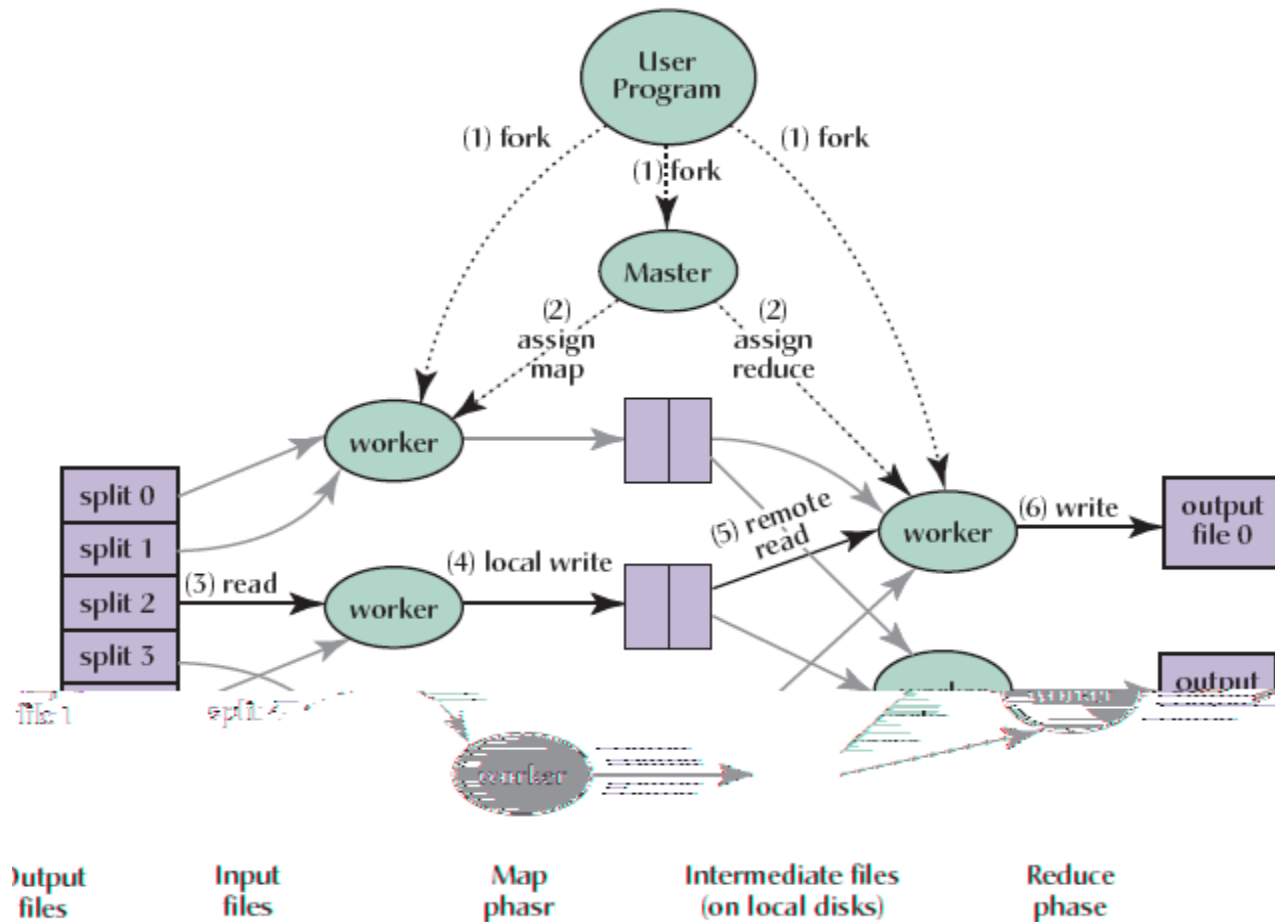
Class of Data-Intensive Applications

- ▶ Many different types of applications
 - ▶ Data-center kind of applications
 - ▶ Data scans, sorting, indexing
 - ▶ More “compute-intensive” data-intensive applications
 - ▶ Machine learning, data mining, NLP
 - ▶ Map-reduce / Hadoop being widely used for this class
 - ▶ Standard Database Operations
 - ▶ Sigmod 2009 paper compares Hadoop with Databases and OLAP systems
- ▶ What is Map-reduce suitable for?
- ▶ What are the alternatives?
 - ▶ MPI/OpenMP/Pthreads – too low level?

This Paper

- ▶ Compares Hadoop with FREERIDE
 - ▶ Developed at Ohio State 2001 – 2003
 - ▶ High-level API than MPI/OpenMP
 - ▶ Supports disk-resident data
- ▶ Comparison for
 - ▶ Data Mining Applications
 - ▶ A simple data-center application, word-count
 - ▶ Compare performance and API
 - ▶ Understand performance overheads
- ▶ Will an alternative API be better for “Map-Reduce”?

Map-Reduce Execution



Hadoop Implementation

▶ HDFS

- ▶ Almost GFS, but no file update
- ▶ Cannot be directly mounted by an existing operating system

▶ Fault tolerance

- ▶ Name node
- ▶ Job Tracker
- ▶ Task Tracker

Hadoop - More Details

- ▶ Locality---schedule a map task near a replica of the corresponding input data
- ▶ Combiner---use local reduction to save the network bandwidth
- ▶ Backup tasks---alleviate the problem of stragglers---not available in Hadoop

FREERIDE: GOALS

- ▶ Framework for Rapid Implementation of Data Mining Engines
- ▶ The ability to rapidly prototype a high-performance mining implementation
- ▶ Distributed memory parallelization
- ▶ Shared memory parallelization
- ▶ Ability to process disk-resident datasets
- ▶ Only modest modifications to a sequential implementation for the above three
- ▶ Developed 2001-2003 at Ohio State

FREEERIDE – Technical Basis

- ▶ Popular data mining algorithms have a common canonical loop
- ▶ Generalized Reduction
- ▶ Can be used as the basis for supporting a common API
- ▶ Demonstrated for Popular Data Mining and Scientific Data Processing Applications

```
While( ) {  
    forall (data instances d) {  
        I = process(d)  
        R(I) = R(I) op f(d)  
    }  
    .....  
}
```

Comparing Processing Structure

- ▶ Similar, but with subtle differences

```
{* Outer Sequential Loop *}
While() {
  {* Reduction Loop *}
  Foreach(element e) {
    (i, val) = Process(e) ;
    RObj(i) = Reduce(RObj(i),val) ;
  }
  Global Reduction to Combine RObj
}
```

```
{* Outer Sequential Loop *}
While() {
  {* Reduction Loop *}
  Foreach(element e) {
    (i, val) = Process(e) ;
  }
  Sort (i,val) pairs using i
  Reduce to compute each RObj(i)
}
```

Processing Structure: FREERIDE (left) and Map-Reduce (right)

Observations on Processing Structure

- ▶ Map-Reduce is based on functional idea
 - ▶ Do not maintain state
- ▶ This can lead to sorting overheads
- ▶ FREERIDE API is based on a programmer managed reduction object
 - ▶ Not as 'clean'
 - ▶ But, avoids sorting
 - ▶ Can also help shared memory parallelization
 - ▶ Helps better fault-recovery

An Example

► KMeans pseudo-code using FREERIDE

```
FREERIDE (k - means)
void Kmeans :: reduction(void *block) {
  for each point  $\in$  block{
    for (i = 0; i < k; i++) {
      dis = distance(point, i);
      if (dis < min) {
        min = dis;
        min_index = i;
      }
    }
    objectID = clusterID[min_index];
    for (j = 0; j < ndim; j++) {
      reductionobject->Accumulate(objectID, j, point[j]);
      reductionobject->Accumulate(objectID, ndim, 1);
      reductionobject->Accumulate(objectID, ndim + 1,
                                  dis);
    }
  }
}

int Kmeans :: finalize() {
  for (i = 0; i < k; i++) {
    objectID = clusterID[i];
    count = (*reductionobject)(objectID, ndim);
    for (j = 0; j < ndim; j++) {
      clusters[i][j] += (*reductionobject)(objectID, j)
                        / (count + 1);
    }
    totaldistance += (*reductionobject)(objectID,
                                         ndim + 1);
  }
}
```

Example – Now with Hadoop

► KMeans pseudo-code using Hadoop

```
HADOOP (k – means)


public void map(LongWritable key, Text point){
    minDistance = Double.MAX_DISTANCE;
    for (i = 0; i < k; i++){
        if (distance(point, clusters[i]) < minDistance){
            minDistance = distance(point, clusters[i]);
            currentCluster = i;
        }
    }
    EmitIntermediate(currentCluster, point);
}

public void reduce(IntWritable key,
                  Iterator < PointWritable > points){
    num = 0;
    while (points.hasNext()){
        PointWritable currentPoint = points.next();
        num += currentPoint.get_Num();
        for (i = 0; i < dim; i++){
            sum[i] += currentPoint.point[i];
        }
    }
    for (i = 0; i < dimension; i++){
        mean[i] = sum[i]/num;
    }
    Emit(key, mean);
}
```

Experiment Design

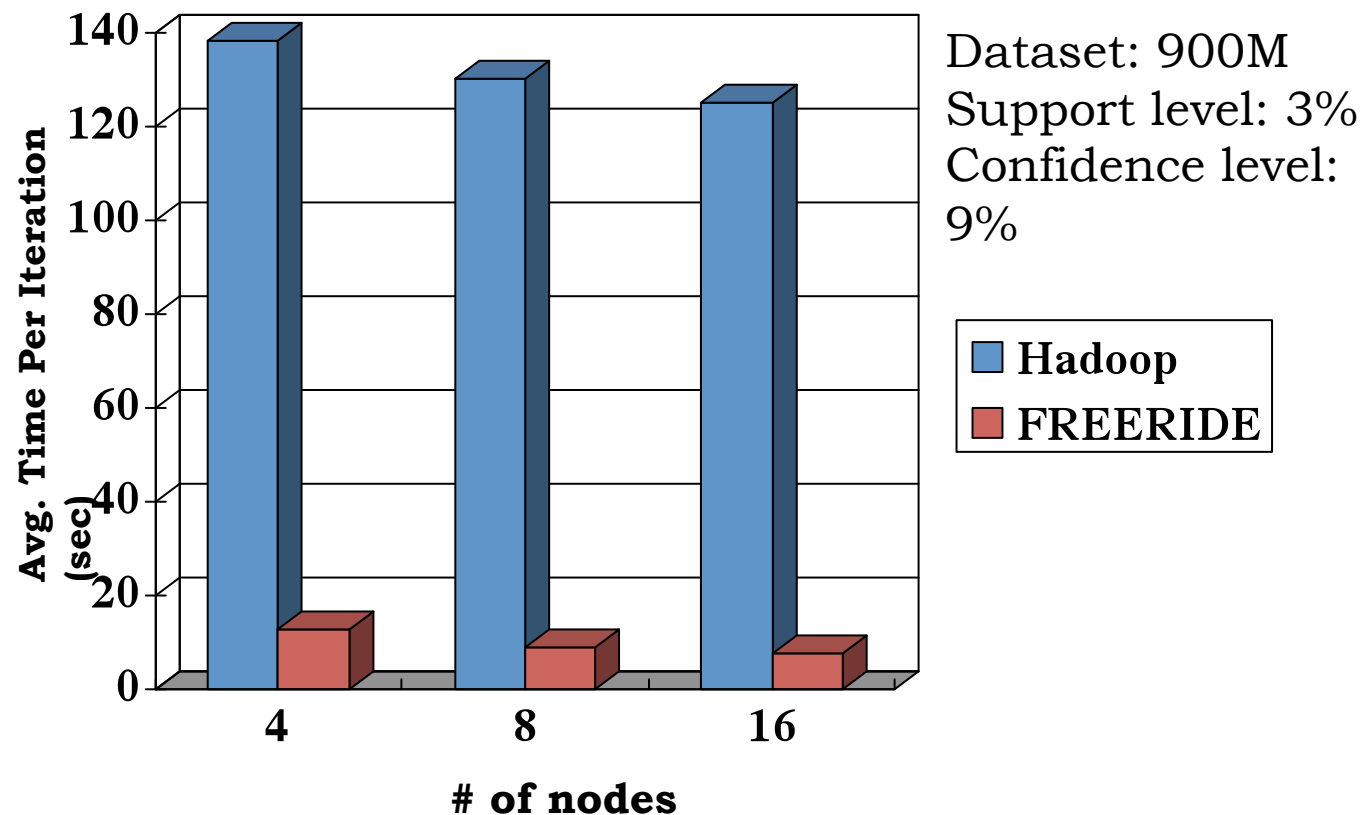
- ▶ Tuning parameters in Hadoop
 - ▶ Input Split size
 - ▶ Max number of concurrent map tasks per node
 - ▶ Number of reduce tasks
- ▶ For comparison, we used four applications
 - ▶ Data Mining: KMeans, KNN, Apriori
 - ▶ Simple data scan application: Wordcount
- ▶ Experiments on a multi-core cluster
 - ▶ 8 cores per node (8 map tasks)

► KMeans: varying # of nodes




Results – Data Mining (II)

► Apriori: varying # of nodes

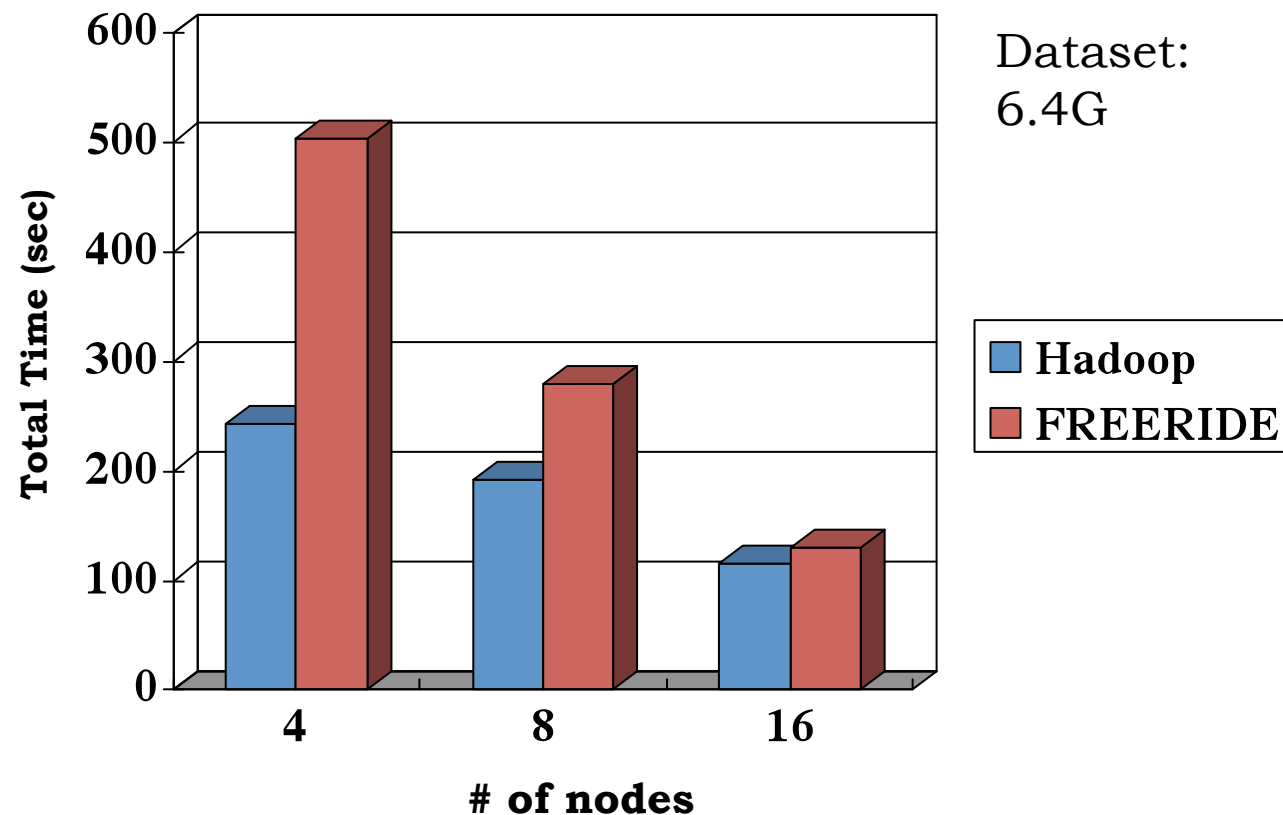


► KNN: varying # of nodes




Results – Datacenter-like Application

- ▶ Wordcount: varying # of nodes



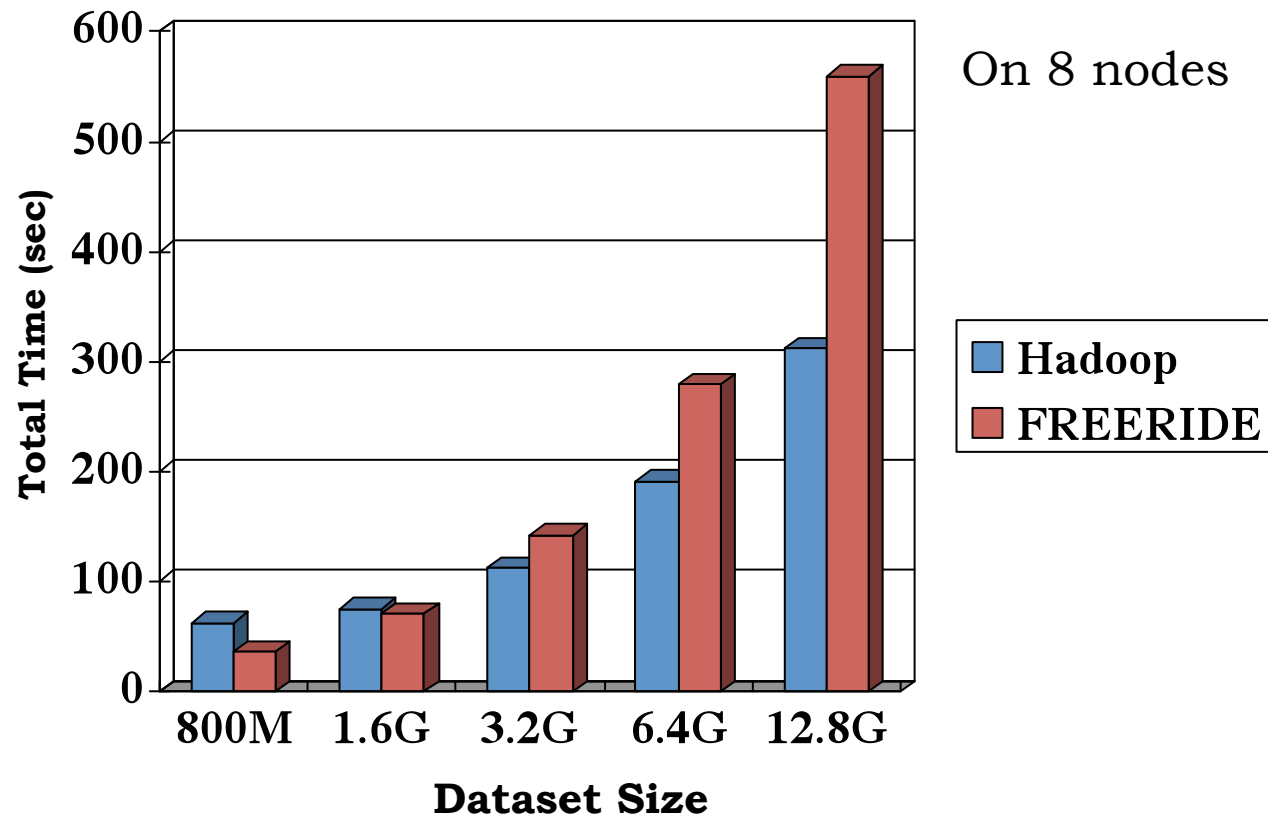
Scalability Comparison

► KMeans: varying dataset size



Scalability – Word Count

- ▶ Wordcount: varying dataset size

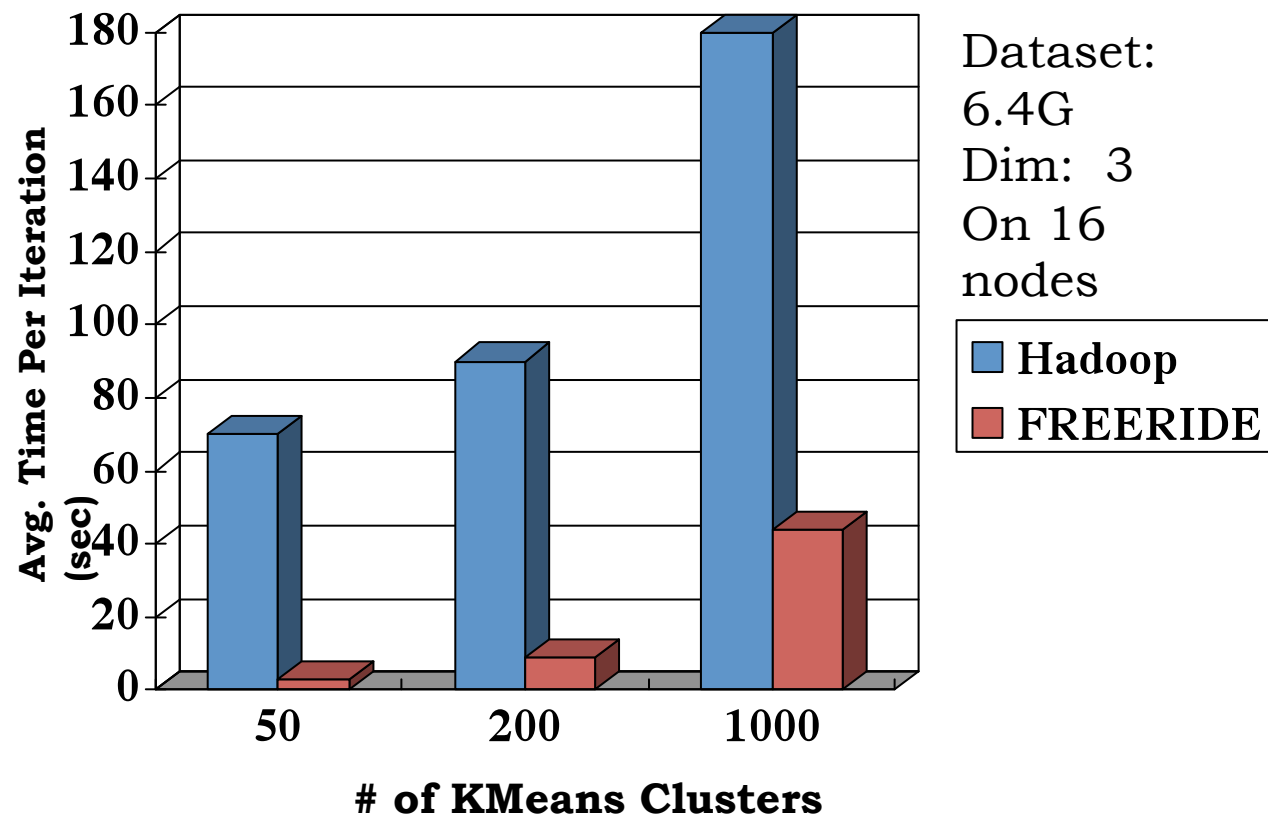


Overhead Breakdown

- ▶ Four components affecting the hadoop performance
 - ▶ Initialization cost
 - ▶ I/O time
 - ▶ Sorting/grouping/shuffling
 - ▶ Computation time
- ▶ What is the relative impact of each ?
- ▶ An Experiment with k-means

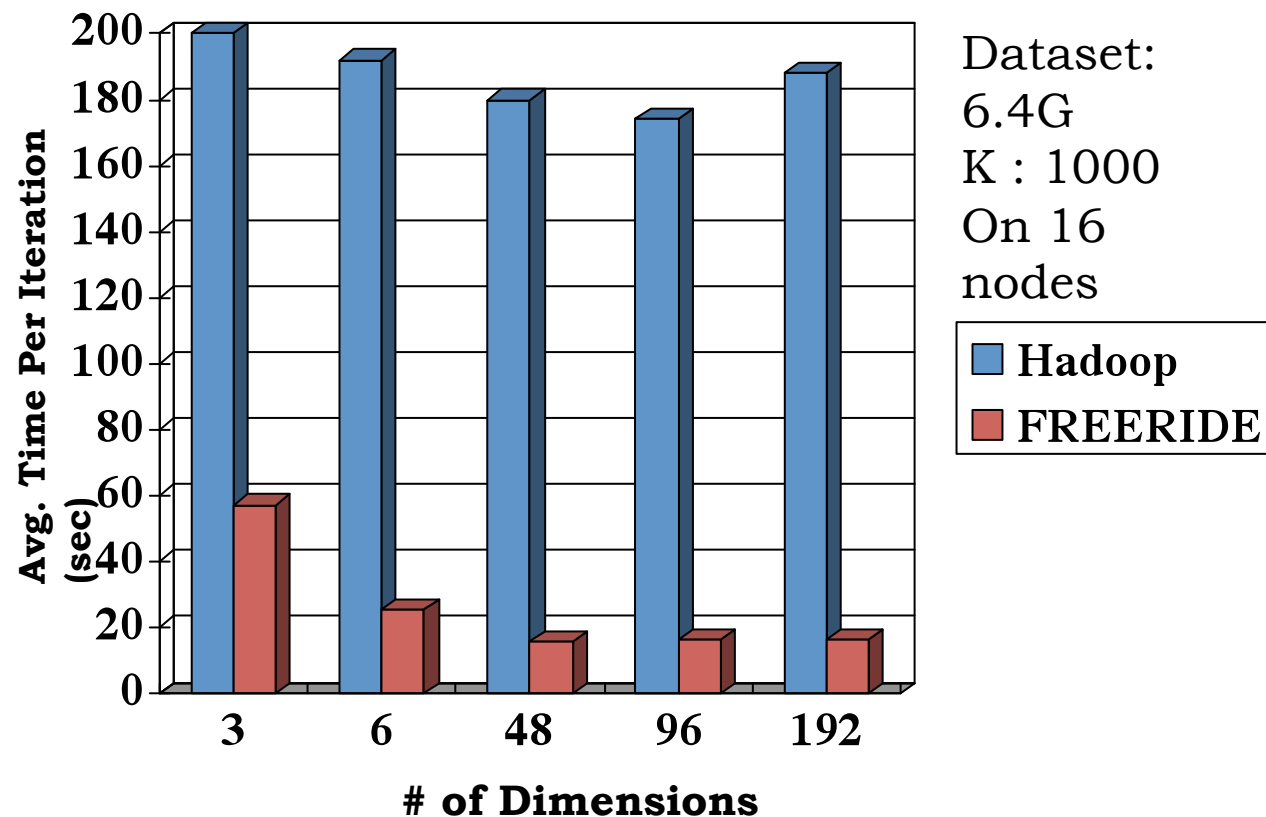
Analysis with K-means

- Varying the number of clusters (k)



Analysis with K-means (II)

- Varying the number of dimensions



Observations

- ▶ Initialization costs and limited I/O bandwidth of HDFS are significant in Hadoop
- ▶ Sorting is also an important limiting factor for Hadoop's performance

Related Work

- ▶ Lots of work on improving and generalizing it...
 - ▶ Dryad/DryadLINQ from Microsoft
 - ▶ Sawzall from Google
 - ▶ Pig/Map-Reduce-Merge from Yahoo!
 - ▶ ...
- ▶ Address MapReduce limitations
 - ▶ One input, two-stage data flow is extremely rigid
 - ▶ Only two high-level primitives

Conclusions

- FREERIDE outperformed Hadoop for three data mining applications
- MapReduce may be not quite suitable for data mining applications
- Alternative API can be supported by 'map-reduce' implementations
 - Current work on implementing different API for Phoenix
 - Should be release in next 2 months